

Agenda for today:

- Using Multi-threaded programming
- Message passing:
 - Parallel primitives
 - Reduce
- Map/Reduce

Upcoming Deadlines



Python Programming 2

Not available until Apr 1 | Due Apr 15 at 10am | -/5 pts



Preparation for Lecture 4/15

Not available until Apr 1 | Due Apr 15 at 10am



Assignment 2

Available until Apr 20 | Due Apr 17 at 10am | -/14 pts

DIT 873 & DAT346

Techniques for Large-Scale Data

Lecture 4

Recap

Monte Carlo: Computing Pi

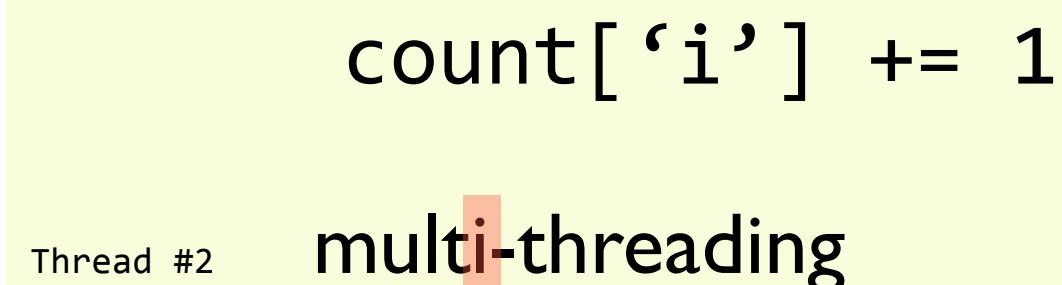
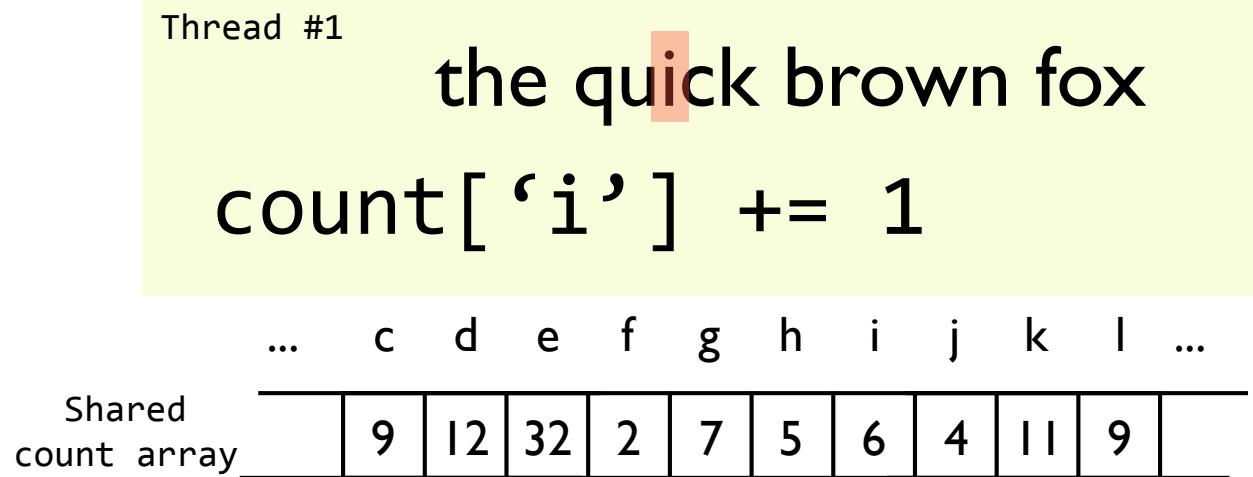
```
p = multiprocessing.Pool(args.workers)
s = p.map(sample_pi, [n]*args.workers)
```

- List of inputs
- Apply the same function to each list element
- Replace list elements by functions output

Here:
only communication
providing argument &
returning result!

Threads writing to memory

Multi-threaded character counting:



Possible Execution

1. Thread 1 reads current value 6 from count array into register
2. Thread 1 increments register
3. Thread 2 reads current value 6 from count array
4. Thread 1 writes incremented value 7
5. Thread 2 increments register
6. Thread 2 writes incremented value 7

⇒ `count['i']` incorrect

See Sec. 3.8.3.1 Rauber & Rünger (2013)

Threads writing to memory

Multi-threaded character counting:

Thread #1

```
acquire lock count['i']
count['i'] += 1
release lock count['i']
```

... c d e f g h i j k l ...

Shared
count array

	9	12	32	2	7	5	6	4	11	9	
--	---	----	----	---	---	---	---	---	----	---	--

Thread #2

```
acquire lock count['i']
count['i'] += 1
release lock count['i']
```

Possible Execution

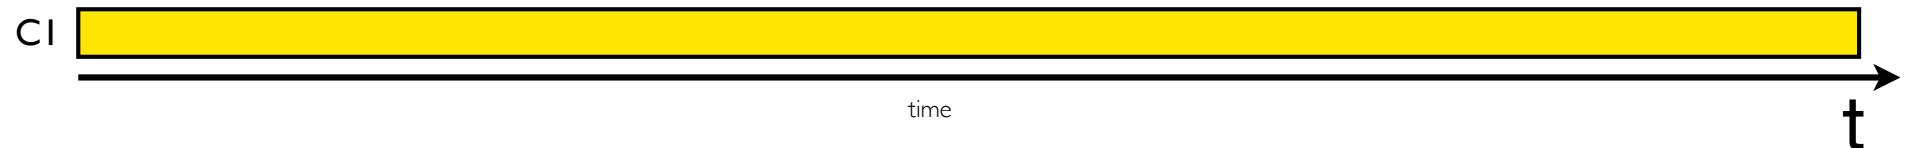
1. Thread 1 locks count[i]
2. Thread 2 tries to lock count[i] and waits
3. Thread 1 reads current value 6 from count array into register
4. Thread 1 increments register
5. Thread 1 writes incremented value 7
6. Thread 1 releases lock
7. Thread 2 locks count[i]
8. Thread 2 reads current value 7 from count array
9. Thread 2 increments register
10. Thread 2 writes incremented value 8
11. Thread 2 releases lock

⇒ count['i'] correct

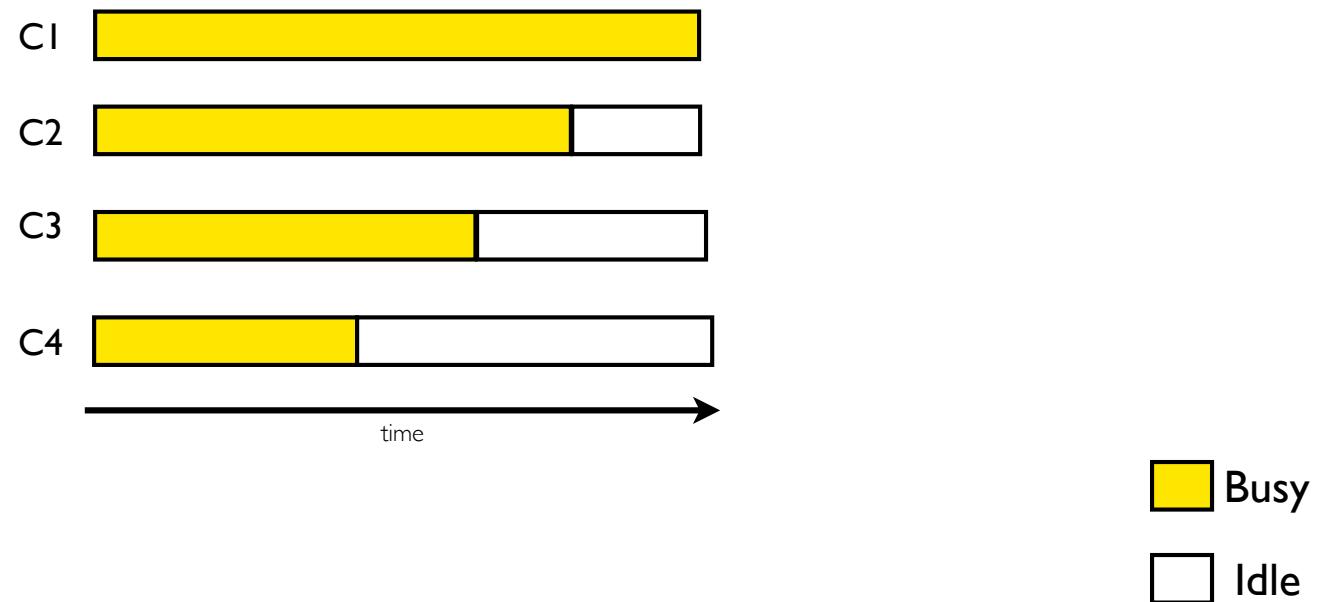
See Sec. 3.8.3.1 Rauber & Rünger (2013)

Load Balancing: k-means

Single core



Serial and parallel sections on 4 cores



Under the hood: Python arrays

- Numpy, Python array module expose the underlying C-arrays
- This allows sharing them in memory
- See `multiprocessing.RawArray`

<https://docs.python.org/2/library/array.html>

Continuing ...

Threads in Python for the lazy

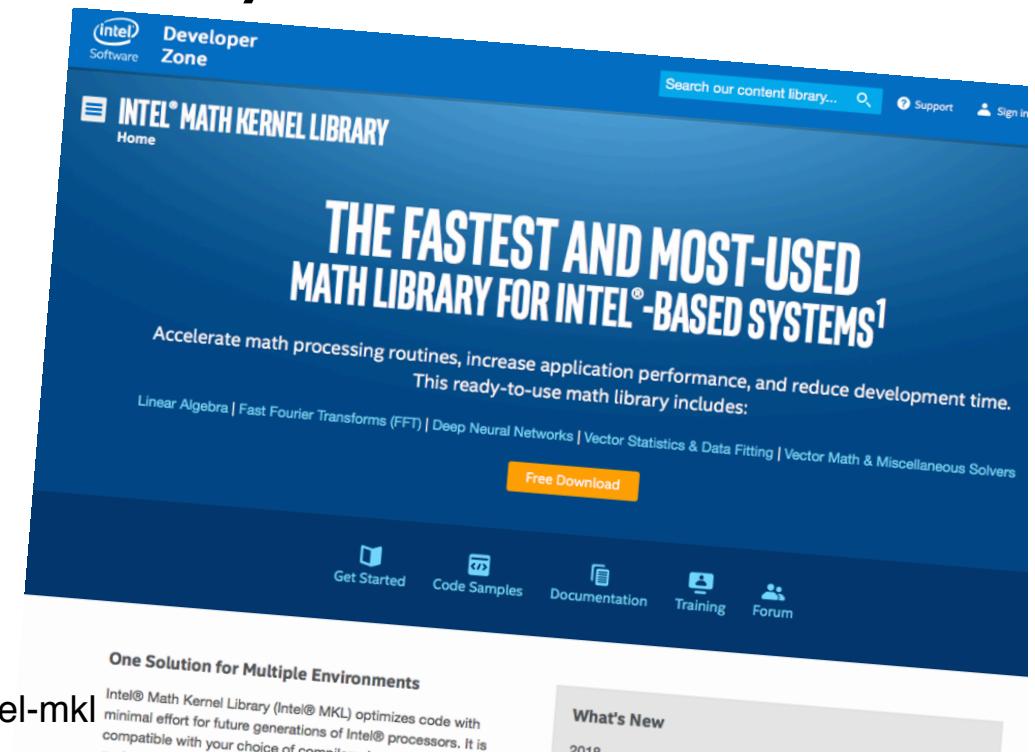
- Intel Math Kernel Library (MKL) is multi-threaded
- used by conda install of numpy
- mostly for linalg, some basic array operations including sum

```
import mkl
import numpy as np

mkl.set_num_threads(4)
x = np.ones((N,N))
y = np.ones((N,N))
np.dot(x,y)
```

<https://software.intel.com/en-us/mkl>

<https://software.intel.com/en-us/articles/numpy-scipy-with-intel-mkl>



Pitfalls

- Multiprocessing + threaded MKL could lead to reduced performance

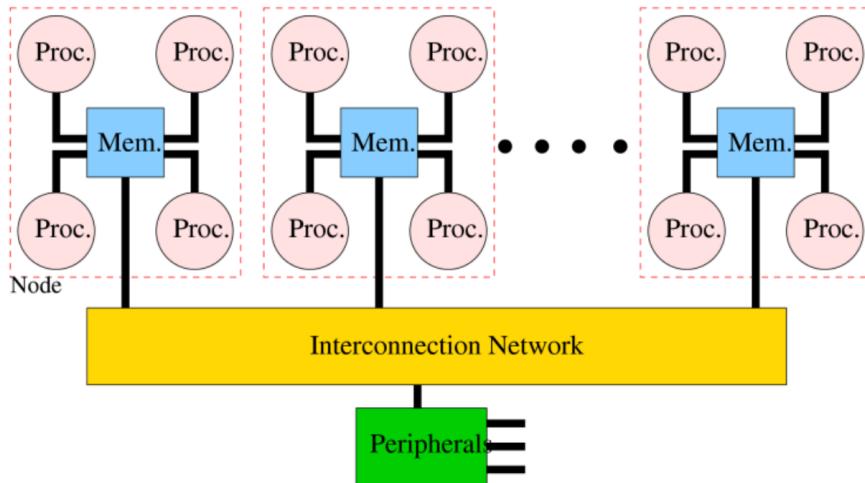
#MP processes + #ML threads > #cores

Parallel programming models

Message Passing

<https://www.scipy.org/topical-software.html#parallel-and-distributed-programming>

Parallel programming models: message passing



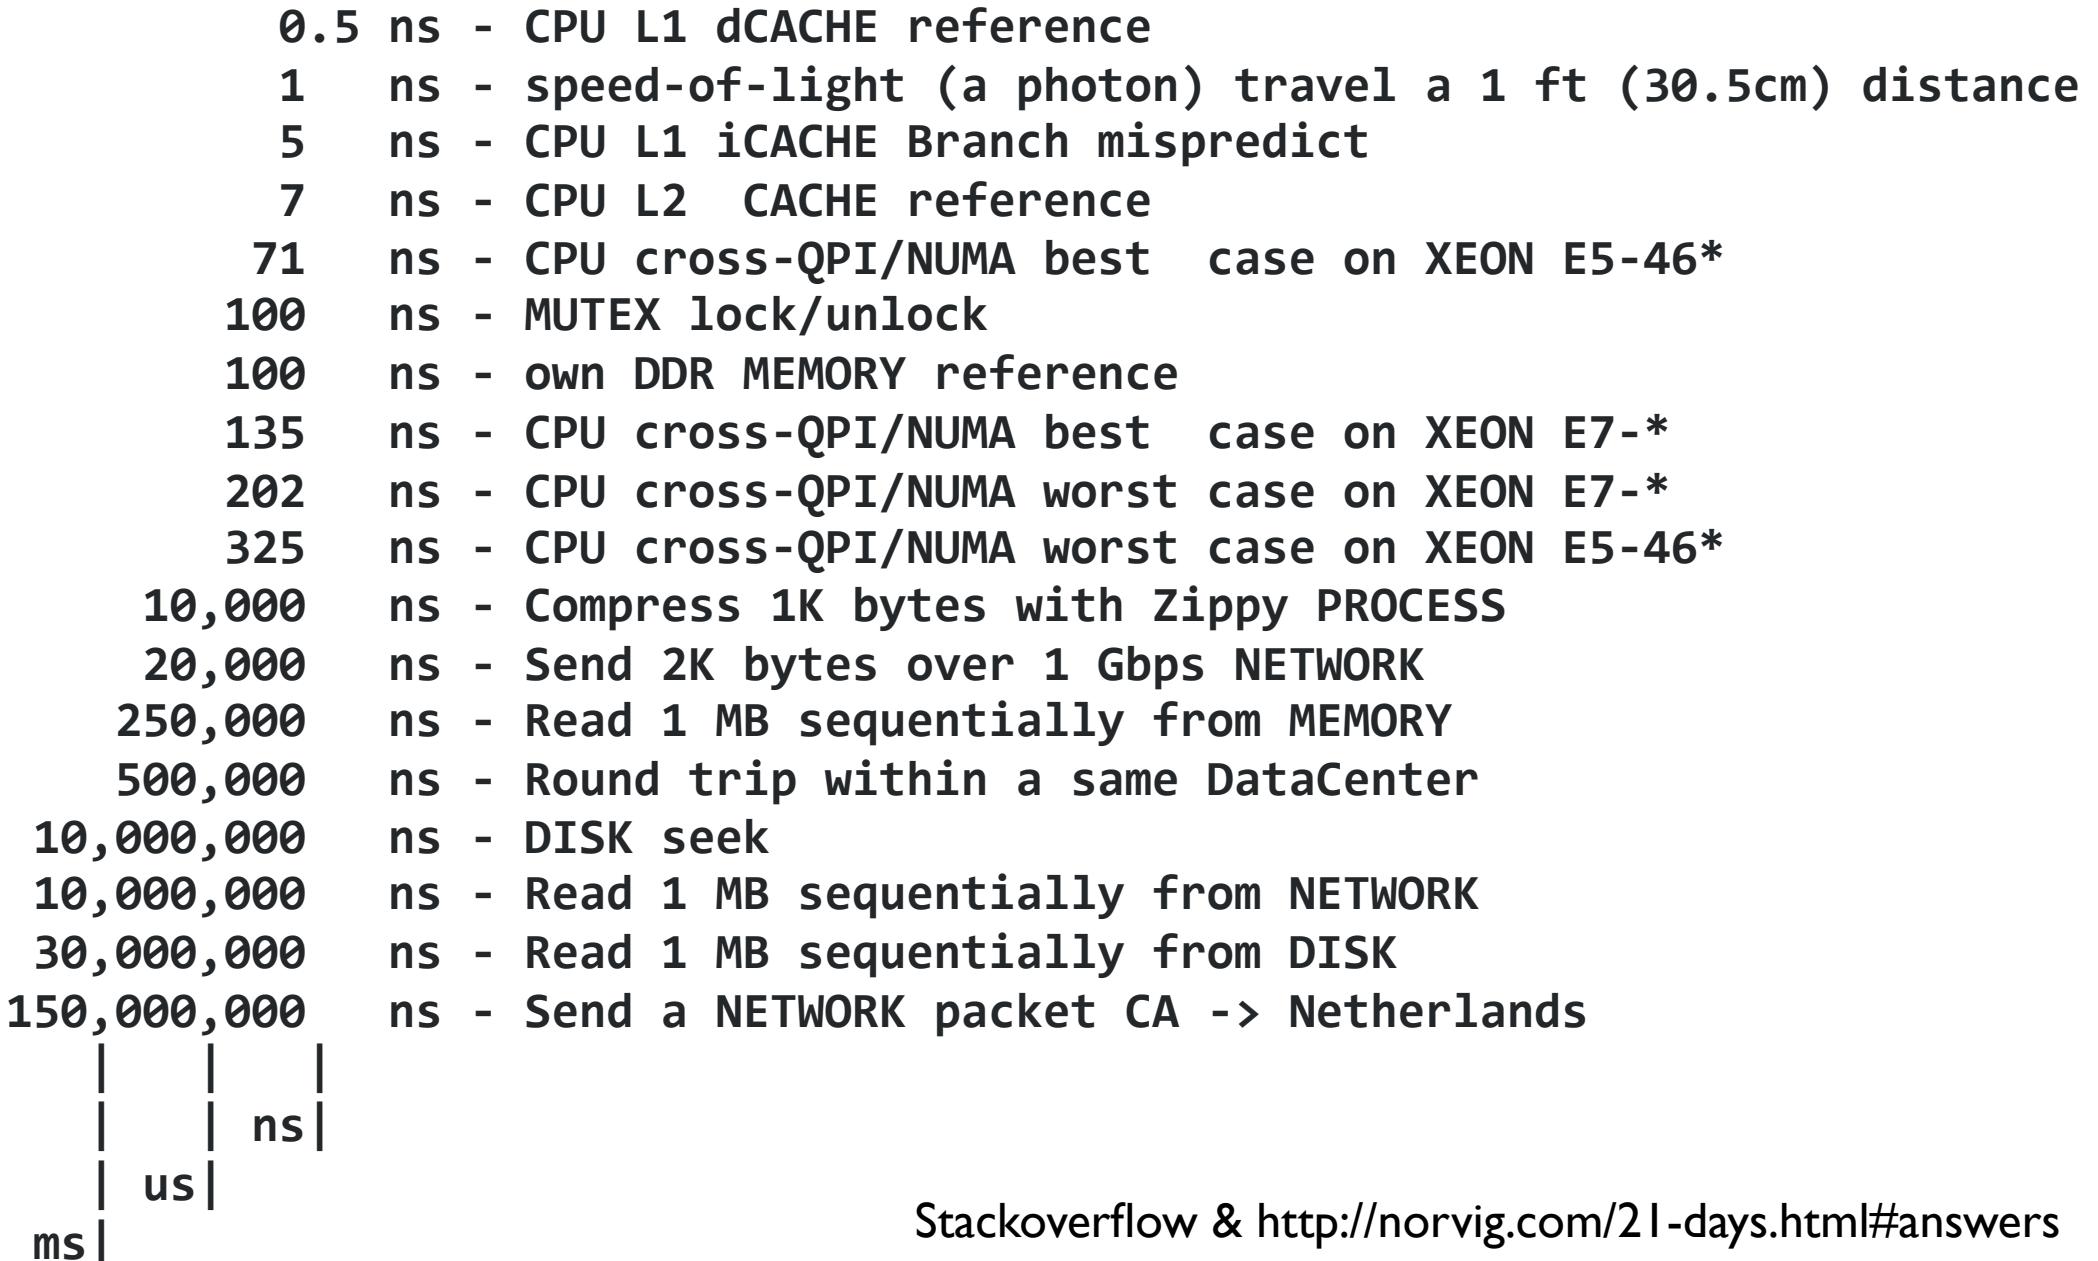
- Communication incl. exchange of data via interconnect. Infiniband 5 μ s (5000 ns + 2 x RAM access)

- Lower throughput for sharing large amount of data vs. RAM
- Explicit parallelization needed
- MPI is library standard

Passing a message over Infiniband takes about 5000 ns. This is how many times slower than a memory access?

A	B	C	D	E
I-10	II-50	5I-250	25I-1000	>1000

Latency



MPI: Message Passing Interface

- See <https://www.open-mpi.org/>
- Runs on large portion of TOP500 super computers
- Mature code base: OpenMPI v1.0 released Nov 2005. MPI 1.0 standard developed 1992-1994
- Python bindings: <http://mpi4py.scipy.org/>



Open MPI v3.0.1 documentation

| Home | Support | FAQ | Search |

[About](#)
[Presentations](#)
[Open MPI Team](#)

Open MPI commands
(section 1 man pages)

**Open MPI general information
(section 7 man pages)**

[mpi_crsp](#) [orte_filem](#) [orte_snapc](#)
[pal_crs](#) [orte_hosts](#) [orte_sstore](#)

MPI API
(section 3 man pages)

uda_support	MPI_Get_library_version
te	MPI_Get_processor_name
_class	MPI_Get_version
_code	MPI_Graph_create
_string	MPI_Graph_get
	MPI_Graph_map
	MPI_Graph_neighbors
	MPI_Graph_neighbors_count

Open MPI v3.0.1 documentation

MPI_Comm_delete_attr	MPI_Info_create	MPI_Type_size	shmem_int_swap
MPI_Comm_disconnect	MPI_Info_delete	MPI_Type_size_x	shmem_int_wait
MPI_Comm_dup	MPI_Info_dup	MPI_Type_struct	shmem_int_wait_until
MPI_Comm_dup_with_info	MPI_Info_env	MPI_Type_ub	shmem_int_xor_to_all
MPI_Comm_f2c	MPI_Info_f2c	MPI_Type_vector	shmem_iout128
MPI_Comm_free	MPI_Info_free	MPI_Unpack	shmem_iout32
MPI_Comm_free_keyval	MPI_Info_get	MPI_Unpack_external	shmem_iout64
MPI_Comm_get_attr	MPI_Info_get_nkeys	MPI_Unpublish_name	shmem_long_add
MPI_Comm_get_errhandler	MPI_Info_get_ntkkey	MPI_Wait	shmem_long_and_to_all
MPI_Comm_get_info	MPI_Info_get_valuenl	MPI_Waitall	shmem_long_cswap
MPI_Comm_get_name	MPI_Info_set	MPI_Waitany	shmem_long_fadd
MPI_Comm_get_parent	MPI_Init	MPI_Waitsome	shmem_long_fetch
MPI_Comm_group	MPI_Intr_thread	MPI_Win_allocate	shmem_long_finc
MPI_Comm_idup	MPI_Initialized	MPI_Win_allocate_shared	shmem_long_q
MPI_Comm_ipn	MPI_Intercomm_create	MPI_Win_attach	shmem_long_qet
MPI_Comm_rank	MPI_Intercomm_merge	MPI_Win_c2f	shmem_long_qet_nbi
MPI_Comm_remote_group	MPI_Iprobe	MPI_Win_call_errhandler	shmem_long_qet_nbi
MPI_Comm_remote_size	MPI_Irecv	MPI_Win_complete	shmem_long_inc
MPI_Comm_set_attr	MPI_Ireduce	MPI_Win_create	shmem_long_iput
MPI_Comm_set_errhandler	MPI_Ireduce_scatter	MPI_Win_create_dynamic	shmem_long_max_to_all
MPI_Comm_set_info	MPI_Ireduce_scatter_block	MPI_Win_create_errhandler	shmem_long_min_to_all
MPI_Comm_set_name	MPI_Isend	MPI_Win_create_keyval	shmem_long_nothing
MPI_Comm_size	MPI_Is_thread_main	MPI_Win_delete_attr	shmem_long_p
MPI_Comm_spawn	MPI_Iscan	MPI_Win_detach	shmem_long_prod_to_all
MPI_Comm_spawn_multiple	MPI_Iscatter	MPI_Win_f2c	shmem_long_put
MPI_Comm_split	MPI_Iscatterv	MPI_Win_fence	shmem_long_put_nbi
MPI_Comm_split_type	MPI_Isend	MPI_Win_flush	shmem_long_set
MPI_Comm_test_inter	MPI_Issend	MPI_Win_flush_all	shmem_long_sum_to_all
MPI_Compare_and_swap	MPI_Keyval_create	MPI_Win_flush_local	shmem_long_swap
MPI_Dims_create	MPI_Keyval_free	MPI_Win_flush_local_all	shmem_long_wait
MPI_Dist_graph_create	MPI_Lookup_name	MPI_Win_free	shmem_long_wait_until
MPI_Dist_graph_create_adjacent	MPI_Message_c2f	MPI_Win_free_keyval	shmem_long_xor_to_all
MPI_Dist_graph_neighbors	MPI_Message_f2c	MPI_Win_get_attr	shmem_longdouble_d
MPI_Dist_graph_neighbors_count	MPI_Mprobe	MPI_Win_get_errhandler	shmem_longdouble_det
MPI_Erhandler_create	MPI_Mrecv	MPI_Win_get_group	shmem_longdouble_iput
MPI_Erhandler_free	MPI_Neighbor_allgather	MPI_Win_get_info	shmem_longdouble_iput
MPI_Erhandler_get	MPI_Neighbor_allgatherv	MPI_Win_get_name	shmem_longdouble_iput
MPI_Erhandler_set	MPI_Neighbor_alltoall	MPI_Win_lock	shmem_longdouble_max_to_all
MPI_Error_class	MPI_Neighbor_alltoallv	MPI_Win_lock_all	shmem_longdouble_min_to_all
MPI_Error_string	MPI_Neighbor_alltoallw	MPI_Win_post	shmem_longdouble_p
MPI_Exscan	MPI_Op_c2f	MPI_Win_set_attr	shmem_longdouble_prod_to_all
MPI_Fetch_and_op	MPI_Op_commutative	MPI_Win_set_errhandler	shmem_longdouble_prod
MPI_File_f2c	MPI_Op_create	MPI_Win_set_info	shmem_longdouble_put
MPI_File_get_errhandler	MPI_Op_f2c	MPI_Win_set_name	shmem_longlong_add
MPI_File_close	MPI_Open	MPI_Win_shared_query	shmem_longlong_and_to_all
MPI_File_create_errhandler	MPI_Open_port	MPI_Win_start	shmem_longlong_cswap
MPI_File_delete	MPI_Pack	MPI_Win_sync	shmem_longlong_fadd
MPI_File_f2c	MPI_Pack_external	MPI_Win_test	shmem_longlong_fetch
MPI_File_get_amode	MPI_Pack_external_size	MPI_Win_unlock	shmem_longlong_finc
MPI_File_get_atomicity	MPI_Pack_size	MPI_Win_unlock_all	shmem_longlong_g
MPI_File_get_byte_offset	MPI_Pcontrol	MPI_Win_wait	shmem_longlong_get
MPI_File_get_errhandler	MPI_Probe	MPI_Wlock	shmem_longlong_get_nbi
MPI_File_get_group	MPI_Publish_name	MPI_Writem	shmem_longlong_ipet
MPI_File_get_info	MPI_Put	OMPI_Attrity_str	shmem_longlong_inc
MPI_File_get_position	MPI_Query_thread	OpenMPJ	shmem_longlong_iout
MPI_File_get_position_shared	MPI_Raccumulate	OpenSHMEM_my_pe	shmem_longlong_max_to_all
MPI_File_get_type_extent	MPI_Recv	num_pes	shmem_longlong_min_to_all
MPI_File_get_view	MPI_Recv_init	intro_shmem	shmem_longlong_nothing
MPI_File_iread	MPI_Reduce	shfree	shmem_longlong_p
MPI_File_iread_all	MPI_Reduce_local	shmalloc	shmem_longlong_put
MPI_File_iread_at	MPI_Reduce_scatter	shmem_addr_accessible	shmem_longlong_put_nbi
MPI_File_iread_at_all	MPI_Register_datagp	shmem_align	shmem_longlong_set
MPI_File_iread_shared	MPI_Request_c2f	shmem_alltoall32	shmem_longlong_sum_to_all
MPI_File_ivwrite	MPI_Request_f2c	shmem_alltoall64	shmem_longlong_swap
MPI_File_ivwrite_all	MPI_Request_free	shmem_alltoall32	shmem_longlong_walt
MPI_File_ivwrite_at	MPI_Request_get_status	shmem_alltoall64	shmem_longlong_walt_until
MPI_File_ivwrite_at_all	MPI_Rget	shmem_barrier	shmem_longlong_xor_to_all
MPI_File_ivwrite_shared	MPI_Rget_accumulate	shmem_barrier_all	shmem_malloc
MPI_File_open	MPI_Rput	shmem_broadcast32	shmem_my_pe
MPI_File_preallocate	MPI_Rsend	shmem_broadcast64	shmem_n_pes
MPI_File_read	MPI_Rsend_init	shmem_charg	shmem_pe_accessible
MPI_File_read_all	MPI_Scan	shmem_charg_get	shmem_ptr
MPI_File_read_all_begin	MPI_Scatter	shmem_charg_get_nbi	shmem_putch128
MPI_File_read_all_end	MPI_Scatterv	shmem_charg_p	shmem_putch128_nbi

<https://www.open-mpi.org/doc>

PI_File_read_at	MPI_Send	shmem_char_put	shmem_put16_nbi
PI_File_read_all	MPI_Send_init	shmem_char_put_nbi	shmem_put32_nbi
PI_File_read_all_begin	MPI_Sndrecv	shmem_clear_cache_inv	shmem_put32_nbi
PI_File_read_all_end	MPI_Sndrecv_replace	shmem_clear_cache_line_inv	shmem_put64_nbi
PI_File_read_ordered	MPI_Sizeof	shmem_clear_lock	shmem_put64_nbi
PI_File_read_ordered_begin	MPI_Ssend	shmem_collect32	shmem_put8_nbi
PI_File_read_ordered_end	MPI_Ssend_init	shmem_collect64	shmem_putmem
PI_File_read_shared	MPI_Start	shmem_complexd Prod_to_all	shmem_putmem_nbi
PI_File_seek	MPI_Startall	shmem_complexd Sum_to_all	shmem_quiet
PI_File_seek_shared	MPI_Status_c2f	shmem_complexd Prod_to_all	shmem_realall
PI_File_Set_atomicty	MPI_Status_F2c	shmem_complexd Sum_to_all	shmem_set Cache_inv
PI_File_Set_errhandler	MPI_Status_Set_cancelled	shmem_double_fetch	shmem_set Cache_line_inv
PI_File_Set_info	MPI_Status_Set_element	shmem_double_q	shmem_set lock
PI_File_Set_size	MPI_Status_Set_elements_x	shmem_double_get	shmem_short_and_to_all
PI_File_Set_view	MPI_T_category_changed	shmem_double_get_nbi	shmem_short_q
PI_File_sync	MPI_T_category_get_categories	shmem_double_get	shmem_short_get
PI_File_write	MPI_T_category_get_cvars	shmem_double_get	shmem_short_get_nbi
writeat_all	MPI_T_category_get_Info	shmem_double_max_to_all	shmem_short_iget
PI_File_write_all_begin	MPI_T_category_get_num	shmem_double_min_to_all	shmem_short_iput
PI_File_write_all_end	MPI_T_category_get_pvars	shmem_double_p	shmem_short_max_to_all
PI_File_write_at	MPI_T_cvar_get	shmem_double_Prod_to_all	shmem_short_min_to_all
PI_File_write_at_all	MPI_T_cvar_get_num	shmem_double_put	shmem_short_or_to_all
PI_File_write_at_all_begin	MPI_T_cvar_handle_alloc	shmem_double_put_nbi	shmem_short_p
PI_File_write_at_all_end	MPI_T_cvar_handle_free	shmem_double_set	shmem_short_prod_to_all
PI_File_write_ordered	MPI_T_cvar_read	shmem_double_sum_to_all	shmem_short_pout
PI_File_write_ordered_begin	MPI_T_cvar_write	shmem_double_swap	shmem_short_pout_nbi
PI_File_write_ordered_end	MPI_T_enum_get_info	shmem_fcollect32	shmem_short_sum_to_all
PI_File_write_shared	MPI_T_enum_get_item	shmem_fcollect64	shmem_short_wait
PI_Finalize	MPI_T_finalize	shmem_fence	shmem_short_wait_until
PI_Finalized	MPI_T_init_thread	shmem_finalize	shmem_short_xor_to_all
PI_Free_mem	MPI_T_pvar_get_info	shmem_float_fetch	shmem_swap
PI_Gather	MPI_T_pvar_get_num	shmem_float_q	shmem_test_lock
PI_Gatherv	MPI_T_pvar_handle_alloc	shmem_float_get	shmem_udflush
PI_Get	MPI_T_pvar_handle_free	shmem_float_get_nbi	shmem_wait
PI_Get_accumulate	MPI_T_pvar_read	shmem_float_iput	shmem_wait_until
PI_Get_address	MPI_T_pvar_redreset	shmem_float_max_to_all	shmemmaligan
PI_Get_count	MPI_T_pvar_reset	shmem_float_min_to_all	shmemalloc
PI_Get_elements	MPI_T_pvar_session_create	shmem_float_p	start_pes
PI_Get_elements_x	MPI_T_pvar_session_free		

[Open MPI is an Associated Project](#) of the [Software in the Public Interest](#) non-profit organization.

in MPI serve

Page last modified: 2
Mar-20
2004-2018 The Open M

```
# Run with mpiexec -n 4 python ./mpi-montecarlo.py -s 1000000
from mpi4py import MPI
import argparse
import random
from math import pi

comm = MPI.COMM_WORLD
# This script will be executed once per process
size = comm.Get_size() # Number of processes available
rank = comm.Get_rank() # Rank of the process executing this process

def compute_pi(args):
    n = int(args.steps / size)

    if rank == 0:
        steps = [n] * size
    else:
        steps = None

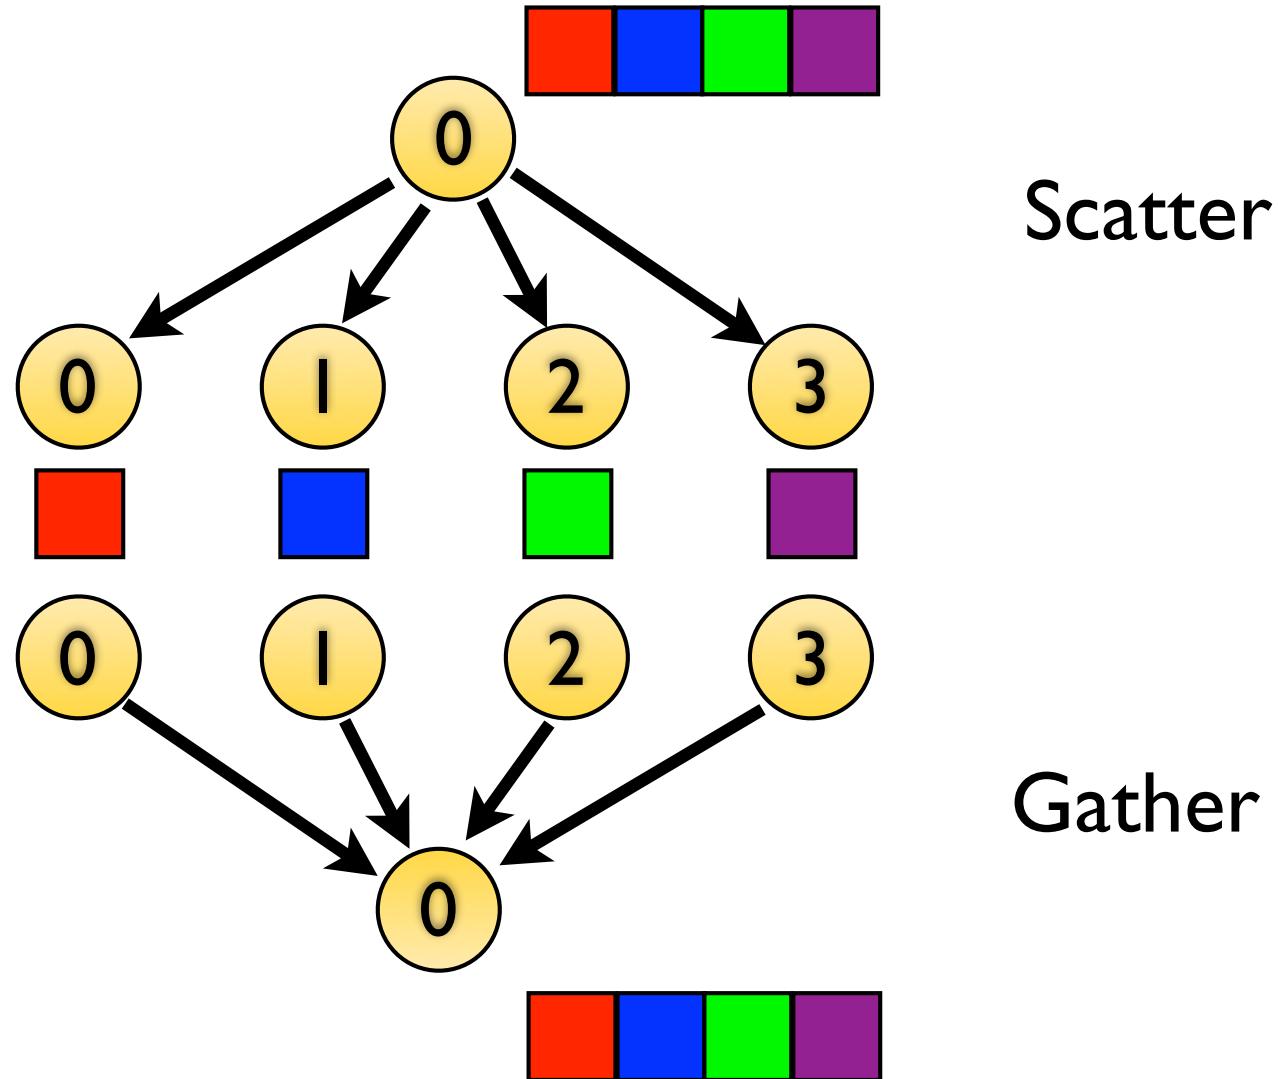
    steps = comm.scatter(steps, root=0)
    print(f"Rank={rank}, steps={steps}")

    s = sample_pi(steps)

    result = comm.gather(s, root=0)
    if rank == 0:
        n_total = n*size
        s_total = sum(result)
        pi_est = (4.0*s_total)/n_total
        print(" Steps\tSuccess\tPi est.\tError")
        print("%6d\t%7d\t%1.5f\t%1.5f" % (n_total, s_total, pi_est, pi-est))
```

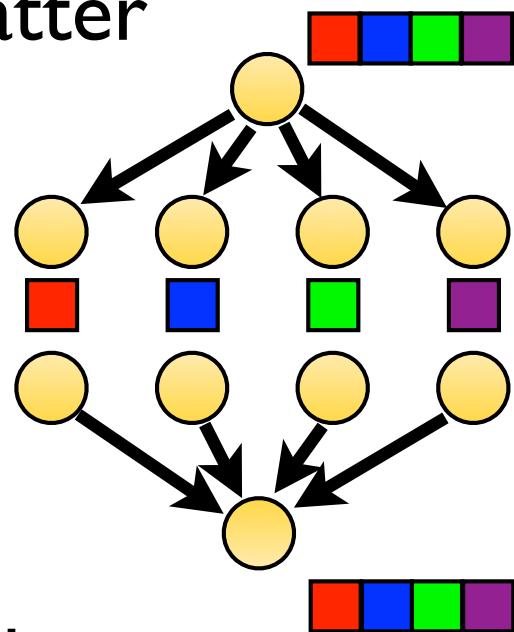
MPI: Scatter/Gather

Convention: rank 0
process is root



MPI scatter/gather vs. mp.map

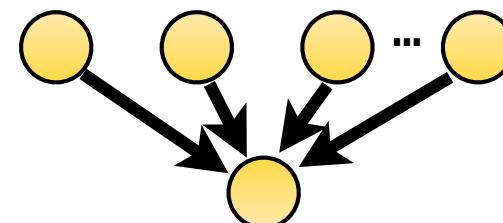
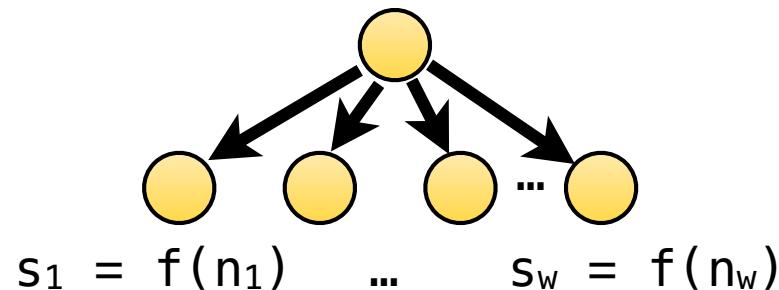
Scatter



Gather

`s = p.map(f, [n] * w)`

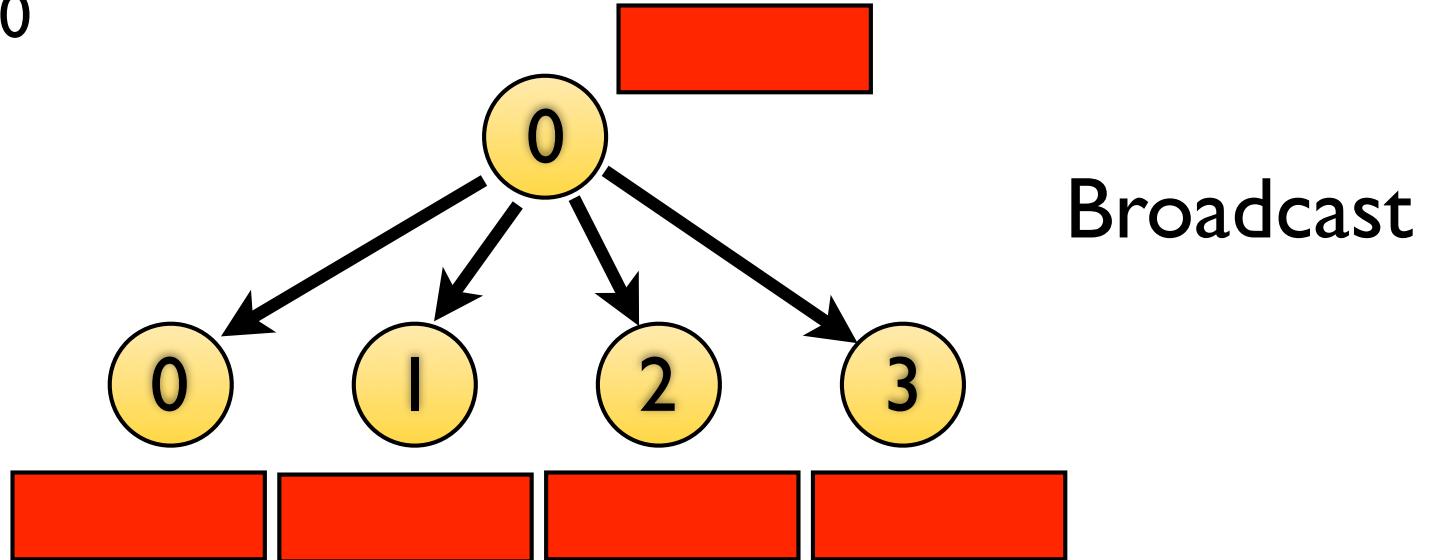
$[n_1, n_2, n_3, \dots, n_w]$



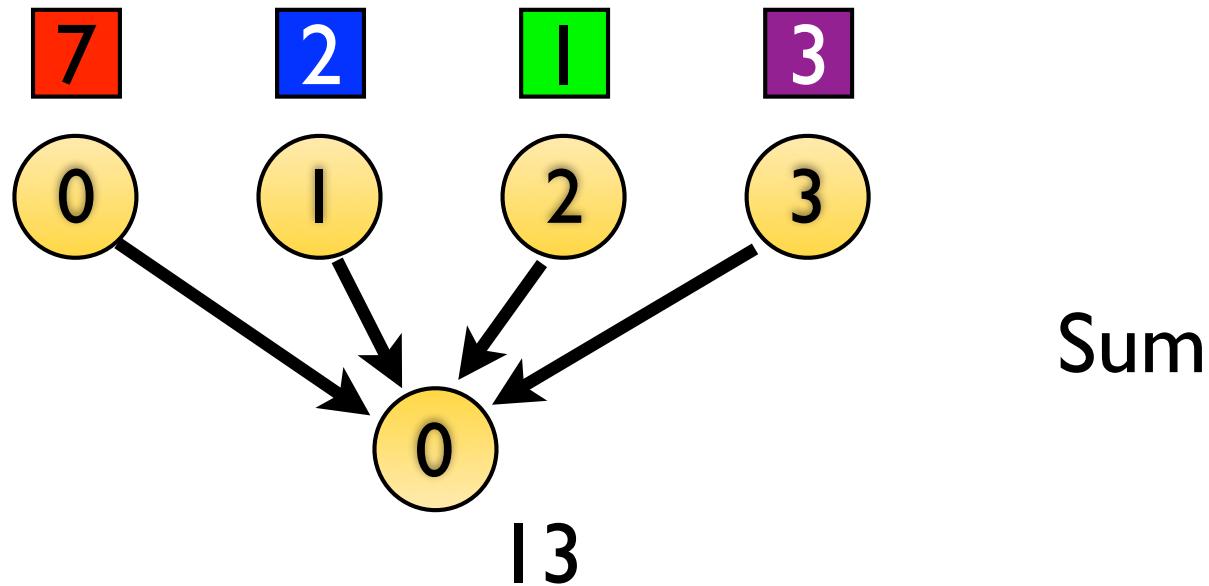
$[s_1, s_2, s_3, \dots, s_w]$

MPI Broadcast

Convention: rank 0
process is root



MPI Reduce



Reduce supports binary operations w/
associative law: Sum, Prod, Max, Min,
Argmax, Argmin, Logic operations,...

```
def compute_pi(args):
    n = int(args.steps / size)

    if rank == 0:
        steps = [n] * size
    else:
        steps = None

    steps = comm.scatter(steps, root=0)
    print(f"Rank={rank}, steps={steps}")

    s = sample_pi(steps)

    s_total = 0
    s = np.asarray(s)
    s_total = np.asarray(s_total)

    comm.Reduce(s, s_total, op=MPI.SUM, root=0)

    if rank == 0:
        n_total = n*size
        pi_est = (4.0*s_total)/n_total
        print(" Steps\tSuccess\tPi est.\tError")
        print("%6d\t%7d\t%1.5f\t%1.5f" % (n_total, s_total, pi_est, pi-est))
```



During a **gather** operation, when one message is passed from every one of the n nodes to the root, the total number of message passing operations is ...

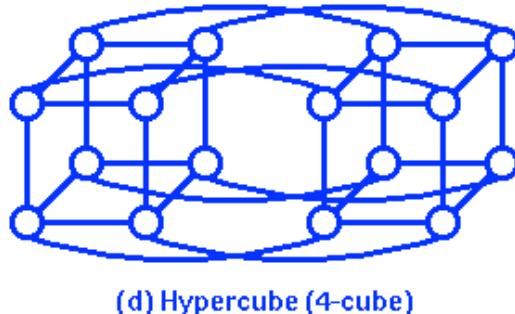
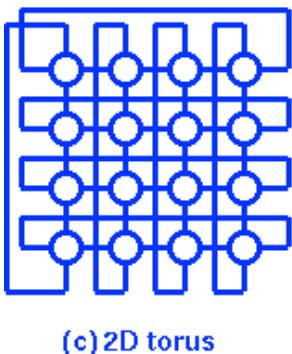
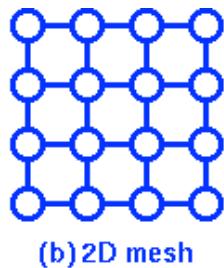
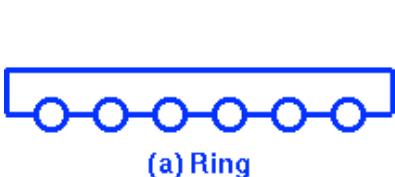
n nodes, n odd

For node farthest from root 1
second-farthest 2, 3rd-farthest 3
 $+2+3+4+5+\dots+(n-1)/2 = 1/8(n^2-1)$
using Gauss

Here a message passing operation is passing the node's own data, or a message received from another node. Assume that no data is aggregated.

A	B	C	D	E
$O(\log n)$	$O(n)$	$O(n \log n)$	$O(n^2)$	$O(n^3)$

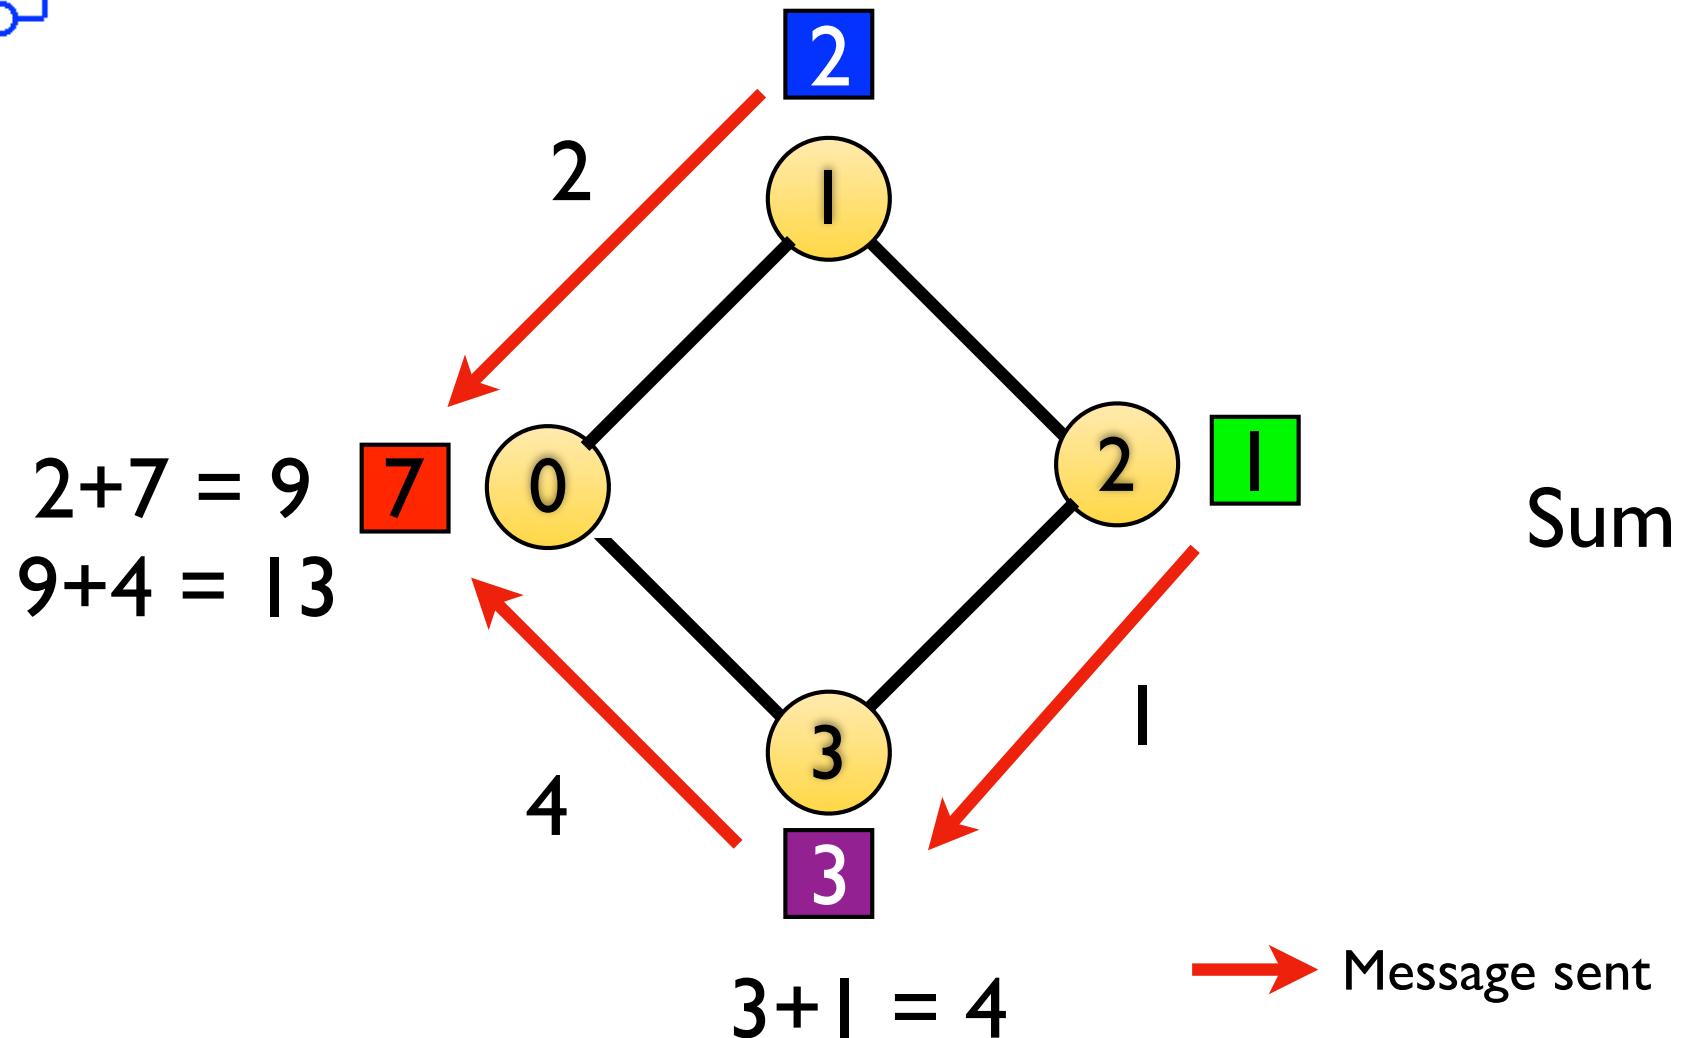
Reduce is *not just* Gather + Function call



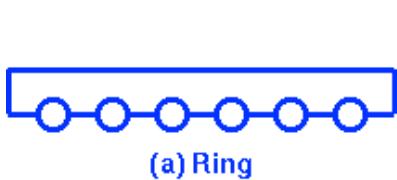
Example Interconnect topologies

- Communication Bottlenecks close to MPI root
- Substantial Latencies for thousands of cores
- Broadcast, Gather, Scatter similarly adapt to interconnect topology

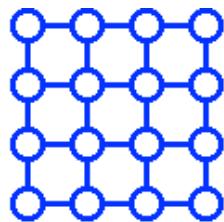
Reduce is *not* Gather + Function call



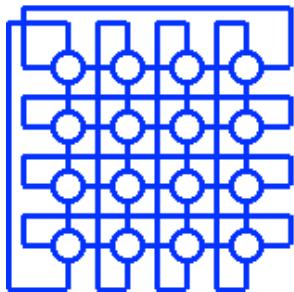
Importance of Reduce



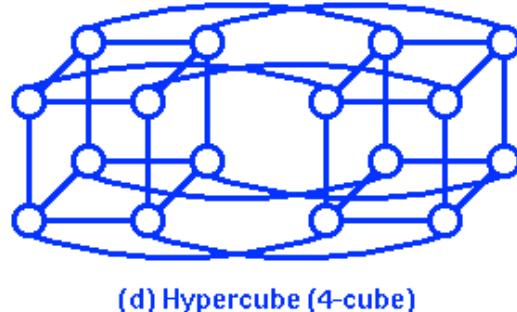
(a) Ring



(b) 2D mesh



(c) 2D torus



(d) Hypercube (4-cube)

Example Interconnect topologies

- Communication Bottlenecks close to root
- Substantial Latencies for thousands of cores
- Reduce operations yield $O(\text{nodes})$ message passing operations instead of $O(\text{nodes}^2)$
- Broadcast, Gather, Scatter similarly adapt to interconnect topology

Reflection

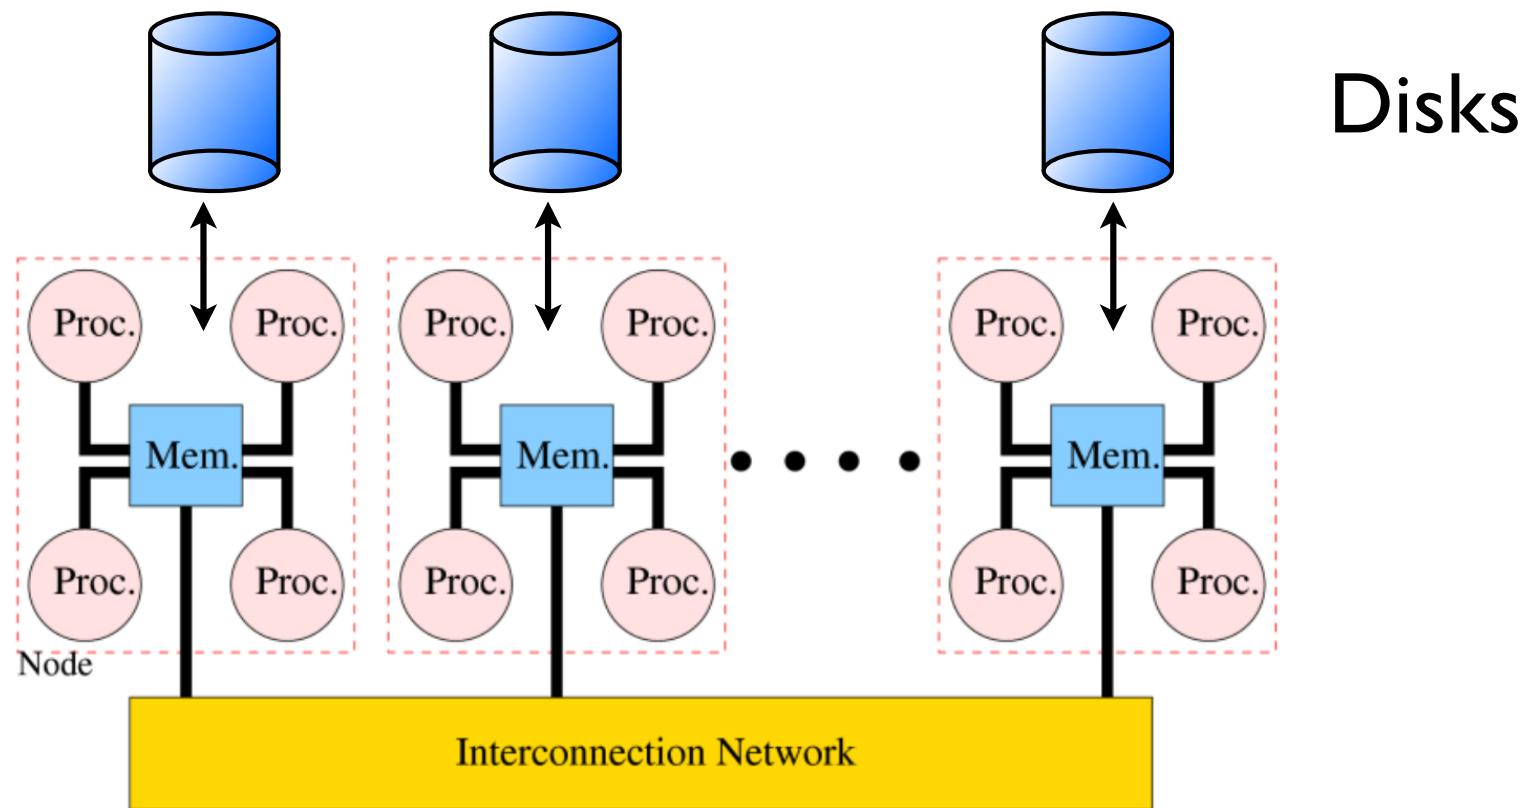
Architectures & Models

- SMP & Multithreading:
 - Shared memory, fast communication between all threads
- Distributed Memory and Message Passing
 - High latency, interconnection topology matters: limited #connections

Workload Comparison

	HPC	Data Science
CPU	Lots pf cycles Floating Point (FP) Little memory	Medium Mixed integer/FP Large memory
Data/File IO	Often very little IO	Massive IO Computations IO-bound
Tasks	Standard Tools <i>(e.g. matrix solvers)</i>	Custom tools and <i>frequently new, changing problems</i>

System architecture for data-centric computation



Aspects of data-centric computation

- Scale of Data: Parallel IO, Data-compute co-location
- Cost of computation: commodity hardware
- Reliability of systems: Design for failures
- Software is run once: support rapid development:
 - automated parallelization
 - restricted model of computation

Parallel programming models

Map Reduce

See Skiena, Data Science Design Manual, Chap 12.6

<http://hadoop.apache.org/docs/current/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>

Map/reduce assumptions

- Input files are distributed over nodes (implicit)
- All data are (key, value) tuples
- Main parallel operations are:
 - Map:
 $(key, value) \rightarrow (key_1, value_1), (key_2, value_2), \dots$
 - Reduce:
 $(key_1, [value_1, value_2, value_3, \dots]) \rightarrow (key_{new}, value_{new}), (key_{new2}, value_{new2}), \dots$

[] optional

Example: Count number of links to webpages

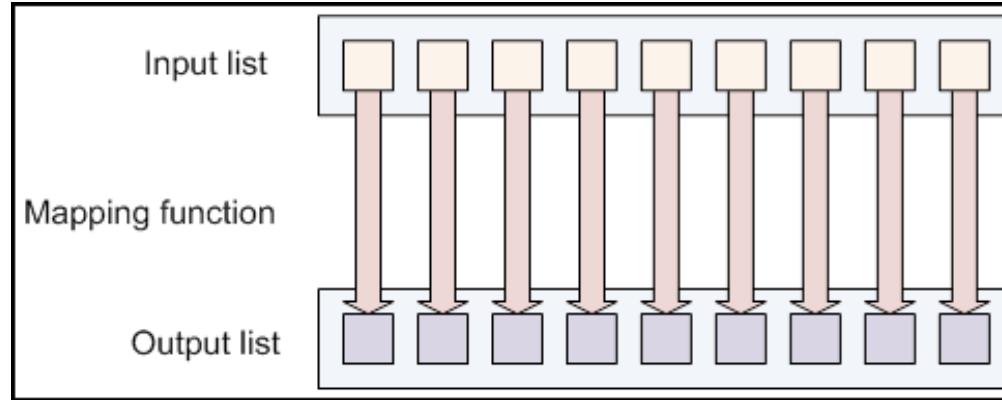
Map(DocumentURL,DocumentHTML):

For each url in DocumentHTML return (url, 1)

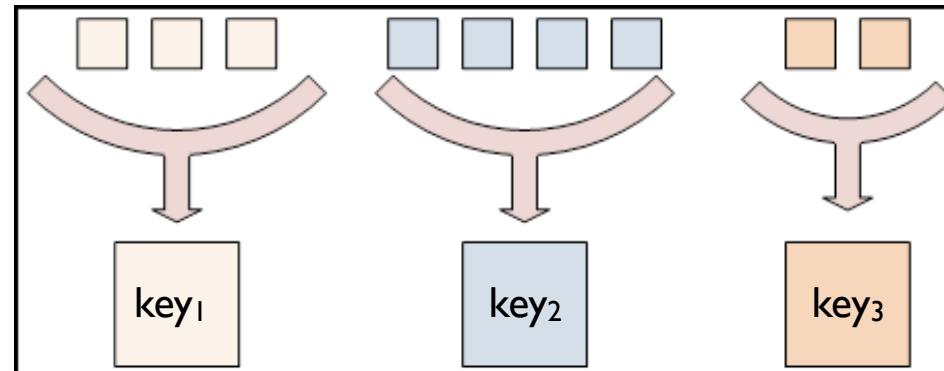
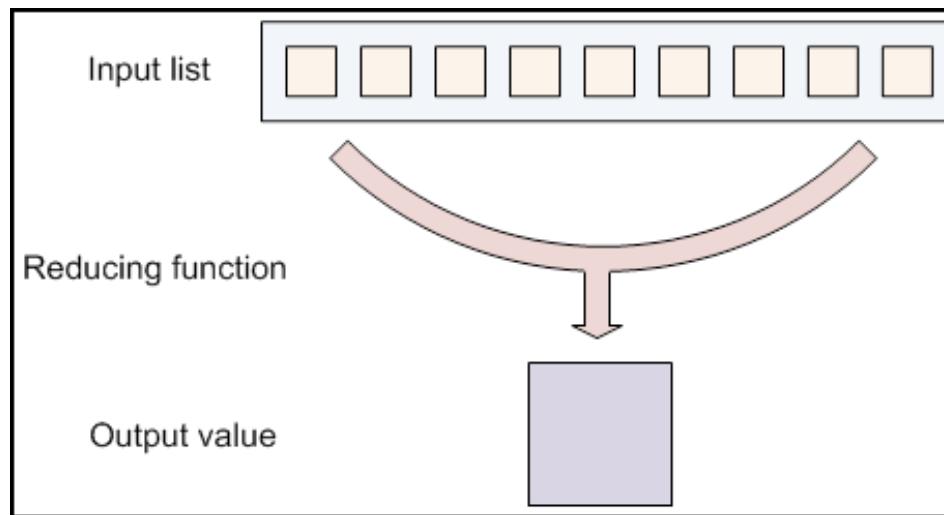
Reduce(url, counts) # counts = [count1, count2,...]

return (url, sum(counts))

Map

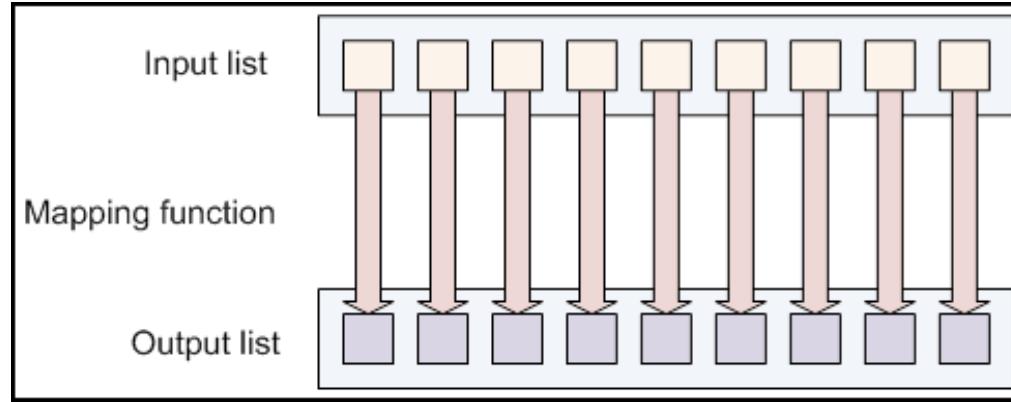


Reduce



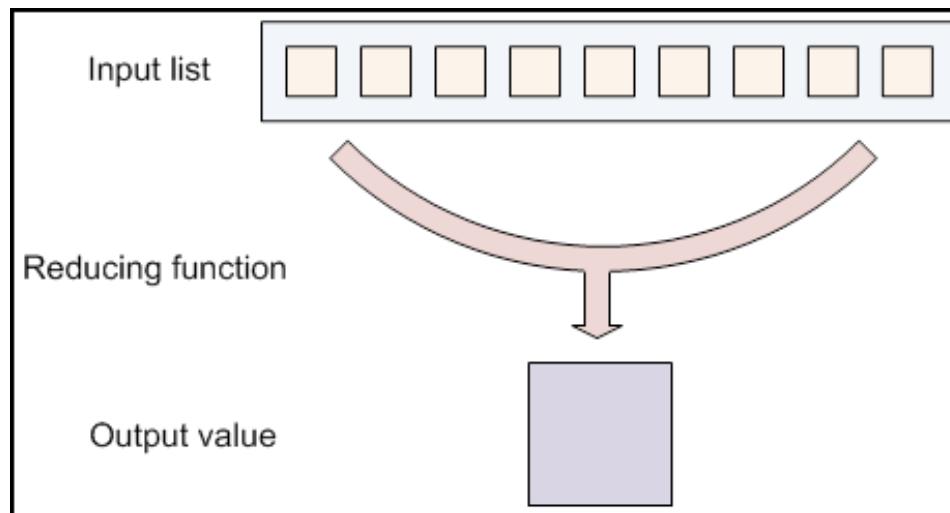
Reduce receives
input aggregated by
key

Map

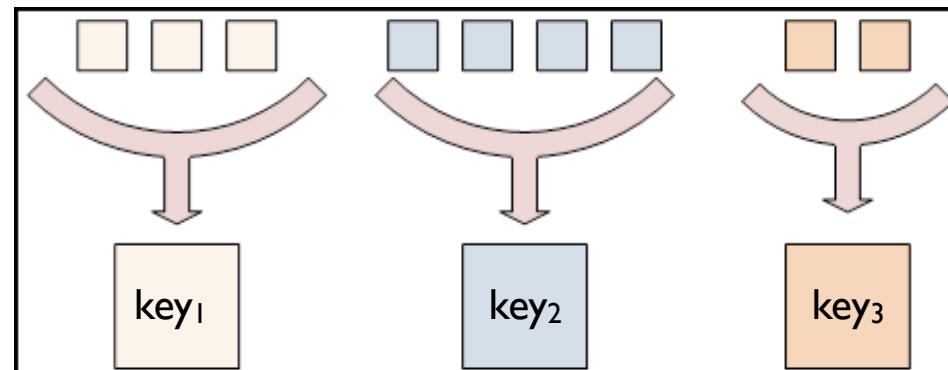


Each map call only needs its input; no assertions about other inputs, execution order, assignment to nodes

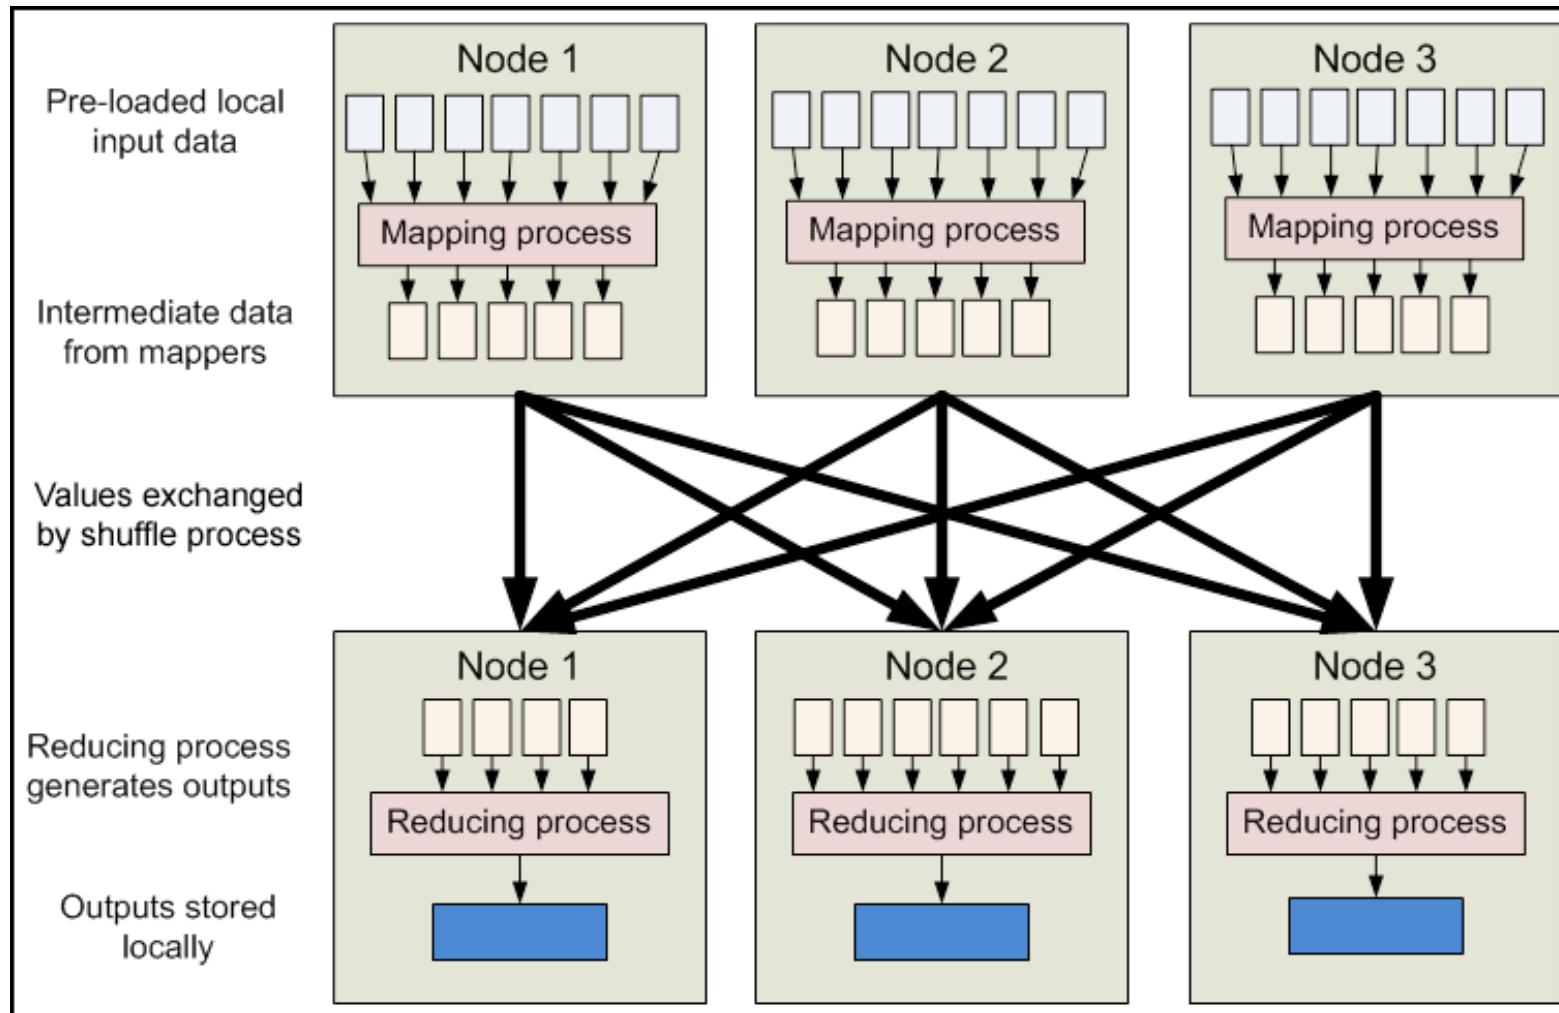
Reduce



Each reduce call receives all values for its key; no assertions about other keys, execution order, assignment to nodes



Map/Reduce Parallelism



MapReduce vs multi-threading and MPI

- No concept of process, thread nor hardware
- No control over data flow
- Reduce obtains all values with the same key
- *Substantial effort behind the scenes*

MrJob: Map/Reduce in Python

- <https://github.com/Yelp/mrjob>
- <https://mrjob.readthedocs.io/en/stable/>
- Pure Python
- Runs locally and in the cloud

```
""" Find duplicate keys in a file containing lines consisting of
    key,value
"""
from mrjob.job import MRJob

class FindDuplicates(MRJob):

    def mapper(self, _, line):
        key, value = line.split(',')
        yield (key, 1)

    def reducer(self, key, counts):
        s = sum(counts)
        if s > 1:
            yield (key, s)

if __name__ == '__main__':
    FindDuplicates.run()
```

Implementing wordcount with Map-reduce

```
schliep@MacBookPro13:> wc mpi-summary-statistics.py
      63      197    1918 mpi-summary-statistics.py
```

```
from mrjob.job import MRJob

class WC(MRJob):

    def mapper(self, _, line):
        nr_words = len(line.split())
        nr_chars = len(line)
        yield ("words", nr_words)
        yield ("chars", nr_chars)

    def reducer(self, key, counts):
        if key == "words":
            yield ("words", sum(counts))
        elif key == "chars":
            yield ("chars", sum(counts))

if __name__ == '__main__':
    WC.run()
```

```

def mapper(self, _, line):
    nr_words = len(line.split())
    nr_chars = len(line)
    yield ("words", nr_words)
    yield ("chars", nr_chars)

```

$n = \# \text{ lines}$

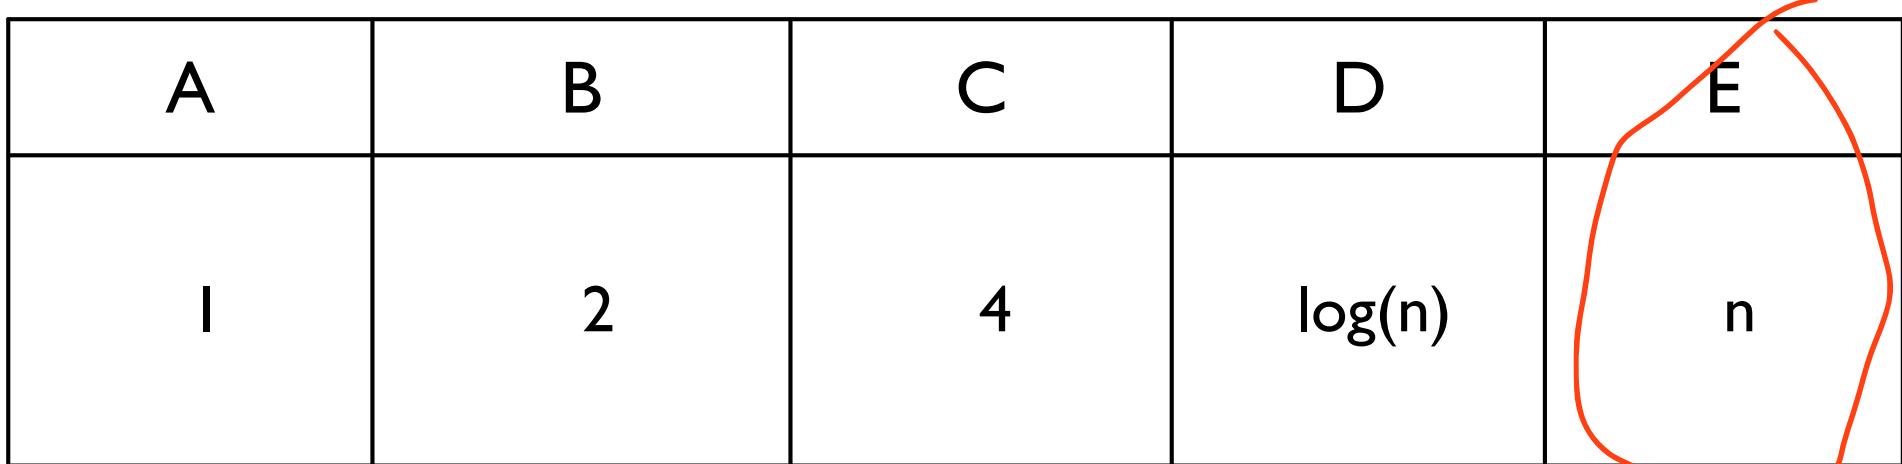
```

def reducer(self, key, counts):
    if key == "words":
        yield ("words", sum(counts))
    elif key == "chars":
        yield ("chars", sum(counts))

```

The maximal speedup achievable by the mapper is ...

Theoretical n
real-world $k = \# \text{nodes}$



$n = \# \text{ lines}$

```
from mrjob.job import MRJob

class WC(MRJob):

    def mapper(self, _, line):
        nr_words = len(line.split())
        nr_chars = len(line)
        yield ("words", nr_words)
        yield ("chars", nr_chars)

    def reducer(self, key, counts):
        if key == "words":
            yield ("words", sum(counts))
        elif key == "chars":
            yield ("chars", sum(counts))

if __name__ == '__main__':
    WC.run()
```

Creates $2n$ tuples

Parallel: up to one node per line

```

def mapper(self, _, line):
    nr_words = len(line.split())
    nr_chars = len(line)
    yield ("words", nr_words)
    yield ("chars", nr_chars)

def reducer(self, key, counts):
    if key == "words":
        yield ("words", sum(counts))
    elif key == "chars":
        yield ("chars", sum(counts))

```

The maximal speedup achievable by the reducer is ...

A	B	C	D	E
I	2	4	$\log(n)$	n

$n = \# \text{ lines}$

```
from mrjob.job import MRJob

class WC(MRJob):

    def mapper(self, _, line):
        nr_words = len(line.split())
        nr_chars = len(line)
        yield ("words", nr_words)
        yield ("chars", nr_chars)

    def reducer(self, key, counts):
        if key == "words":
            yield ("words", sum(counts))
        elif key == "chars":
            yield ("chars", sum(counts))

if __name__ == '__main__':
    WC.run()
```

Is called twice.
 $O(n)$ input

Parallel: up to
two nodes

$O(n)$ communication

```
class WC(MRJob):

    def mapper(self, _, line):
        nr_words = len(line.split())
        nr_chars = len(line)
        yield ("words", nr_words)
        yield ("chars", nr_chars)

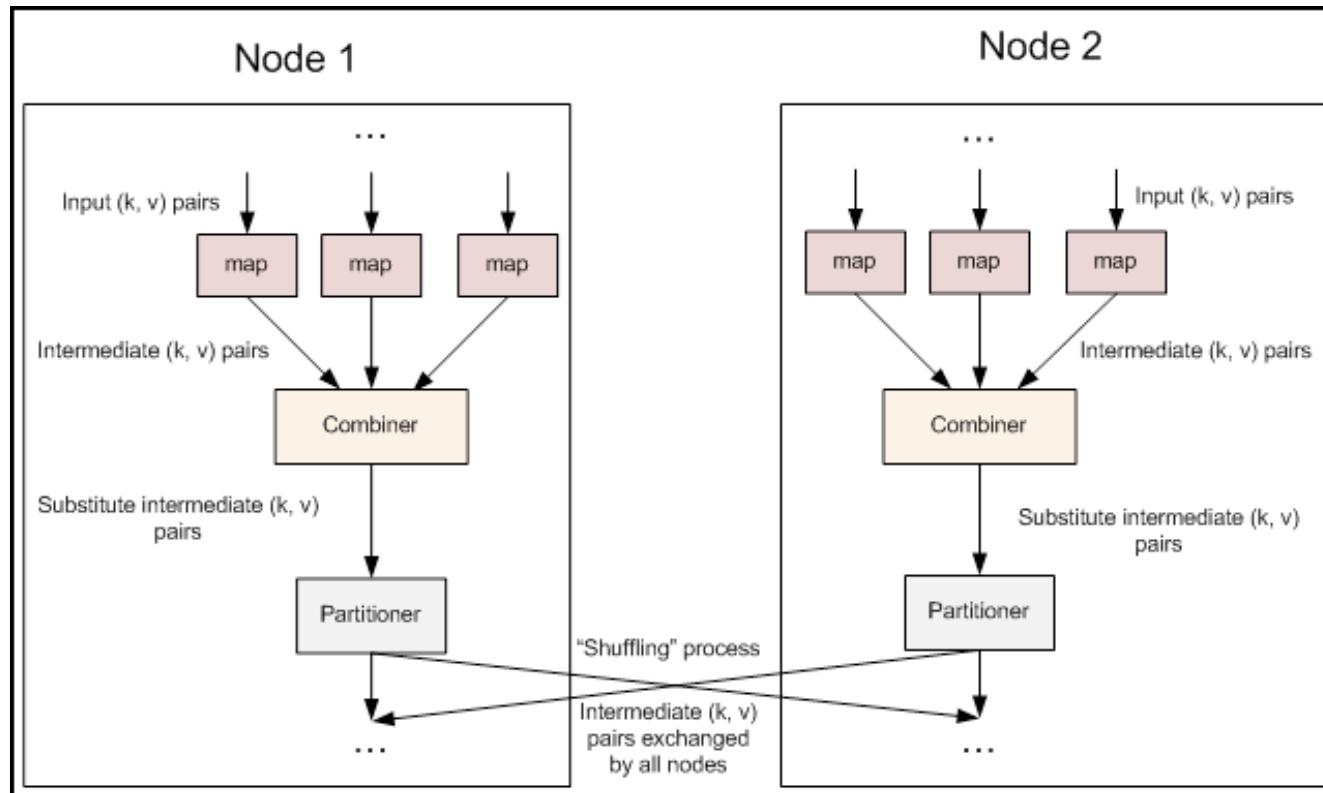
    def combiner(self, key, counts):
        if key == "words":
            yield ("words", sum(counts))
        elif key == "chars":
            yield ("chars", sum(counts))

    def reducer(self, key, counts):
        if key == "words":
            yield ("words", sum(counts))
        elif key == "chars":
            yield ("chars", sum(counts))

if __name__ == '__main__':
    WC.run()
```

n = # lines
k= #nodes for
mapper

Combiners



```
class WC(MRJob):  
  
    def mapper(self, _, line):  
        nr_words = len(line.split())  
        nr_chars = len(line)  
        yield ("words", nr_words)  
        yield ("chars", nr_chars)  
  
    def combiner(self, key, counts):  
        if key == "words":  
            yield ("words", sum(counts))  
        elif key == "chars":  
            yield ("chars", sum(counts))  
  
    def reducer(self, key, counts):  
        if key == "words":  
            yield ("words", sum(counts))  
        elif key == "chars":  
            yield ("chars", sum(counts))  
  
if __name__ == '__main__':  
    WC.run()  
  
n = # lines  
k= #nodes for  
mapper  
  
Creates 2k tuples  
  
As parallel as  
mapper  
  
Same function as  
reducer on a single  
node
```

```

class WC(MRJob):

    def mapper(self, _, line):
        nr_words = len(line.split())
        nr_chars = len(line)
        yield ("words", nr_words)
        yield ("chars", nr_chars)

    def combiner(self, key, counts):
        if key == "words":
            yield ("words", sum(counts))
        elif key == "chars":
            yield ("chars", sum(counts))

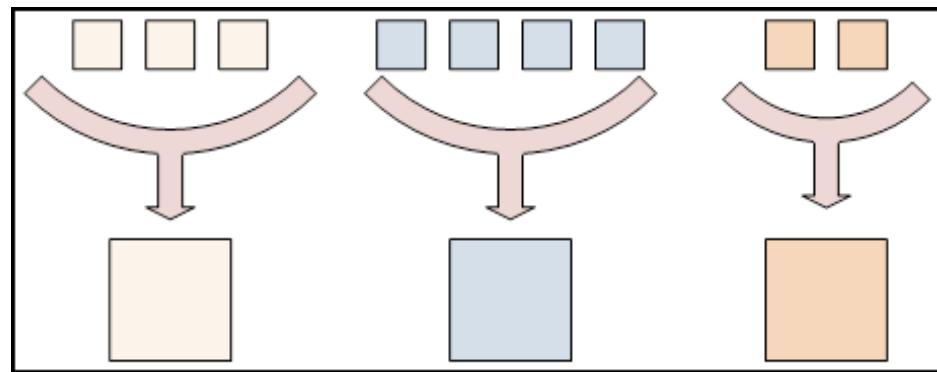
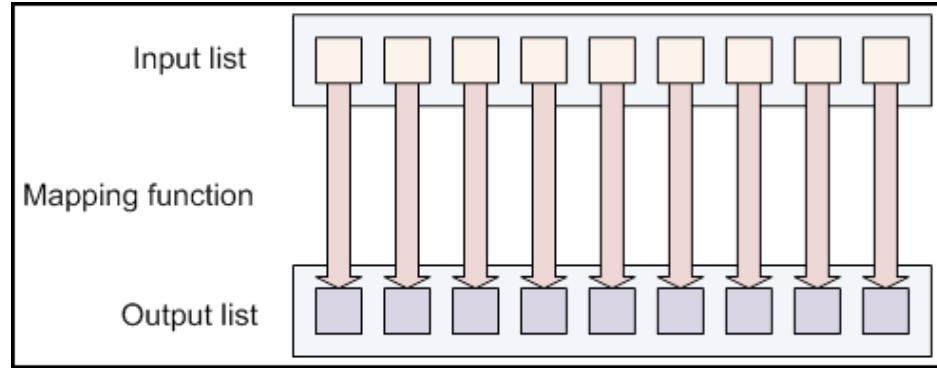
    def reducer(self, key, counts):
        if key == "words":
            yield ("words", sum(counts))
        elif key == "chars":
            yield ("chars", sum(counts))

if __name__ == '__main__':
    WC.run()

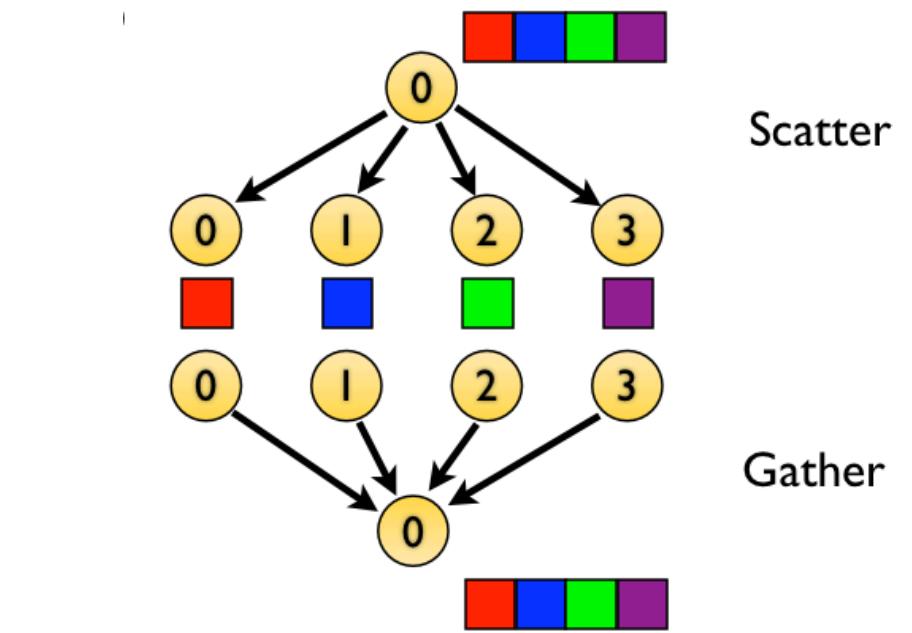
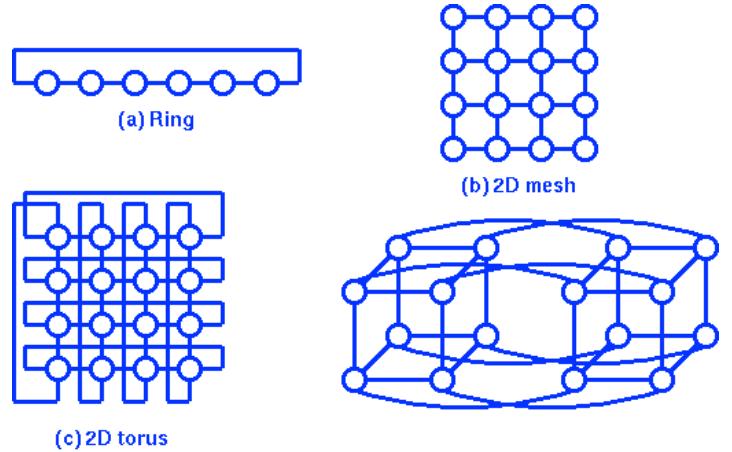
```

$n = \# \text{ lines}$
 $k = \# \text{nodes for mapper}$

Called twice with $O(k)$ input
Parallel: two nodes
 $O(k)$ communication



keys define
communication structure
and effort implicitly



Explicit parallelization in
MPI and threads

Questions & Answers from Chat

Q: Do mapper and combiner always use the same number of nodes? A: A combiner is run on all (key, value) pairs produced by a mapper on one node. So conceptually, yes.

Q: What is the speedup of mapper? A: The theoretical maximal speedup is $n = \# \text{lines}$ if executed on n nodes. Real-world: the speedup is roughly equal to the number of nodes the mapper is running on

Q: What does reduce do? A: Mapper produces ("words", 2), ("words", 4), ("words", 5) [on some machines, maybe in parallel] then reduce gets called with "words", [2,4,5] and thus can count the total number of words.

Q: Does that mean that the reducer has to sum n values together? A: If there is no combiner, yes. This is clearly bad, so that is why the combiner has been introduced. So on each of the nodes the combiner has to add up n/k values (this is parallel on k nodes; on each node multiple cores could be used). The reducer only get k values to add up.

Agenda for today:

- Using Multi-threaded programming
- Message passing:
 - Parallel primitives
 - Reduce
- Map/Reduce

Upcoming Deadlines

 **Python Programming 2**
Not available until Apr 1 | Due Apr 15 at 10am | -/5 pts

 **Preparation for Lecture 4/15**
Not available until Apr 1 | Due Apr 15 at 10am

 **Assignment 2**
Available until Apr 20 | Due Apr 17 at 10am | -/14 pts