

## Agenda for today:

- Map/Reduce
- Spark

## Upcoming Deadlines

 **Assignment 2**  
Available until Apr 20 | Due Apr 17 at 10am | -/14 pts

 **Preparation for Lecture 4/20**  
Not available until Apr 15 | Due Apr 20 at 10am

 **Preparation for Lecture 4/22**  
Not available until Apr 20 | Due Apr 22 at 10am

 **Python Programming 3**  
Not available until Apr 15 | Due Apr 29 at 10am | -/5 pts

 **Assignment 3**  
Available until May 7 | Due May 4 at 10am | -/17 pts


# Announcements

- Bayes:
  - Only run very short test jobs interactively (e.g. < 8 threads, a few seconds)
  - Make sure your that jobs in the queue are quite granular (e.g., not one giant computation)
  - Monitor your job ... if it takes too long, it might be better to kill it

# Recap

# MPI scatter/gather vs. mp.map



Scatter



Gather


`s = p.map(f, [n] * w)`

`[n1, n2, n3, ..., nw]`




`[s1, s2, s3, ..., sw]`

# MPI Reduce



Reduce supports: Sum, Prod, Max, Min, Argmax, Argmin, Logic operations

# Reduce is *not* Gather + Function call




# Map/reduce assumptions

- Input files are distributed over nodes (implicit)
- All data are (key, value) tuples
- Main parallel operations are:
  - Map:  
 $(\text{key}, \text{value}) \rightarrow (\text{key}_1, \text{value}_1), (\text{key}_2, \text{value}_2), \dots$
  - Reduce:  
 $(\text{key}_1, [\text{value}_1, \text{value}_2, \text{value}_3, \dots]) \rightarrow (\text{key}_{\text{new}}, \text{value}_{\text{new}}), (\text{key}_{\text{new}2}, \text{value}_{\text{new}2}), \dots$


[ ] optional

# Map




Each map call only needs its input; no assertions about other inputs, execution order, assignment to nodes

# Reduce



*keys control amount of parallelism*



Each reduce call receives all values for its key; no assertions about other keys, execution order, assignment to nodes

```
class WC(MRJob):
```

$n = \# \text{ lines}$   
 $k = \# \text{ nodes for mapper}$

```
def mapper(self, _, line):
    nr_words = len(line.split())
    nr_chars = len(line)
    yield ("words", nr_words)
    yield ("chars", nr_chars)
```

Creates  $2n$  tuples

Parallel: up to one node per line

```
def combiner(self, key, counts):
    if key == "words":
        yield ("words", sum(counts))
    elif key == "chars":
        yield ("chars", sum(counts))
```

Creates  $2k$  tuples

As parallel as mapper

---

```
def reducer(self, key, counts):
    if key == "words":
        yield ("words", sum(counts))
    elif key == "chars":
        yield ("chars", sum(counts))
```

$O(k)$  communication in Shuffle

Called twice with  $O(k)$  input

Parallel: two nodes

```
if __name__ == '__main__':
    WC.run()
```

Continuing ...

Implementing wordcount with Map-reduce

**Q: how to get line count?**

```
from mrjob.job import MRJob

class WC(MRJob):

    def mapper(self, _, line):
        nr_words = len(line.split())
        nr_chars = len(line)
        yield ("words", nr_words)
        yield ("chars", nr_chars)

    def reducer(self, key, counts):
        if key == "words":
            yield ("words", sum(counts))
        elif key == "chars":
            yield ("chars", sum(counts))

if __name__ == '__main__':
    WC.run()
```

```
from mrjob.job import MRJob

class WC(MRJob):

    def mapper(self, _, line):
        nr_words = len(line.split())
        nr_chars = len(line)
        yield ("words", nr_words)
        yield ("chars", nr_chars)

    def reducer(self, key, counts):
        if key == "words":
            total = 0
            for i, c in enumerate(counts):
                total += c
            yield ("words", total)
            yield ("lines", i+1)
        elif key == "chars":
            yield ("chars", sum(counts))

if __name__ == '__main__':
    WC.run()
```

One  
(words, nr\_words)  
tuple per line

**line count with combiner?**

```
def mapper(self, _, line):
    nr_words = len(line.split())
    nr_chars = len(line)
    yield ("words", nr_words)
    yield ("chars", nr_chars)

def combiner(self, key, counts):
    if key == "words":
        yield ("words", sum(counts))
    elif key == "chars":
        yield ("chars", sum(counts))

def reducer(self, key, counts):
    if key == "words":
        total = 0
        for i, c in enumerate(counts):
            total += c
        yield ("words", total)
        yield ("lines", i+1)
    elif key == "chars":
        yield ("chars", sum(counts))

if __name__ == '__main__':
    WC().run()
```

One  
(words, nr\_words)  
tuple per line on node

One  
(words, nr\_words)  
tuple per node

```
def combiner(self, key, counts):
    if key == "words":
        total = 0
        for i, c in enumerate(counts):
            total += c
        yield ("words", (i+1, total))
    elif key == "chars":
        yield ("chars", sum(counts))
```

```
def reducer(self, key, counts):
    if key == "words":
        total_lines = 0
        total_words = 0
        for c in counts:
            total_lines += c[0]
            total_words += c[1]
        yield ("lines", total_lines)
        yield ("words", total_words)
    elif key == "chars":
        yield ("chars", sum(counts))
```

```
__name__ == '__main__':
```

Creates k tuples  
w/ aggregated  
word, line counts

```
def combiner(self, key, counts):
    if key == "words":
        total = 0
        for i, c in enumerate(counts):
            total += c
        yield ("words", (i+1, total))
    elif key == "chars":
        yield ("chars", sum(counts))
```

```
def reducer(self, key, counts):
    if key == "words":
        total_lines = 0
        total_words = 0
        for c in counts:
            total_lines += c[0]
            total_words += c[1]
        yield ("lines", total_lines)
        yield ("words", total_words)
    elif key == "chars":
        yield ("chars", sum(counts))

if __name__ == '__main__':
```

Sums over the k tuples

Implementing wordcount with Map-reduce

**Q: how to get average #words and  
#character per line??**

# The following code is correct

```
def reducer(self, key, counts):
    if key == "words":
        total_lines = 0
        total_words = 0
        for c in counts:
            total_lines += c[0]
            total_words += c[1]
        yield ("lines", total_lines)
        yield ("words", total_words)
        yield ("wordsperline",
              float(total_words)/total_lines)
    elif key == "chars":
        yield ("chars", sum(counts))
        charsperline = float(sum(counts)) / total_lines
```

A	B	C	D	E
Yes	No			

```
def reducer(self, key, counts):
    if key == "words":
        total_lines = 0
        total_words = 0
        for c in counts:
            total_lines += c[0]
            total_words += c[1]
            yield ("lines", total_lines)
            yield ("words", total_words)
            yield ("wordspерline",
                  float(total_words)/total_lines)
    elif key == "chars":
        yield ("chars", sum(counts))
```

Not available at same time!

Possibly not available on the same machine

**average #characters per word?**

```
class WC(MRJob):

    def mapper(self, _, line):
        nr_words = len(line.split())
        nr_chars = len(line)
        yield ("linestats", (nr_words, nr_chars))

    def combiner(self, key, counts):
        nr_words = 0
        nr_chars = 0
        for i, c in enumerate(counts):
            nr_words += c[0]
            nr_chars += c[1]
        nr_lines = i + 1
        yield ("stats", (nr_lines, nr_words, nr_chars))

    def reducer(self, key, counts):
        total_lines = 0
        total_words = 0
        total_chars = 0
        for c in counts:
            total_lines += c[0]
```

```

def mapper(self, _, line):
    nr_words = len(line.split())
    nr_chars = len(line)
    yield ("linestats", (nr_words, nr_chars))

def combiner(self, key, counts):
    nr_words = 0
    nr_chars = 0
    for i, c in enumerate(counts):
        nr_words += c[0]
        nr_chars += c[1]
    nr_lines = i + 1
    yield ("stats", (nr_lines, nr_words, nr_chars))

```

The maximal speedup achievable by the reducer (e.g. on how many machines can it be run) is ...

A	B	C	D	E
1	2	4	$\log(n)$	n

```
yield ("stats", (nr_lines, nr_words, nr_chars))

def reducer(self, key, counts):
    total_lines = 0
    total_words = 0
    total_chars = 0
    for c in counts:
        total_lines += c[0]
        total_words += c[1]
        total_chars += c[2]

    yield ("lines", total_lines)
    yield ("words", total_words)
    yield ("chars", total_chars)

    yield ("wordsperline", float(total_words)/total_lines)
    yield ("charsperline", float(total_chars)/total_lines)
    yield ("charsperword", float(total_chars)/total_words)
```

Working with Map-reduce

**Q: how to pass parameters?**

```
""" Find duplicate keys in a file containing lines consisting of
key,value
"""
from mrjob.job import MRJob

class FindDuplicates(MRJob):

    def mapper(self, _, line):
        key, value = line.split(',')
        yield (key, 1)

    def combiner(self, key, counts):
        yield (key, sum(counts))

    def reducer(self, key, counts):
        s = sum(counts)
        if s > 1: _____
            yield (key, s)

if __name__ == '__main__':
    FindDuplicates.run()
```

Set threshold from  
command line

```
from mrjob.job import MRJob

class FindDuplicates(MRJob):

    def mapper(self, _, line):
        key, value = line.split(',')
        yield (key, 1)

    def combiner(self, key, counts):
        yield (key, sum(counts))
```

```
def reducer_init(self):
    self.threshold = 42
```

Executed once (per node)  
before running reducer

```
def reducer(self, key, counts):
    s = sum(counts)
    if s > self.threshold:
        yield (key, s)
```

Used in every call

```
if __name__ == '__main__':
    FindDuplicates.run()
```

## ss FindDuplicates(MRJob):

```
def mapper(self, _, line):
    key, value = line.split(',')
    yield (key, 1)

def combiner(self, key, counts):
    yield (key, sum(counts))
```

Define command line arguments for program

```
def configure_args(self):
    super(FindDuplicates, self).configure_args()
    self.add_passthru_arg('--threshold', default=1,
                          help="Frequency threshold")
```

```
def reducer_init(self):
    self.threshold = int(self.options.threshold)
```

```
def reducer(self, key, counts):
    s = sum(counts)
    if s > self.threshold:
        yield (key, s)
```

```
name == '__main__':
```

## ss FindDuplicates(MRJob):

```
def mapper(self, _, line):
    key, value = line.split(',')
    yield (key, 1)

def combiner(self, key, counts):
    yield (key, sum(counts))

def configure_args(self):
    super(FindDuplicates, self).configure_args()
    self.add_passthru_arg('--threshold', default=1,
                          help="Frequency threshold")

def reducer_init(self):
    self.threshold = int(self.options.threshold)

def reducer(self, key, counts):
    s = sum(counts)
    if s > self.threshold:
        yield (key, s)

name == 'main':
```

Configure arguments for  
super class

## ss FindDuplicates(MRJob):

```
def mapper(self, _, line):
    key, value = line.split(',')
    yield (key, 1)

def combiner(self, key, counts):
    yield (key, sum(counts))

def configure_args(self):
    super(FindDuplicates, self).configure_args()
    self.add_passthru_arg('--threshold', default=1,
                          help="Frequency threshold")

def reducer_init(self):
    self.threshold = int(self.options.threshold)

def reducer(self, key, counts):
    s = sum(counts)
    if s > self.threshold:
        yield (key, s)

if __name__ == '__main__':
    pass
```

Add a new argument. Make sure does not clash with MrJob

## ss FindDuplicates(MRJob):

```
def mapper(self, _, line):
    key, value = line.split(',')
    yield (key, 1)

def combiner(self, key, counts):
    yield (key, sum(counts))

def configure_args(self):
    super(FindDuplicates, self).configure_args()
    self.add_passthru_arg('--threshold', default=1,
                          help="Frequency threshold")

def reducer_init(self):
    self.threshold = int(self.options.threshold)

def reducer(self, key, counts):
    s = sum(counts)
    if s > self.threshold:
        yield (key, s)
```

`self.options` is a namespace holding all command line arguments

`threshold = 1 unless specified on the command line`

`name == 'main'.`

```
MacBookPro13:Sandbox schliep$ python mrjob-duplicates-combiner-2.py --threshold 121 test.data
No configs found; falling back on auto-configuration
No configs specified for inline runner
Running step 1 of 1...
Creating temp directory /var/folders/wx/9cjwzycx77j7dxc0nh5zbn8m0000gn/T/mrjob-duplicates-
combiner-2.schliep.20180425.213944.509621
Streaming final output from /var/folders/wx/9cjwzycx77j7dxc0nh5zbn8m0000gn/T/mrjob-
duplicates-combiner-2.schliep.20180425.213944.509621/output...
"0000000129" 126
"0000000189" 126
"0000000454" 131
"0000000480" 123
"0000000551" 133
"0000000591" 126
"0000000639" 140
"0000000669" 124
"0000000847" 130
"0000000934" 128
"0000000945" 133
```

Working with Map-reduce

**Q: Multiple steps or how to find the  
most frequent item?**

**most frequent item?**

```
from mrjob.job import MRJob

class FindDuplicates(MRJob):

    def mapper(self, _, line):
        key, value = line.split(',')
        yield (key, 1)

    def combiner(self, key, counts):
        yield (key, sum(counts))
```

```
def reducer(self, key, counts):
    s = sum(counts)
    if s > self.threshold:
        yield (key, s)
```

After reducer has completed all words and their frequencies are available

```
if __name__ == '__main__':
    FindDuplicates.run()
```

Need one more step!

```
from mrjob.job import MRJob, MRStep

class FindDuplicates(MRJob):

    def mapper(self, _, line):
    ...
    def combiner(self, key, counts):
        ...
    def reducer(self, key, counts):      Produces (freq,word)
        s = sum(counts)                  pairs
        yield None, (s, key)

    def findmax(self, _, count_word_tuples):
        yield max(count_word_tuples)

    def steps(self):
        return [MRStep(mapper=self.mapper,
                      combiner=self.combiner,
                      reducer=self.reducer),
                MRStep(reducer=self.findmax)]

if __name__ == '__main__':
    FindDuplicates.run()
```

```
from mrjob.job import MRJob, MRStep

class FindDuplicates(MRJob):

    def mapper(self, _, line):
        ...
        ...

    def combiner(self, key, counts):
        ...

    def reducer(self, key, counts):
        s = sum(counts)
        yield None, (s, key)

    def findmax(self, _, count_word_tuples):
        yield max(count_word_tuples)          max looks at first
                                                item in tuple

    def steps(self):
        return [MRStep(mapper=self.mapper,
                      combiner=self.combiner,
                      reducer=self.reducer),
                MRStep(reducer=self.findmax)]

if __name__ == '__main__':
    FindDuplicates.run()
```

```
from mrjob.job import MRJob, MRStep
```

```
class FindDuplicates(MRJob):
```

```
    def mapper(self, _, line):
```

```
    ...
```

```
    def combiner(self, key, counts):
```

```
        ...
```

```
    def reducer(self, key, counts):
```

```
        s = sum(counts)
```

```
        yield None, (s, key)
```

```
    def findmax(self, _, count_word_tuples):
```

```
        yield max(count_word_tuples)
```

```
    def steps(self):
```

```
        return [MRStep(mapper=self.mapper,
                        combiner=self.combiner,
                        reducer=self.reducer),
                MRStep(reducer=self.findmax)]
```

Manually define  
multi-step  
execution

```
if __name__ == '__main__':
```

```
from mrjob.job import MRJob, MRStep
```

```
class FindDuplicates(MRJob):
```

```
    def mapper(self, _, line):
```

.. Default MRJob steps method

```
    def steps(self):  
        return [MRStep(mapper=self.mapper,  
                       combiner=self.combiner,  
                       reducer=self.reducer)]
```

define  
P  
n

Can specify init functions. Steps can have one or all of the commands

```
        combiner=self.combiner,  
        reducer=self.reducer),  
        MRStep(reducer=self.findmax)]
```

```
if __name__ == '__main__':
```

# Parallel programming models

## Spark

# Apache Spark

- More general model of computation
- Automated parallelization
- In memory (with disk caching) computation
- Explicit program flow

Working with Spark

Computing pi with Monte Carlo

# Monte Carlo: Computing Pi

- Random variable  $X \sim \pi$


- When  $x_i$  drawn from  $\pi$

$$\frac{1}{n} \sum f(x_i) \rightarrow E_{\pi}[f(X)]$$

- Example:

$$\frac{1}{n} \sum f(x_i) \rightarrow \frac{\pi}{4} \quad x_i \sim U[0,1]^2$$

$$\frac{1}{n} \sum f(x_i) \rightarrow E_{\pi}[f(X)] = \int_0^1 \sqrt{1 - x^2} dx$$



# Reminder: Computing Pi with multiprocessing

```
def sample_pi(n):
    """ Perform n steps of Monte Carlo simulation for estimating Pi/4.
    Returns the number of successes."""
    s = 0
    for i in range(n):
        x,y = random.random(), random.random()
        if x**2 + y**2 <= 1.0:
            s += 1
    return s

def compute_pi(args):
    n = args.steps / args.workers

    p = multiprocessing.Pool(args.workers)
    s = p.map(sample_pi, [n]*args.workers)

    n_total = n*args.workers
    s_total = sum(s)
    pi_est = (4.0*s_total)/n_total
    print " Steps\tSuccess\tPi est.\tError"
    print "%6d\t%7d\t%1.5f\t%1.5f" % (n_total, s_total, pi_est, pi-pi_est)
```

# Computing Pi with multiprocessing

```
def sample_pi(n):
    """ Perform n steps of Monte Carlo simulation for estimating Pi/4.
    Returns the number of successes."""
    s = 0
    for i in range(n):
        x,y = random.random(), random.random()
        if x**2 + y**2 <= 1.0:
            s += 1
    return s
```

Sample x,y uniform and  
check whether it is within  
the quarter circle

```
def compute_pi(args):
    n = args.steps / args.workers

    p = multiprocessing.Pool(args.workers)
    s = p.map(sample_pi, [n]*args.workers)

    n_total = n*args.workers
    s_total = sum(s)
    pi_est = (4.0*s_total)/n_total
    print " Steps\tSuccess\tPi est.\tError"
    print "%6d\t%7d\t%1.5f\t%1.5f" % (n_total, s_total, pi_est, pi-est)
```

# Computing Pi with multiprocessing

```
def sample_pi(n):
    """ Perform n steps of Monte Carlo simulation for estimating Pi/4.
    Returns the number of successes."""
    s = 0
    for i in range(n):
        x,y = random.random(), random.random()
        if x**2 + y**2 <= 1.0:
            s += 1
    return s

def compute_pi(args):
    n = args.steps / args.workers

    p = multiprocessing.Pool(args.workers)
    s = p.map(sample_pi, [n]*args.workers)
                [n,...,n]
    n_total = n*args.workers
    s_total = sum(s)
    pi_est = (4.0*s_total)/n_total
    print " Steps\tSuccess\tPi est.\tError"
    print "%6d\t%7d\t%1.5f\t%1.5f" % (n_total, s_total, pi_est, pi-est)
```

Create workers and let  
each worker process  
sample\_pi(n)

# Computing Pi with PySpark

```
import random
from pyspark import SparkContext
sc = SparkContext(master = 'local[4]')

NUM_SAMPLES = 100000000

def inside(p):
    x, y = random.random(), random.random()
    return x*x + y*y < 1

count = sc.parallelize(xrange(0, NUM_SAMPLES)) \
        .filter(inside).count()

print "Pi is roughly %f" % (4.0 * count / NUM_SAMPLES)
```

# Computing Pi with PySpark

```
import random
from pyspark import SparkContext
sc = SparkContext(master = 'local[4]')
```

Boilerplate for  
computations on 4  
local cores

```
NUM_SAMPLES = 100000000
```

```
def inside(p):
    x, y = random.random(), random.random()
    return x*x + y*y < 1
```

```
count = sc.parallelize(xrange(0, NUM_SAMPLES)) \
    .filter(inside).count()
```

```
print "Pi is roughly %f" % (4.0 * count / NUM_SAMPLES)
```

# Computing Pi with PySpark

```
import random
from pyspark import SparkContext
sc = SparkContext(master = 'local[4]')

NUM_SAMPLES = 100000000

def inside(p):
    x, y = random.random(), random.random()
    return x*x + y*y < 1

count = sc.parallelize(xrange(0, NUM_SAMPLES)) \
    .filter(inside).count()

print "Pi is roughly %f" % (4.0 * count / NUM_SAMPLES)
```

Sample x,y uniform  
and check whether it  
is within the quarter  
circle. p is ignored

# Computing Pi with PySpark

```
import random
from pyspark import SparkContext
sc = SparkContext(master = 'local[4]')

NUM_SAMPLES = 100000000

def inside(p):
    x, y = random.random(), random.random()
    return x*x + y*y < 1

count = sc.parallelize(xrange(0, NUM_SAMPLES)) \
        .filter(inside).count()

print "Pi is roughly %f" % (4.0 * count / NUM_SAMPLES)
```

# Computing Pi with PySpark

```
count = sc.parallelize(xrange(0, NUM_SAMPLES)) \  
    .filter(inside).count()
```

# Computing Pi with PySpark

```
count = sc.parallelize(xrange(0, NUM_SAMPLES)) \  
    .filter(inside).count()
```

## Example

```
>>> def odd(x):  
...     return x % 2  
...  
>>> a = [2,3,5,8,9]  
>>> b = filter(odd, a)  
>>> b  
[3, 5, 9]  
>>> len(b)  
3
```

# Computing Pi with PySpark

```
count = sc.parallelize(xrange(0, NUM_SAMPLES)) \  
    .filter(inside).count()
```

## Example

```
>>> def odd(x):  
...     return x % 2  
...  
>>> a = [2,3,5,8,9]  
>>> b = filter(odd, a)  
>>> b  
[3, 5, 9]  
>>> len(b)  
3
```

[0,1,2,...,NUM\_SAMPLES-1]



[4,5,9,...]

count

len([4,5,9,...])

Note: array not created

# Computing Pi with PySpark

```
data = [0,1,2,3,4,5]
distData = sc.parallelize(data)
```

- RDD = Resilient Distributed Dataset
- Fault-tolerant collection of records which can be processed in parallel
  - distData.reduce(lambda a, b: a + b)
- Obtained from Driver program (your Python script), read from HDFS or other data sources

# Find duplicates with MrJob

```
""" Find duplicate keys in a file containing lines consisting of
key,value
"""
from mrjob.job import MRJob

class FindDuplicates(MRJob):

    def mapper(self, _, line):
        key, value = line.split(',')
        yield (key, 1)

    def combiner(self, key, counts):
        yield (key, sum(counts))

    def reducer(self, key, counts):
        s = sum(counts)
        if s > 1:
            yield (key, s)

if __name__ == '__main__':
    FindDuplicates.run()
```

# Find duplicates with MrJob

```
""" Find duplicate keys in a file containing lines consisting of
key,value
"""
from mrjob.job import MRJob

class FindDuplicates(MRJob):

    def mapper(self, _, line):
        key, value = line.split(',')
        yield (key, 1)

    def combiner(self, key, counts):
        yield (key, sum(counts))

    def reducer(self, key, counts):
        s = sum(counts)
        if s > 1:
            yield (key, s)

if __name__ == '__main__':
    FindDuplicates.run()
```

Create one tuple per line of file

# Find duplicates with MrJob

```
""" Find duplicate keys in a file containing lines consisting of
key,value
"""
from mrjob.job import MRJob

class FindDuplicates(MRJob):

    def mapper(self, _, line):
        key, value = line.split(',')
        yield (key, 1)

    def combiner(self, key, counts):
        yield (key, sum(counts))

    def reducer(self, key, counts):
        s = sum(counts)
        if s > 1:
            yield (key, s)

if __name__ == '__main__':
    FindDuplicates.run()
```

Collect all tuples  
with same key per  
node and sum

# Find duplicates with MrJob

```
""" Find duplicate keys in a file containing lines consisting of
key,value
"""
from mrjob.job import MRJob

class FindDuplicates(MRJob):

    def mapper(self, _, line):
        key, value = line.split(',')
        yield (key, 1)

    def combiner(self, key, counts):
        yield (key, sum(counts))

    def reducer(self, key, counts):
        s = sum(counts)
        if s > 1:
            yield (key, s)

if __name__ == '__main__':
    FindDuplicates.run()
```

Collect all tuples with same key across nodes and sum.  
Output duplicates.

# Find duplicates with PySpark

```
import random
from pyspark import SparkContext
sc = SparkContext(master = 'local[4]')
```

```
distFile = sc.textFile("test.data")
```

Distribute Local File

```
counts = distFile.map(lambda l: l.split(',')) \
    .map(lambda t:(int(t[0]),1)) \
    .reduceByKey(lambda a, b: a + b) \
    .filter(lambda t:t[1] > 1)
```

```
cc = counts.collect()
print cc
```

Collect output

# Find duplicates with PySpark

```
import random
from pyspark import SparkContext
sc = SparkContext(master = 'local[4]')

distFile = sc.textFile("test.data")

counts = distFile.map(lambda l: l.split(',')) \
    .map(lambda t:(int(t[0]),1)) \
    .reduceByKey(lambda a, b: a + b) \
    .filter(lambda t:t[1] > 1)
```

```
cc = counts.collect()
print cc
```

Collect output

# Lambda expressions

```
>>> def square(x):
...     return x*x
...
>>> map(square, range(10))
[0, 1, 4, 9, 16, 25, 36, 49, 64, 81]
>>> map(lambda x:x**2, range(10))
[0, 1, 4, 9, 16, 25, 36, 49, 64, 81]
```

- Lambda expression are anonymous (nameless) functions
- Often used in sort ...
- lambda x,y,z,...:<expression>

# Find duplicates with PySpark

```
counts = distFile.map(lambda l: l.split(',')) \      Split key,value string  
    .map(lambda t:(int(t[0]),1)) \  
    .reduceByKey(lambda a, b: a + b) \  
    .filter(lambda t:t[1] > 1)
```

# Find duplicates with PySpark

```
counts = distFile.map(lambda l: l.split(',')) \
    .map(lambda t:(int(t[0]),1)) \           Create (key,l) tuples
    .reduceByKey(lambda a, b: a + b) \
    .filter(lambda t:t[1] > 1)
```

# Find duplicates with PySpark

```
counts = distFile.map(lambda l: l.split(',')) \  
    .map(lambda t:(int(t[0]),1)) \  
    .reduceByKey(lambda a, b: a + b) \  
    .filter(lambda t:t[1] > 1)
```

Collect all tuples  
with same key and  
reduce.


# Find duplicates with PySpark

```
counts = distFile.map(lambda l: l.split(',')) \  
    .map(lambda t:(int(t[0]),1)) \  
    .reduceByKey(lambda a, b: a + b) \  
    .filter(lambda t:t[1] > 1)
```

Filter

# Find duplicates with PySpark

```
counts = distFile.map(lambda l: l.split(',')) \  
    .map(lambda t:(int(t[0]),1)) \  
    .reduceByKey(lambda a, b: a + b) \  
    .filter(lambda t:t[1] > 1)
```



# Find duplicates with PySpark

```
counts = distFile.map(lambda l: l.split(',')) \  
    .map(lambda t:(int(t[0]),1)) \  
    .reduceByKey(lambda a, b: a + b) \  
    .filter(lambda t:t[1] > 1)
```

PySpark == MapReduce ???

Assume you want a bit more than duplicates, e.g.

There are 1001 unique keys of which 461 appear more than 100 times (sum = 50100). The following keys appear over 125 times (sum = 1173)

```
[(00454, 131), (00934, 128), (00639, 140),  
(00189, 126), (00847, 130), (00591, 126),  
(00129, 126), (05051, 133), (00945, 133)]
```

How would you implement this in MapReduce?

Assume you want a bit more than duplicates, e.g.

There are 1001 unique keys of which 461 appear more than 100 times (sum = 50100). The following keys appear over 125 times (sum = 1173)

```
[ (00454, 131), (00934, 128), (00639, 140),  
  (00189, 126), (00847, 130), (00591, 126),  
  (00129, 126), (05051, 133), (00945, 133) ]
```

Can the output above be produced effectively using MapReduce?

A	B	C	D	E
True	False			

```
import random
from pyspark import SparkContext
sc = SparkContext(master = 'local[4]')

distFile = sc.textFile("test.data")

keytuples = distFile.map(lambda l: l.split(',')) \
    .map(lambda t:(int(t[0]),1))

countsPerKey = keytuples.reduceByKey(lambda a, b: a + b)

unique_keys_count = countsPerKey.count()

frequent_keys = countsPerKey.filter(lambda t:t[1] > 100)

frequent_keys_count = frequent_keys.count()
frequent_keys_sum = frequent_keys.values().sum()

very_frequent_keys = frequent_keys.filter(lambda t:t[1] > 125)
very_frequent_keys_sum = very_frequent_keys.values().sum()
vfk = very_frequent_keys.collect()

print(f"There are {unique_keys_count} unique keys in the RDD of which {frequent_keys_count} "
      f"appear more than 100 times (sum = {frequent_keys_sum}). The following keys appear "
      f"over 125 times (sum = {very_frequent_keys_sum})")
print(vfk)
```

# Q&A from Chat

Q: Speedup for reduce? A: Reduce is run once for each distinct key used in tuples yielded by combine. If there are only keys A, B, C, D, then reduce is called once each for A, B, C, D, most likely in parallel on 4 nodes. If there is only one key, say A, then reduce is called exactly once.

Q: multiprocessing starmap vs map. A: Starmap is exactly like map in multiprocessing, only allowing for multiple arguments to the function applied. E.g. for  $f(x) = \sin(x)$  use `map[f, [0.1,0.2]]` and for  $g(x,y) = x^2 + y^2$  use `starmap(g, [(0.0,1.0), (0.5, 1.2), ...])`

Q: Why collect? A: In the last example, counts is a RDD, with data likely existing on different nodes, not a basic Python data type or numpy array, and with `collect()` we transform that into something we can process in Python.

Q: Where will filter (or generally RDD transformations) be executed? That is not specified. We will briefly touch on scheduling computations in Spark. Generally, Spark is more flexible and dynamic than MapReduce, so where computations will be executed is less predictable, but probably more efficient than MapReduce. Note that combiners are a “poor-man’s” version of reduce in some sense. Spark embraces the efficiency of reduce at all levels.

## Agenda for today:

- Map/Reduce
- Spark

## Upcoming Deadlines

 **Assignment 2**  
Available until Apr 20 | Due Apr 17 at 10am | -/14 pts

 **Preparation for Lecture 4/20**  
Not available until Apr 15 | Due Apr 20 at 10am

 **Preparation for Lecture 4/22**  
Not available until Apr 20 | Due Apr 22 at 10am

 **Python Programming 3**  
Not available until Apr 15 | Due Apr 29 at 10am | -/5 pts

 **Assignment 3**  
Available until May 7 | Due May 4 at 10am | -/17 pts