

Agenda for today:

- Bloom filters
- Suffix trees
- Birgit & Shirin: Information about project courses
- Sketching

Upcoming Deadlines

 **Python Programming 3**
Available until May 2 | Due Apr 29 at 10am | -/5 pts

 **Assignment 3**
Available until May 7 | Due May 4 at 10am | -/17 pts

Background material for 2nd part of the course
<https://chalmers.instructure.com/courses/9360/pages/advanced-databases-background>

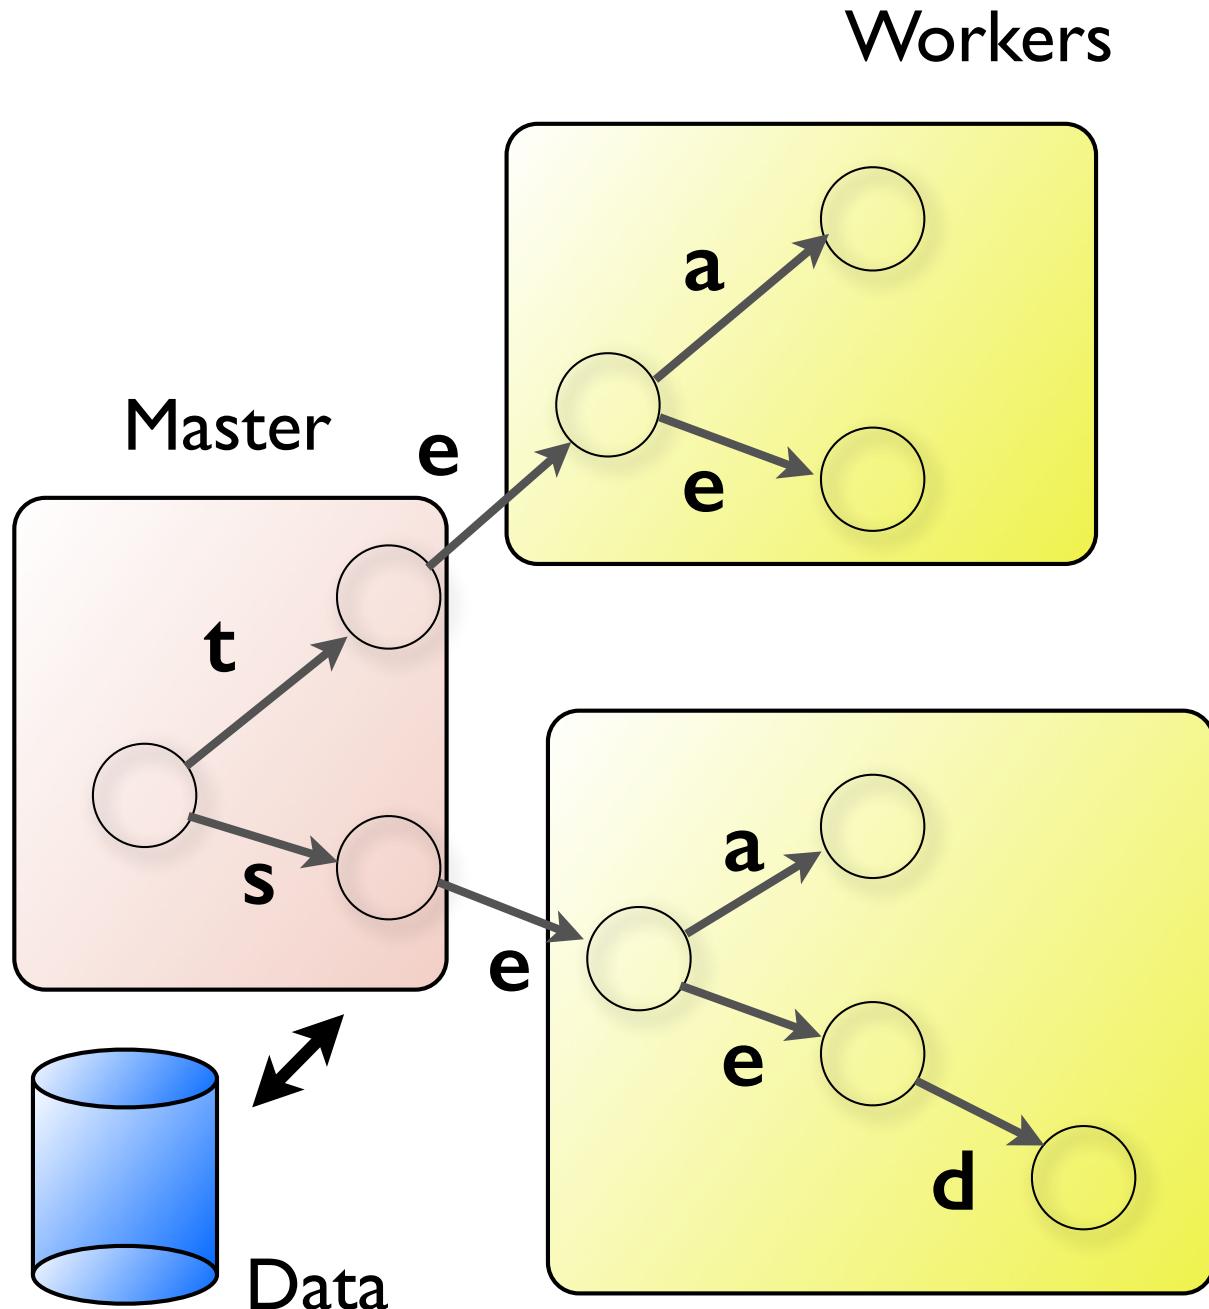
Next lecture: April 29

Recap

Application: Count Frequent Items

- Frequent items:
 - Heavy hitters
 - Top search queries
 - Most frequently requested items in a database
- Very important in Genomics/DNA sequencing

Trie: applications & parallelization



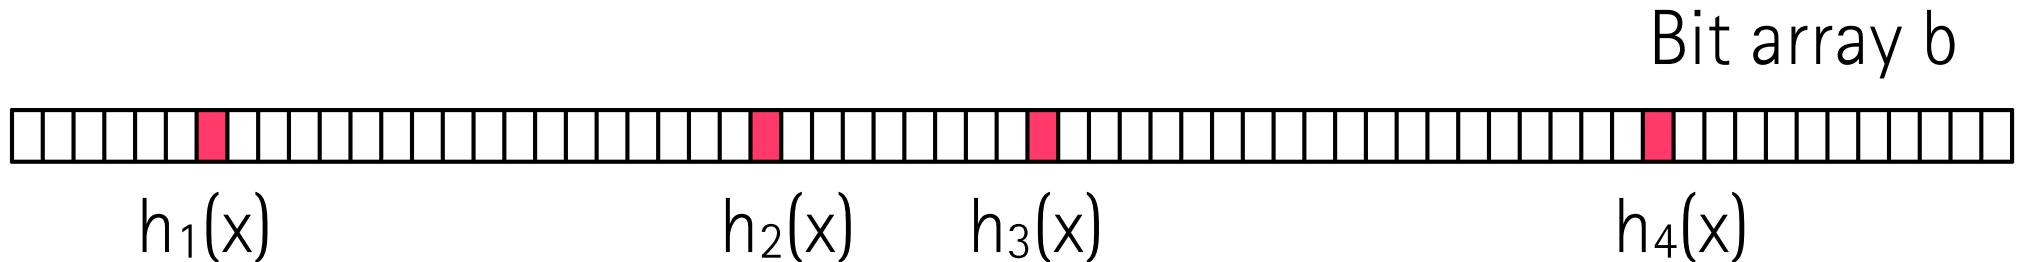
- Duplicates: Second visit to leaf
- Counting: Integer variables in leaf

Trie (Retrieval Tree)

Frequent items: problem sizes

- Billions of items (\sim coverage x genome size)
E.g. 264 Billions for an 80x Human sequencing data set
- Billions of frequent items (\sim genome size)
> 3 Billion for Human
- Billions of infrequent items
100s of Billions

Bloom filter



- Given k hash functions h_1, h_2, \dots, h_k
- Insert: for item x set bits $h_1(x), h_2(x), \dots, h_k(x)$ in a bit array b
- Query: Check whether bits $h_1(x), h_2(x), \dots, h_k(x)$ are set in b
- No false negatives

Bloom filter example

	Text	Python Dict	Trie	Bloom Filter $p = 0.01$	Bloom filter $p = 0.001$	Bloom filter $p = 0.0001$
10 000	146 kB	2.5 MB	1.17 MB	12 KB	18 kB	24 kB
100 000	1.43 MB	19.2 MB	~11.44 MB	117 KB	175 kB	243 kB
1M	14.4 MB	179.9 MB	~112 MB	1.14 MB	1.71 MB	2.29 MB

Assuming 15 unique characters per string. E.g URL

Continuing ...

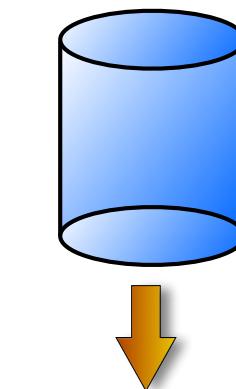
Counting with Bloom Filters?

Counting with Bloom Filters

```
from pybloom import BloomFilter  
f = BloomFilter(capacity=len(X), error_rate=p)  
from collections import defaultdict  
d = defaultdict(int)
```

```
for x in X:  
    if x in f:  
        d[x] += 1  
    else:  
        f.add(x)
```

```
for x in d.keys():  
    if (d[x] + 1) > tau:  
        print x, d[x]
```



Have we seen the k-mer before?

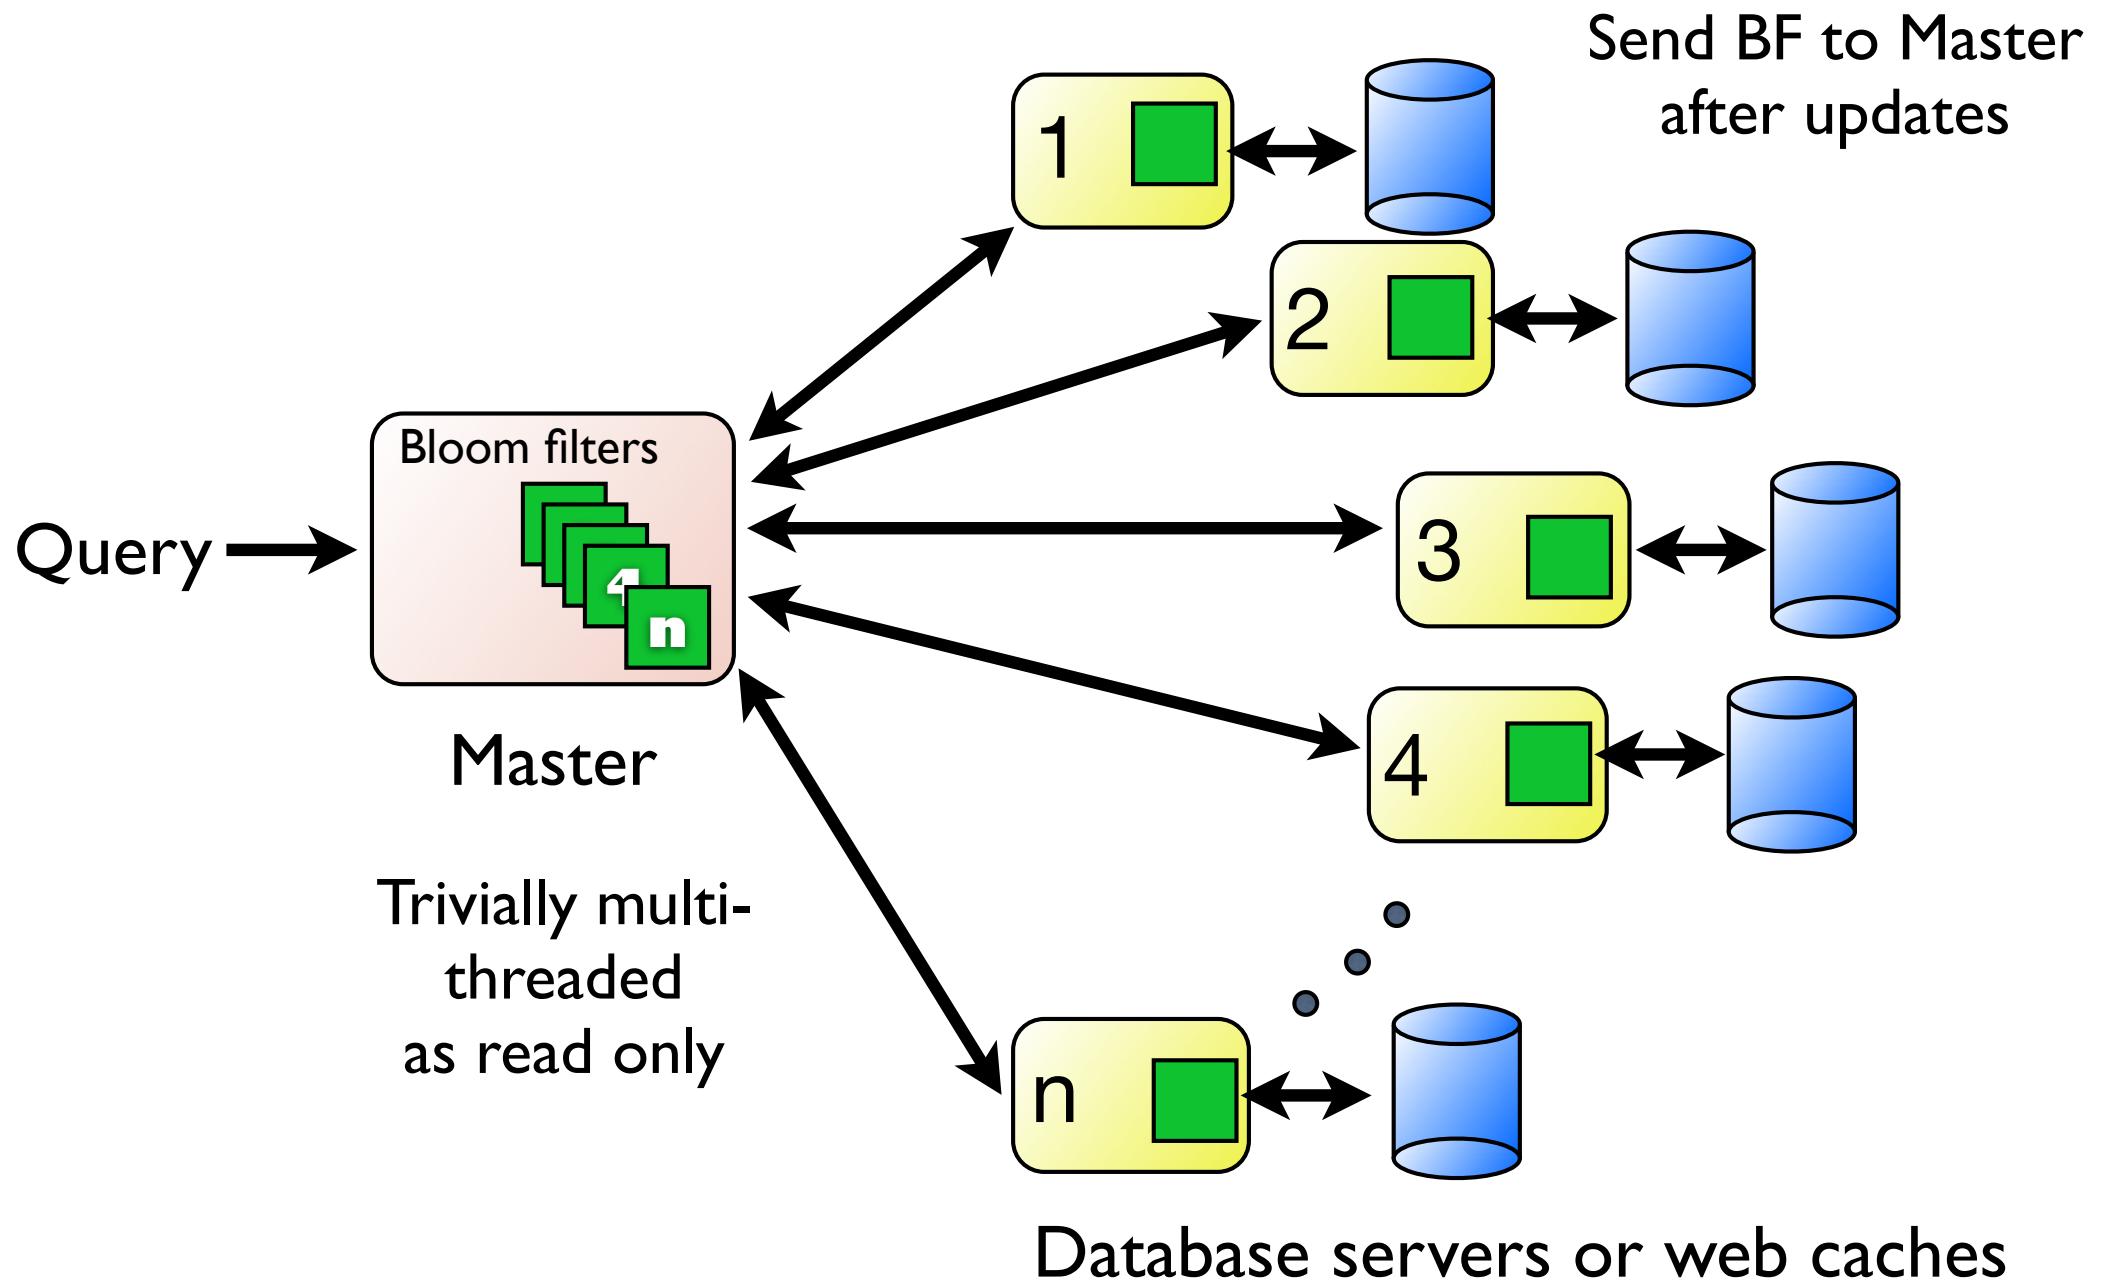


Only count k-mers seen at least twice

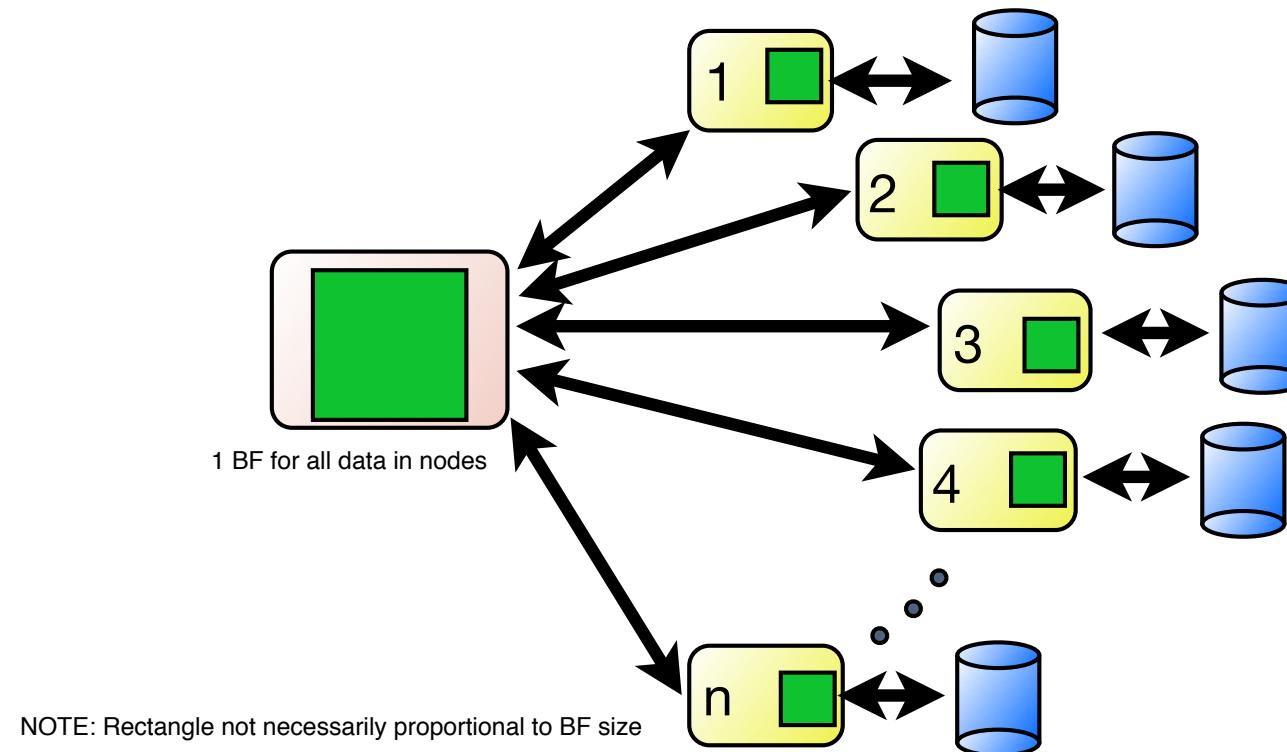
Counting Data Structure

See <https://hadoop.apache.org/docs/r1.0.4/api/org/apache/hadoop/io/BloomMapFile.Reader.html>
for the same idea for files

Distributed Data Store



Can the creation of Bloom filters be efficiently distributed; e.g. with Map Reduce?



A	B	C	D	E
True	False			

Inset X



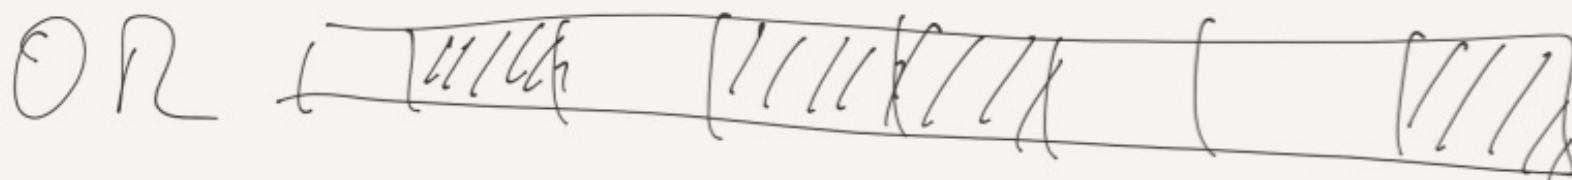
Node 1

Inset Y



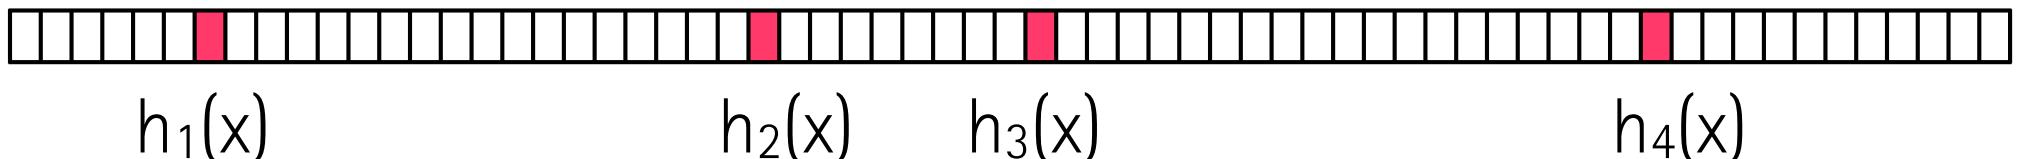
Node 2

Bitwide



Inset x,y

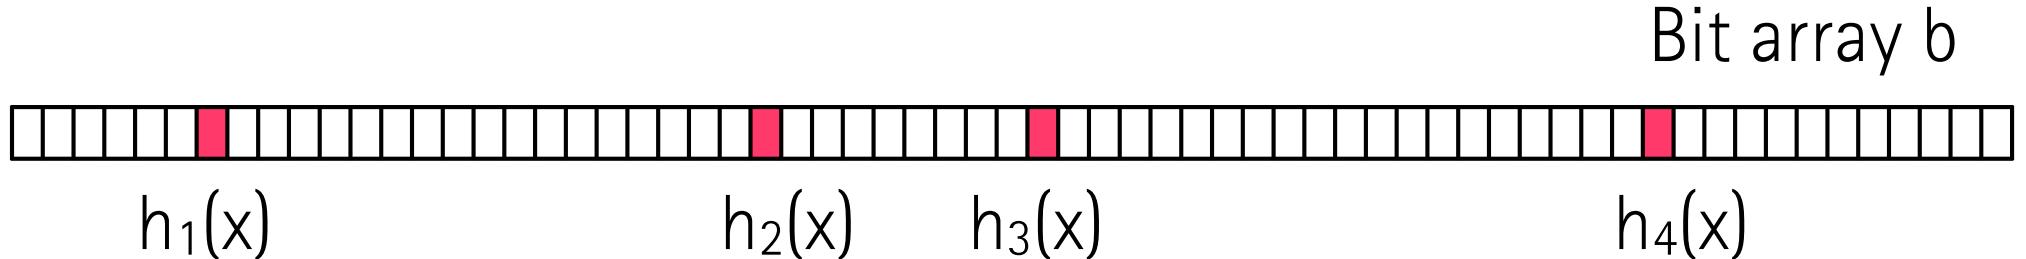
Bit array b



For sufficiently large b , accessing each bit will likely cause a cache miss. The expected number of cache misses for insert and query operations, when the item is not in the Bloom filter is the same.

A	B	C	D	E
True	False			

Bloom filter: Cache Behaviour



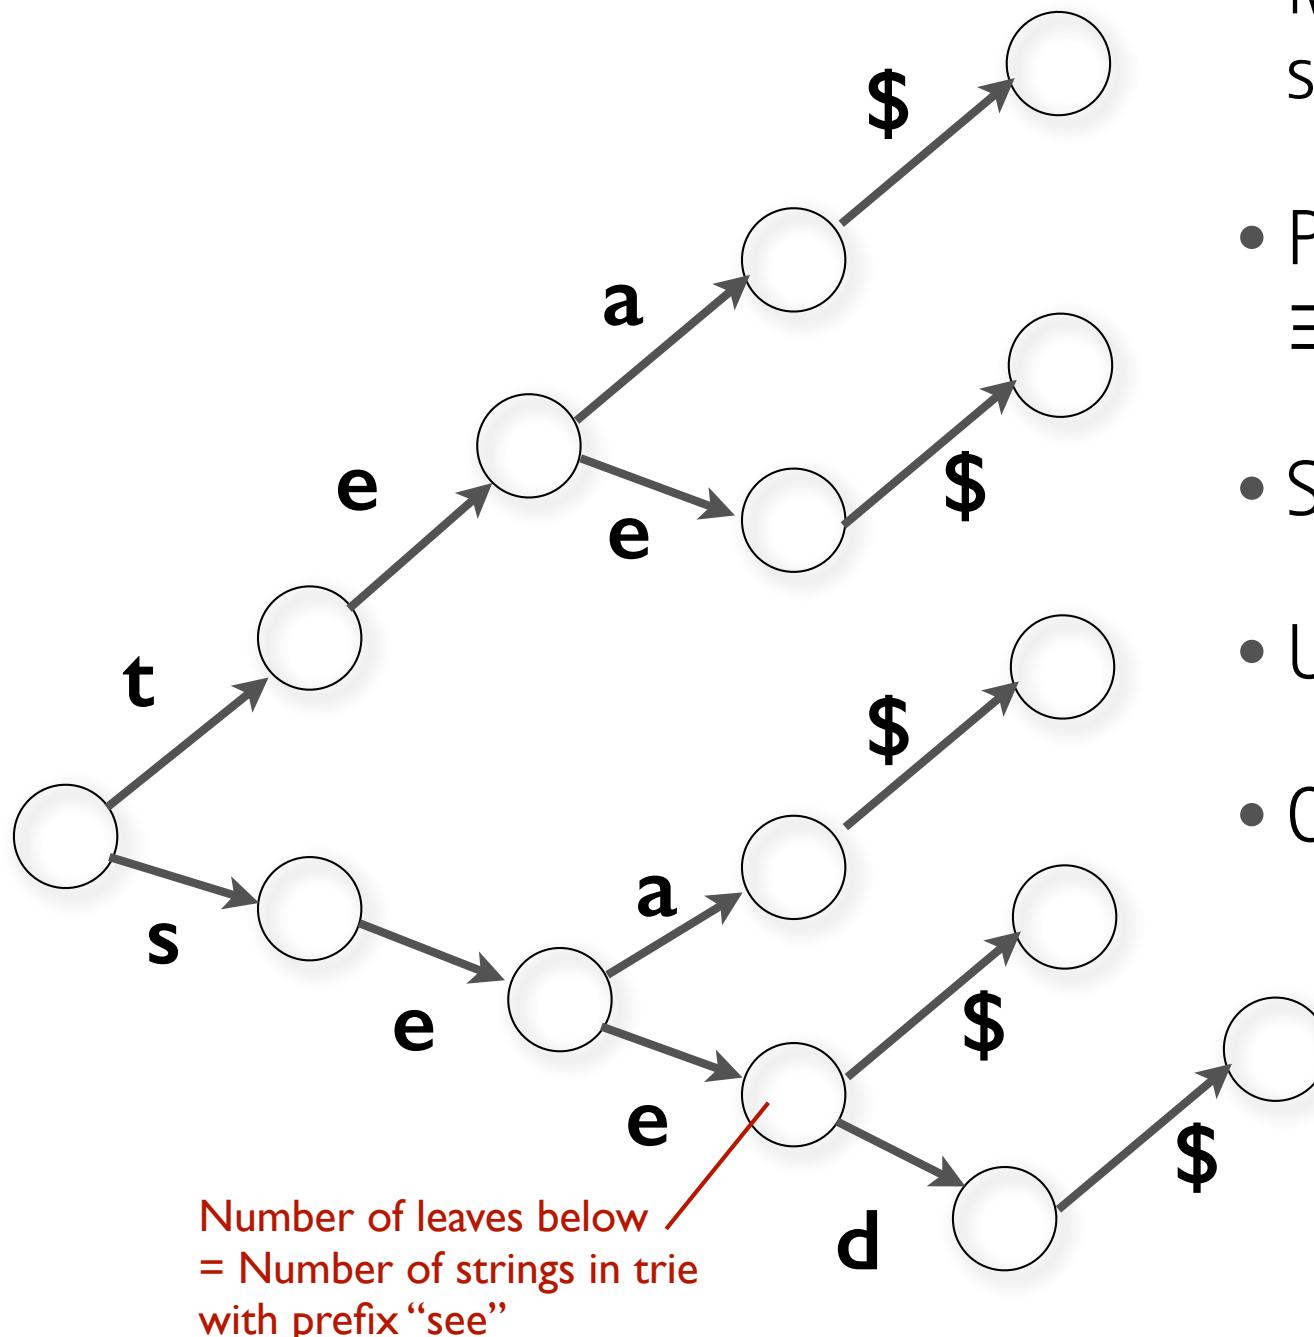
- Insert:
 - Worst-case: n cache misses (average close) per Insert
- Query:
 - item present / false positive: like insert
 - item not-present: average much lower than worst-case

Suffix trees

Data analysis on large texts

- Text T:
 - Natural language
 - Error or event logs (error codes/event types = alphabet)
 - Biological sequences, any other sequence of discrete observations
- Questions:
 - Is s a substring of T?
 - How often appears s in T?

Recall: Tries

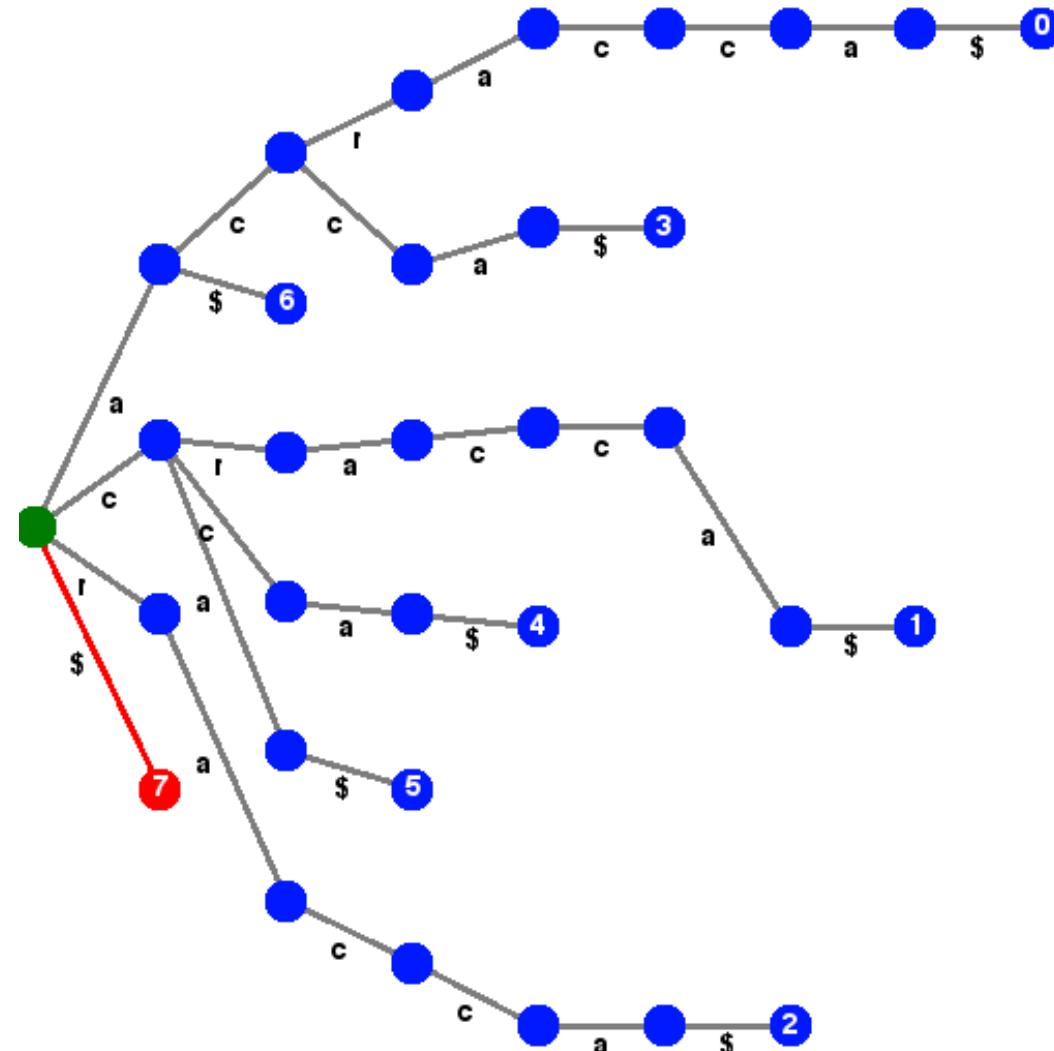


- Membership query:
 $s \in S$
- Prefix query:
 $\exists t \in S: s$ is prefix of t
- Sorting string
- Unique strings
- Count strings

Example: $T = \text{acracca\$}$

$\text{acracca\$}$
 $\text{cracca\$}$
 $\text{racca\$}$
 $\text{acca\$}$
 $\text{cca\$}$
 $\text{ca\$}$
 $\text{a\$}$
 $\$$

Trie of all
suffixes of T



$\overline{T} = \text{b} \overbrace{\text{a n a n a}}^{\text{a}} \text{a}$

Is "na" a substring of \overline{T}

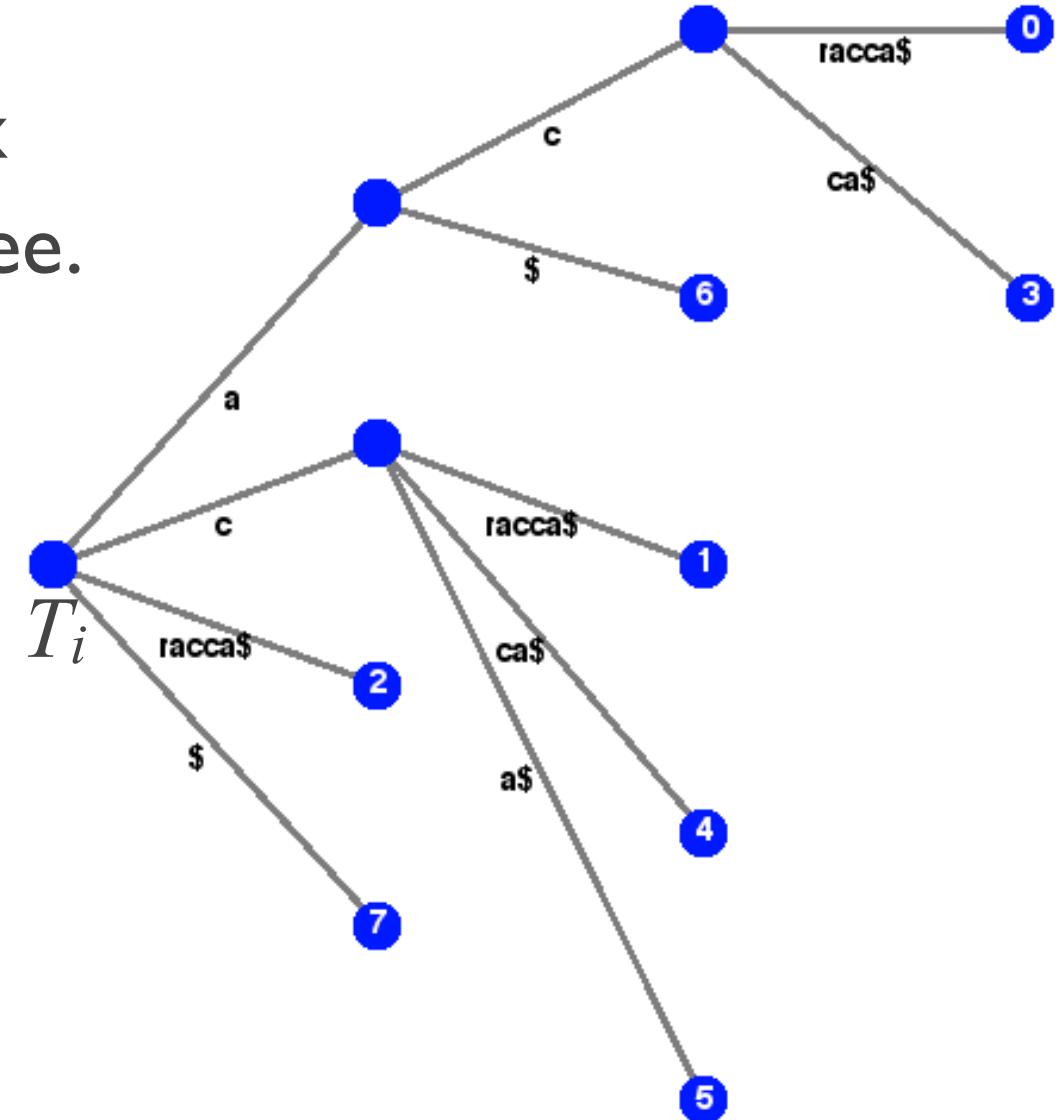
manana are suffixes
na of banana

Definition: Suffix tree

- Given $T=t_1t_2\dots t_n$ its suffix tree \mathcal{T}_T is a rooted tree.

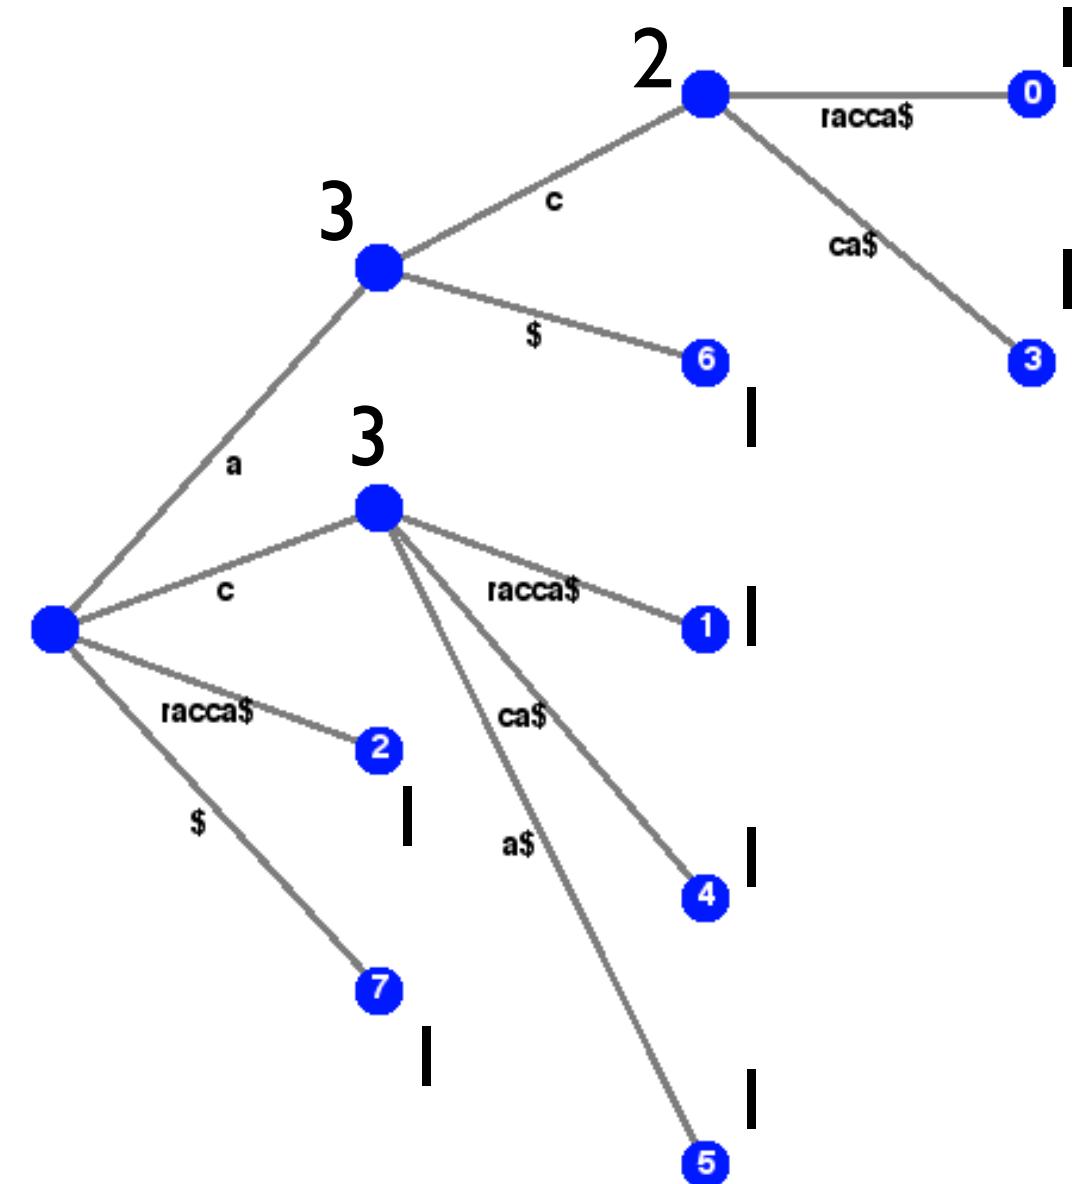
\mathcal{T}_T has

- n leaves; leaf l_i corresponds to suffix T_i
- edge label $\text{label}(u,v)$: sub-string
- Can be constructed in $O(|T|)$



Application: #occurrences of s in T?

- Use tree-traversal to populate leaf-counts for all internal nodes (once!)
 - If for some explicit or implicit vertex u , $L[u] = s$ then s is a substring and its count is given by u (or its least explicit predecessor)
 - Complexity: $O(|s|)$
 - Independent of $|T|$!



Outlook:

- Suffix tree for Human Genome ~40GB
- Burrows-Wheeler-Transform (BWT) and FM-index support same operations with same complexity but index is compressed to ~4 GB

Accelerating computations

Sketches

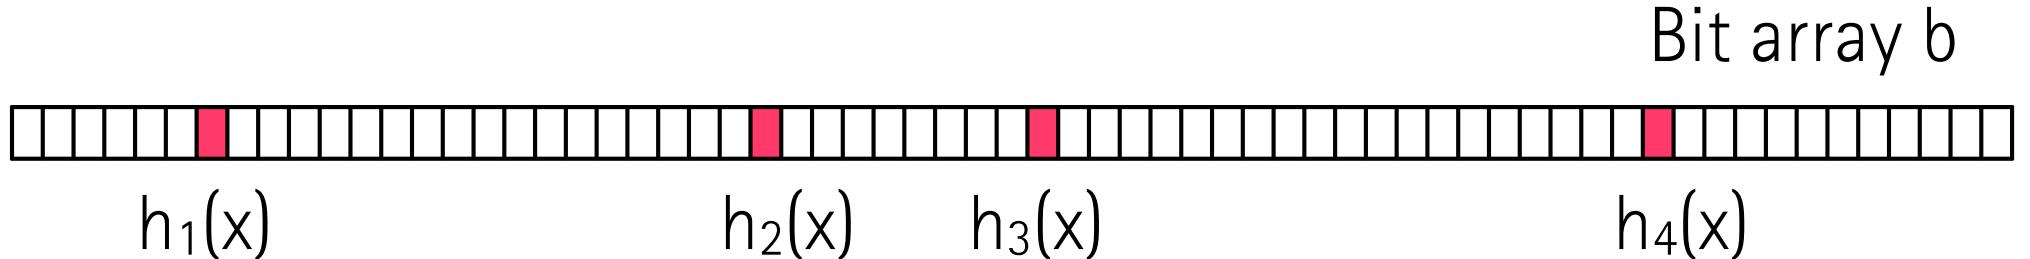
Counting observations

- Data: m elements from $\{1, 2, \dots, n\}$, $m \gg n$
- f_i frequency of i
- Query types:
 - Point query: frequency f_i of i
 - Range Query: for i, j estimate $f_i + f_{i+1} + \dots + f_j$
 - Quantile query: find i s.t. $f_1 + f_2 + \dots + f_i \approx x$
 - Heavy hitters: find i s.t. $f_i \geq m \times$

Subsampling

- Randomly select a small subset of the data
- Compute summary statistics on this

Bloom filter



How to change a Bloom Filter so that it can count?

Count-min sketch

Image of hash functions

w

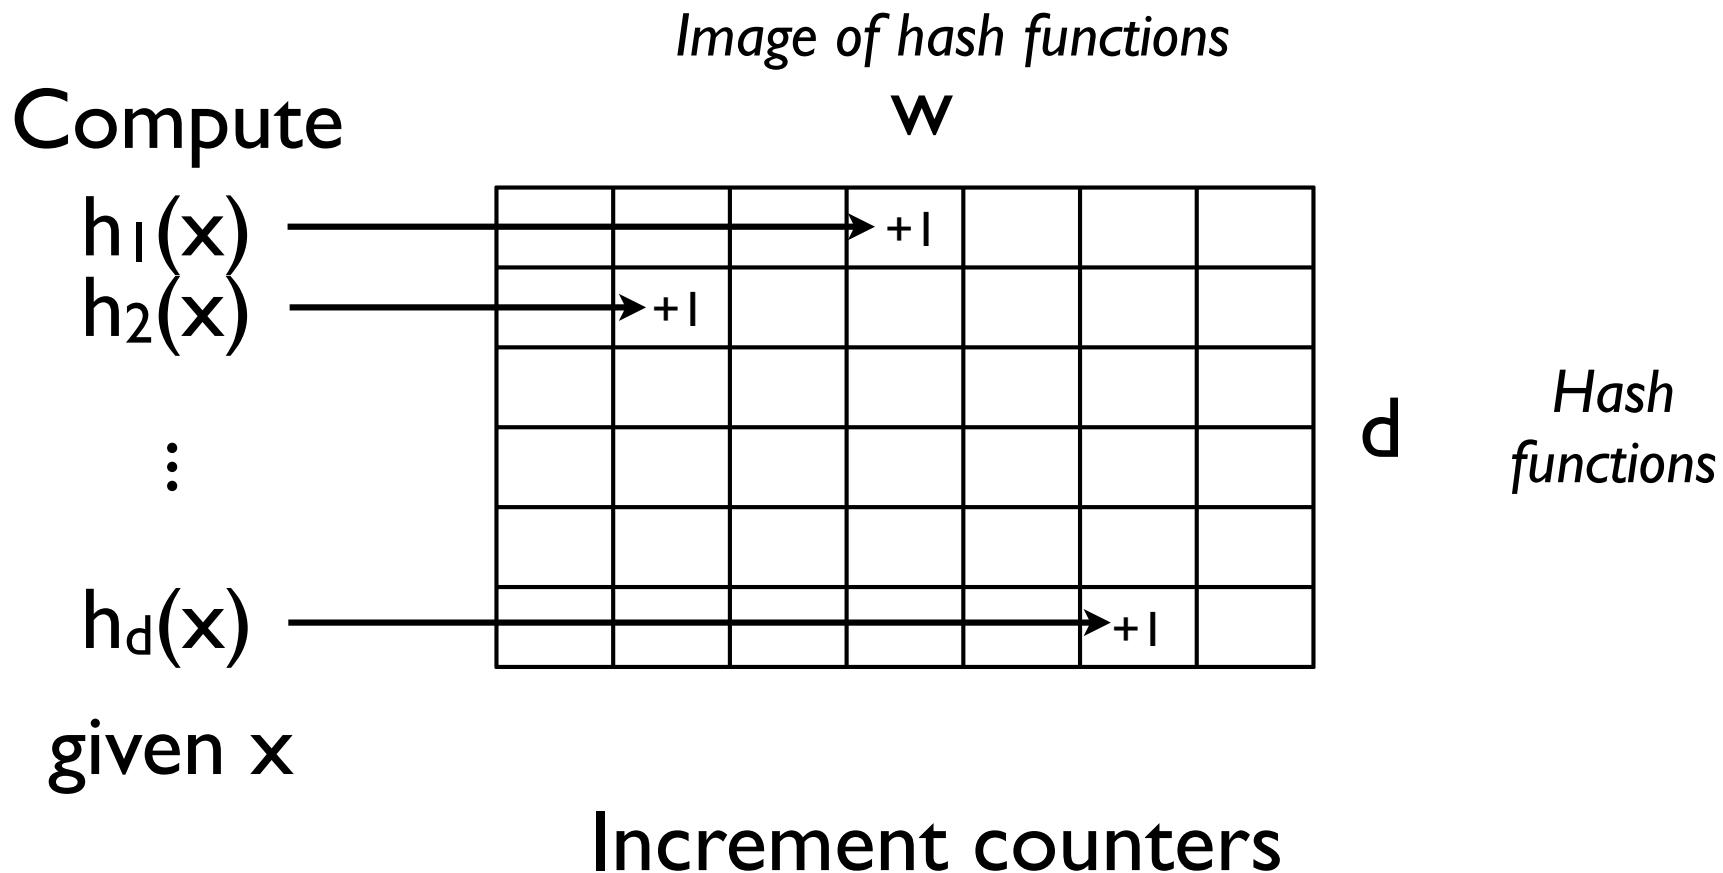
d

*Hash
functions*

$$c_{i,j} = \# \text{ of elements with } h_i(x) = j$$

e element in the data

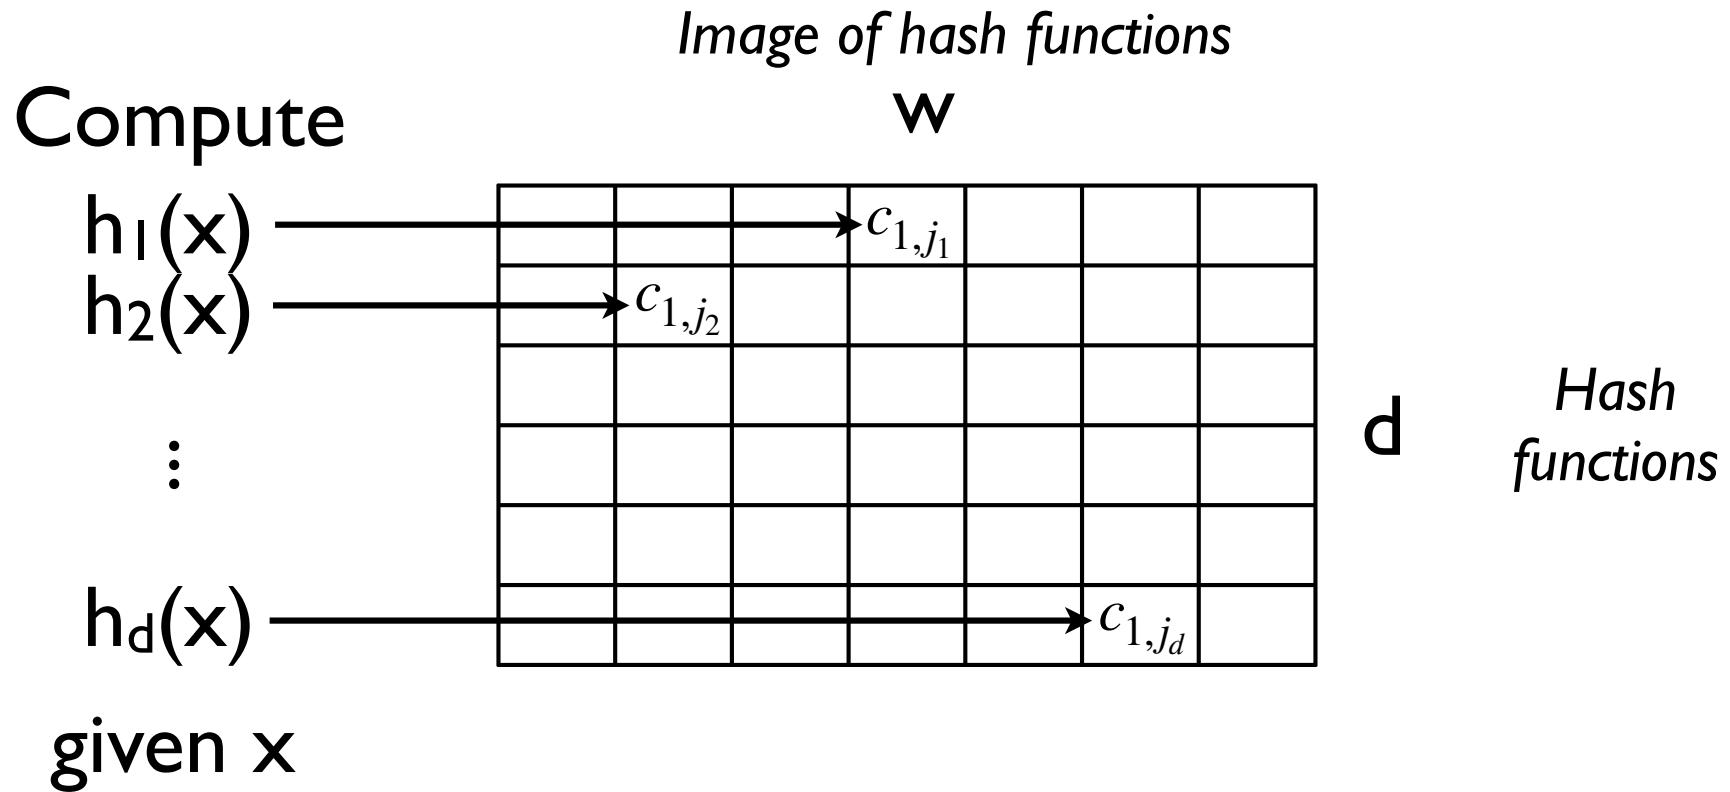
Count-min sketch: Update



Count min update can effectively be parallelized.

A	B	C	D	E
True	False			

Count-min sketch: Query



$$\tilde{f}_x = \min\{c_{1,j_1}, c_{2,j_2}, \dots, c_{d,j_d}\}$$

$$= \min\{c_{1,H_1(x)}, c_{2,H_2(x)}, \dots, c_{d,H_d(x)}\}$$

Count-min sketch:observations

- For all $e, i: c_{i,H_i(x)} \geq f_x$ so $\tilde{f}_x \geq f_x$
- $d=1, w=n, H_1(x) = x$: Exact counting using space linear in n
- Space: $d \times w \times \text{sizeof(counter)}$
- Increasing d ?
- Decreasing w ?

Count-min sketch: analysis

For $w = 2/\epsilon$ and $d = \log_2(1/\delta)$:

$$P[f_x \leq \tilde{f}_x \leq f_x + \epsilon m] \geq 1 - \delta$$

Need 2-way independent hash functions

$$P[h_i(x) = h_i(y)] \leq \frac{1}{w}$$

Count-min sketch: analysis

For $w = 2/\epsilon$ and $d = \log_2(1/\delta)$:

$$P[f_x \leq \tilde{f}_x \leq f_x + \epsilon m] \geq 1 - \delta$$

Error bound of estimate

Larger w , smaller ϵ

E.g. for error bound of
 $\epsilon m = 2000$ for $m = 10^6$
elements we need $w = 1,000$
at $\epsilon = 0.002$

Count-min sketch: analysis

For $w = 2/\epsilon$ and $d = \log_2(1/\delta)$:

$$P[f_x \leq \tilde{f}_x \leq f_x + \epsilon m] \geq 1 - \delta$$

Error bound of estimate
Larger w , smaller ϵ

E.g. for error bound of
 $\epsilon m = 2000$ for $m = 10^6$
elements we need $w = 1,000$
at $\epsilon = 0.002$

Probability of
staying within
error bound
Larger d ,
smaller δ

E.g. $d=7$ is needed
for $\delta = 0.01$

Memory usage: $w * d * \text{sizeof(counter)} = 1000 * 7 * 4 \text{ bytes} = 28\text{kB}$

Accelerating computations

Index data structures

NOTE: The kd-trees + k-means part is for reference only. I mostly skipped it in the lecture. See the original nice paper.

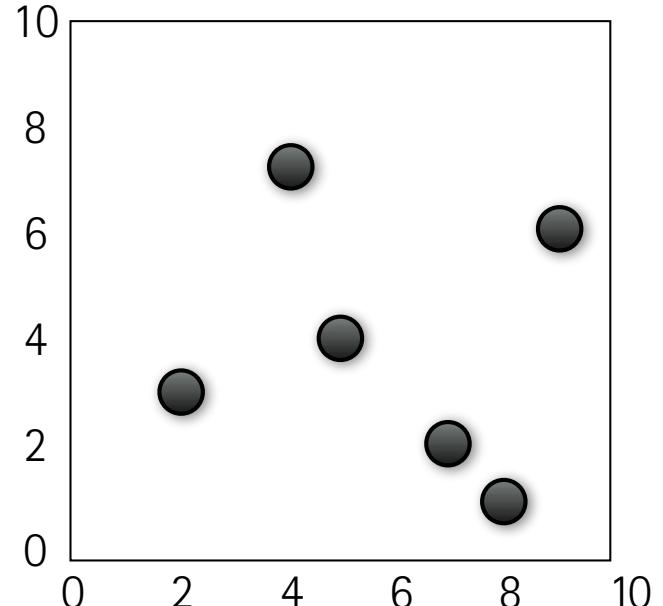
Pelleg, Moore: Accelerating exact k-means algorithms with geometric reasoning. KDD 1999.

Nearest neighbors

- Given a data base of n records with coordinates in Euclidean space (e.g. (x,z) or (x,y,z))
- For a query q , return the nearest neighbor in the database
- *Naive approach:* $O(n)$
- Application: GIS, computer games, dating websites/apps,...

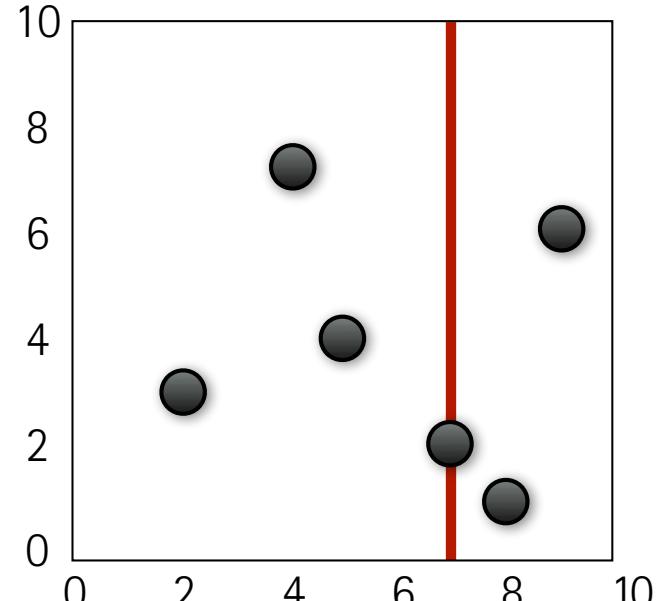
kd-trees

- Spatial index structure
- Balanced binary tree
- Fast nearest neighbor queries (dependent on instances)
- $O(n \log n)$ construction



kd-trees

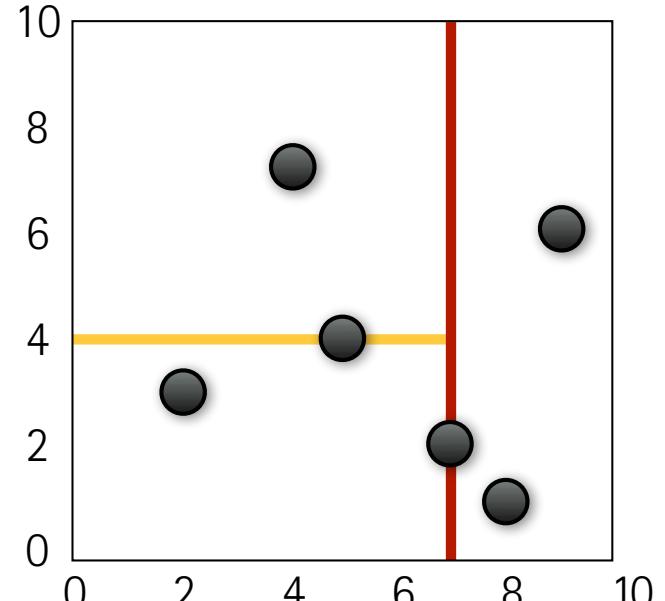
- Spatial index structure
- Balanced binary tree
- $O(\log n)$ nearest neighbor queries (dependent on instances)
- $O(n \log n)$ construction



(7,2)

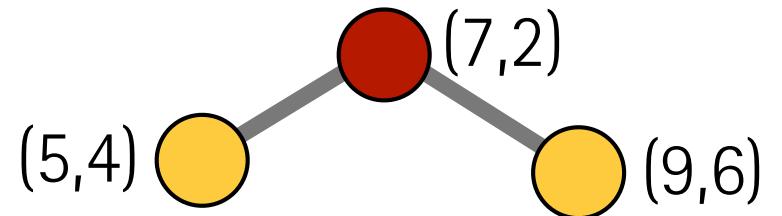
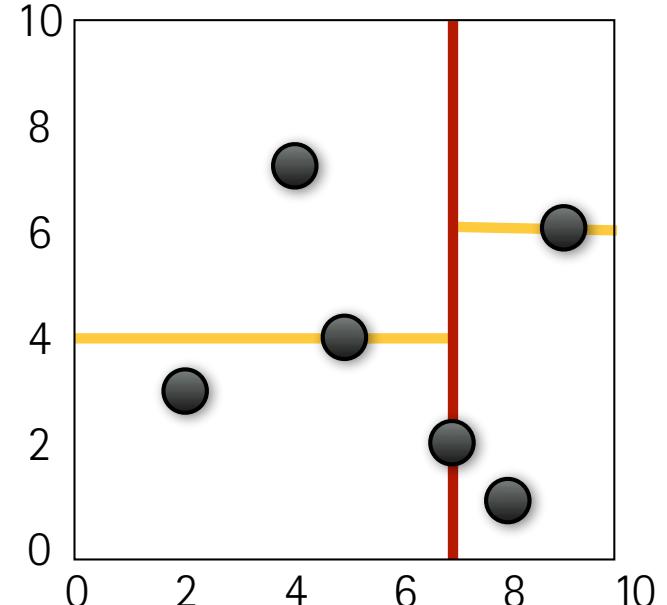
kd-trees

- Spatial index structure
- Balanced binary tree
- $O(\log n)$ nearest neighbor queries (dependent on instances)
- $O(n \log n)$ construction



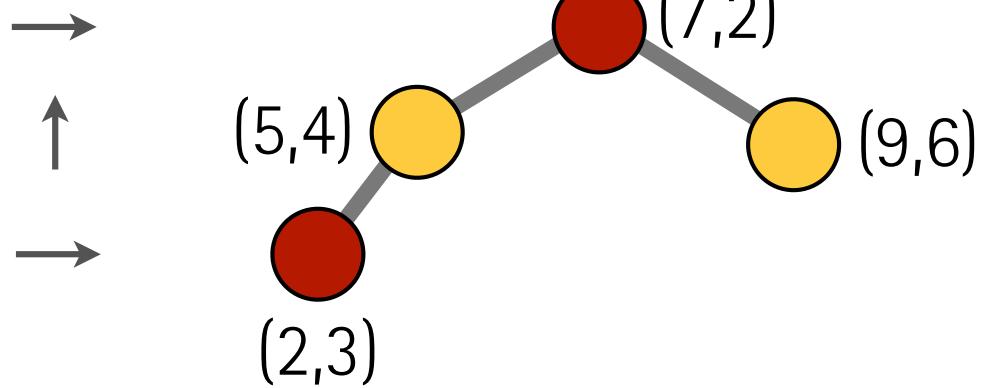
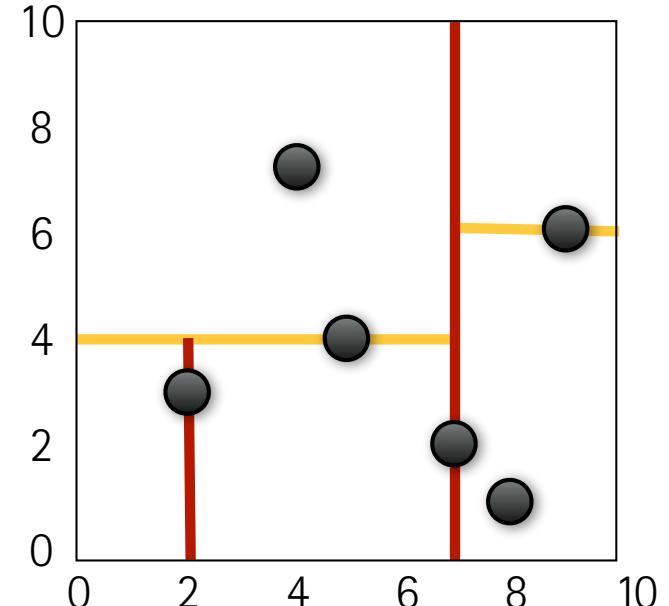
kd-trees

- Spatial index structure
- Balanced binary tree
- $O(\log n)$ nearest neighbor queries (dependent on instances)
- $O(n \log n)$ construction



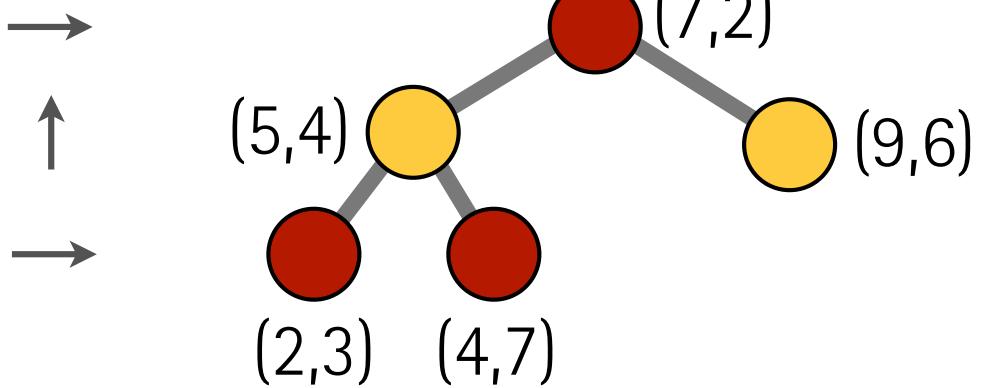
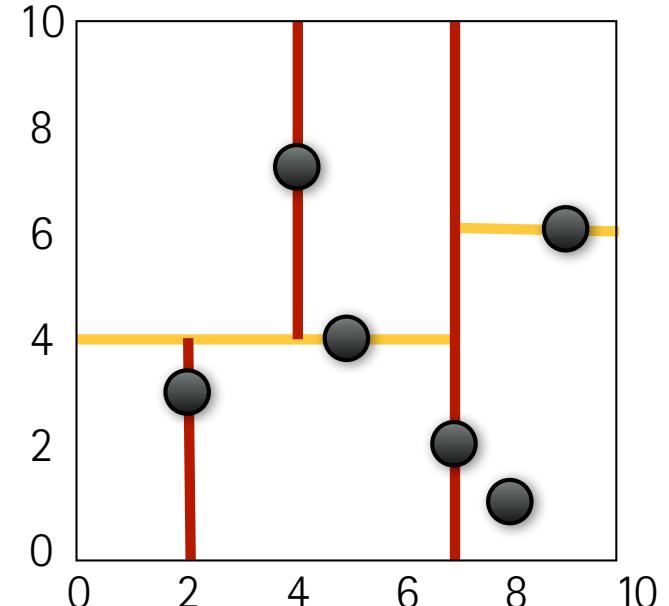
kd-trees

- Spatial index structure
- Balanced binary tree
- $O(\log n)$ nearest neighbor queries (dependent on instances)
- $O(n \log n)$ construction



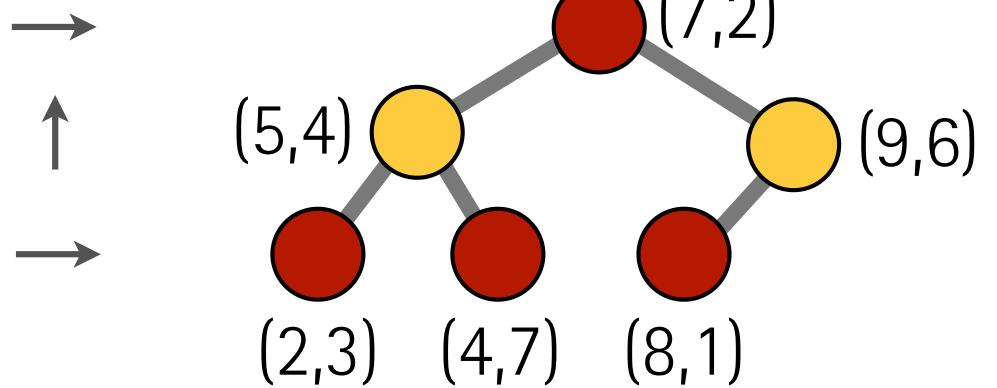
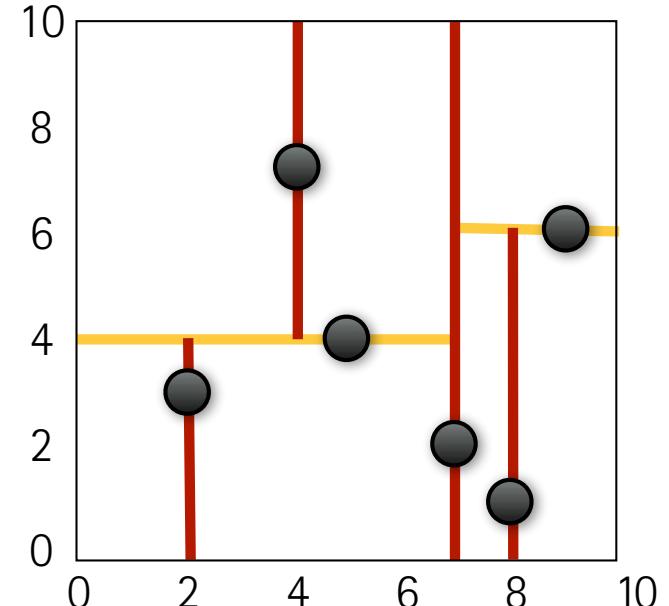
kd-trees

- Spatial index structure
- Balanced binary tree
- $O(\log n)$ nearest neighbor queries (dependent on instances)
- $O(n \log n)$ construction



kd-trees

- Spatial index structure
- Balanced binary tree
- $O(\log n)$ nearest neighbor queries (dependent on instances)
- $O(n \log n)$ construction



k-means clustering

$$\min_{C_1, C_2, \dots, C_k} \left\{ \sum_{k=1}^K W(C_k) \right\}$$

$$W(C_k) = \frac{1}{|C_k|} \sum_{i,j \in C_k} \|x_i - x_j\|^2$$

$\|x - y\|^2$ Euclidean distance

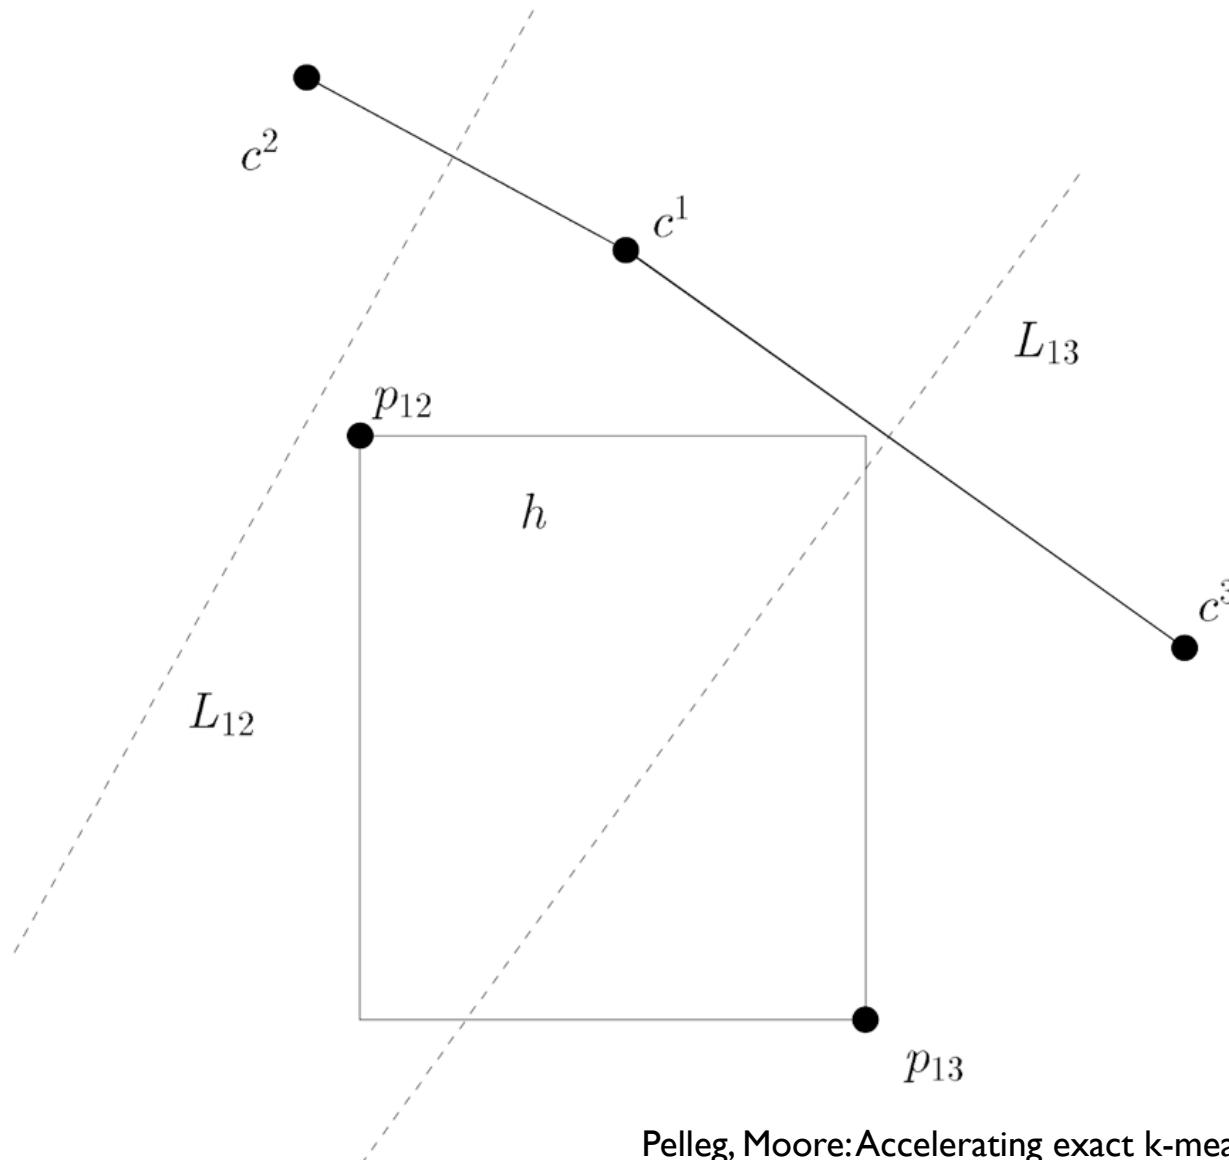
k-means clustering

- Choose k centroids among data points
- Iterate:
 - assign data points to closest centroid $O(n \times k)$
 - recompute centroids $O(n)$

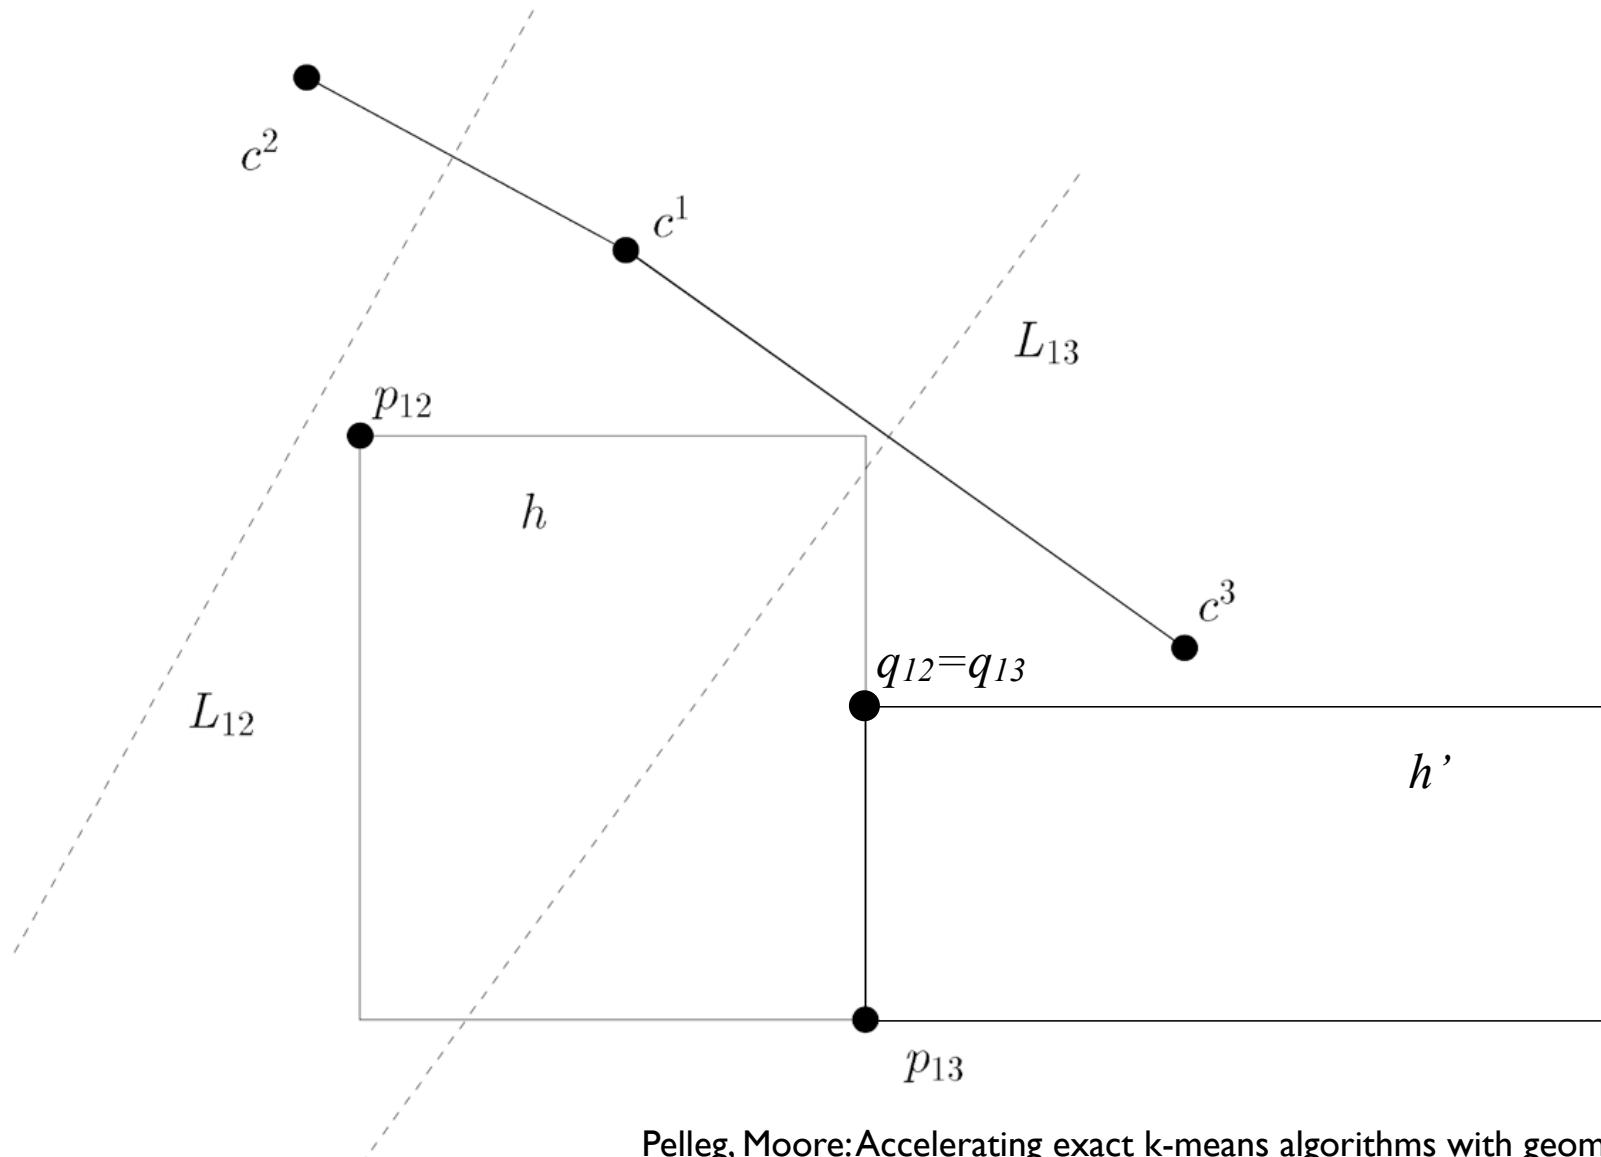
Accelerating k-means with kd-trees

- Idea: Avoid computing distance between every point and every centroid
- kd-trees partition space into hyper-rectangles

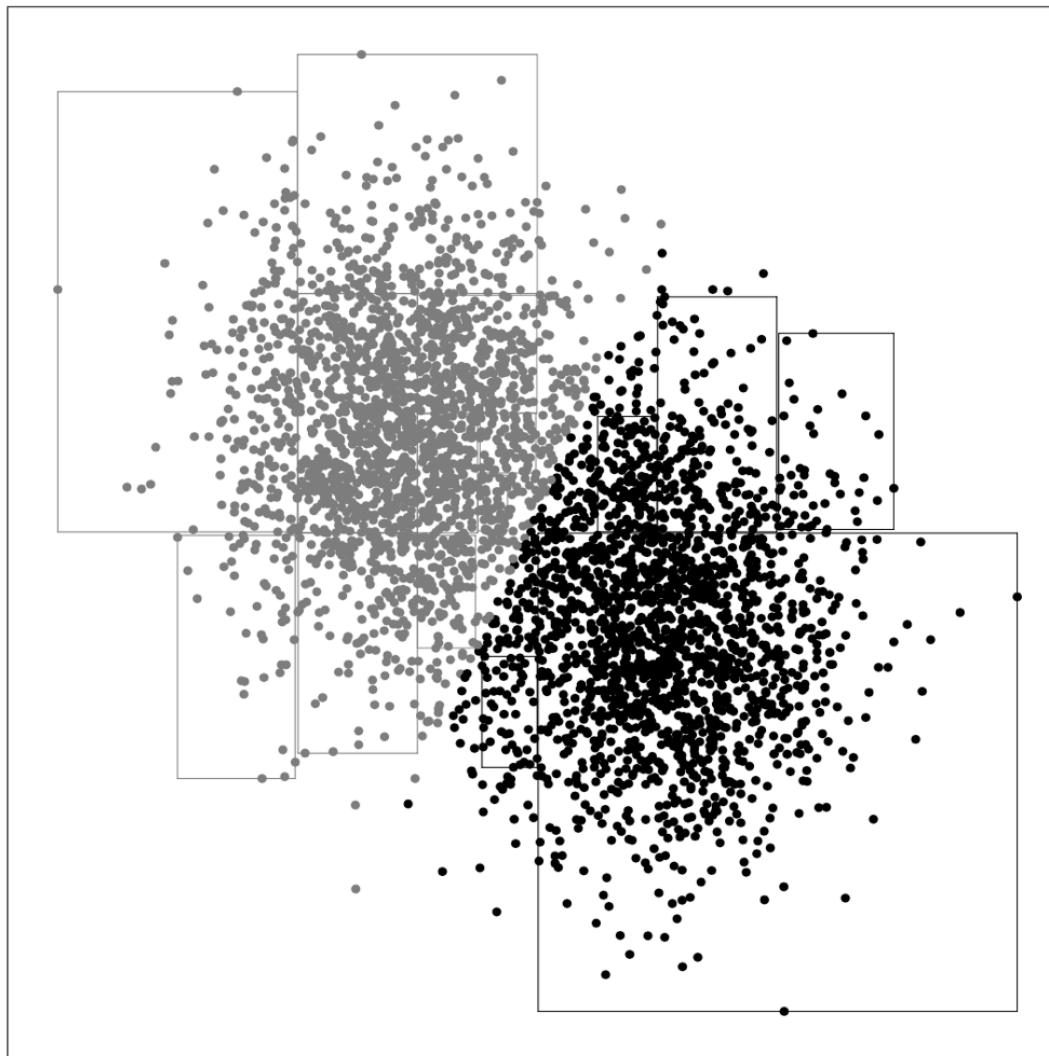
Accelerating k-means with kd-trees



Accelerating k-means with kd-trees



Accelerating k-means with kd-trees



Accelerating k-means with kd-trees

points	blacklisting	naive	speedup
50000	2.02	52.22	25.9
100000	2.16	134.82	62.3
200000	2.97	223.84	75.3
300000	1.87	328.80	176.3
433208	3.41	465.24	136.6

Table 1: Comparative results on real data.

Run-times of the naive and blacklisting algorithm, in seconds per iteration. Run-times of the naive algorithms also shown as their ratio to the running time of the blacklisting algorithm, and as a function of number of points. Results were obtained on random samples from the 2-D “petro” file using 5000 centers.

Use in methods

- For dimension up to 10
 - k-means
 - mixture estimation
 - k-nearest neighbor classification
- Active research topic: higher dimensions; specialization for specific combinations of n , k , d

Elkan, Using the triangle inequality to accelerate k-means, ICML, 2003

Kurban and Dalkilic: A novel approach to optimization of iterative machine learning algorithms: over heap structure. Big Data 2017

How to apply all of this ...?

Answer: reluctantly

General approach

- Implement Data Science solutions the framework you are most familiar with (e.g. Python / Pandas) without worrying about efficiency (or parallelism)
- Probably will be okay for < 100 GB on Laptop and < 1 TB of data on desktop
(assuming SSDs / RAID, enough RAM, multi-cores)

Cost of optimization

- Qualified technical employee: cost to employer 1.5M SEK/year for 1600 hours/year
⇒ ~1000 SEK/hour
- Most expensive Amazon EC2 on-demand instance is 250 SEK/hour (+storage).
- Not worth to invest an hour, if not saving at least four.
- Q: How frequently does a workload run?

What if my data is big?

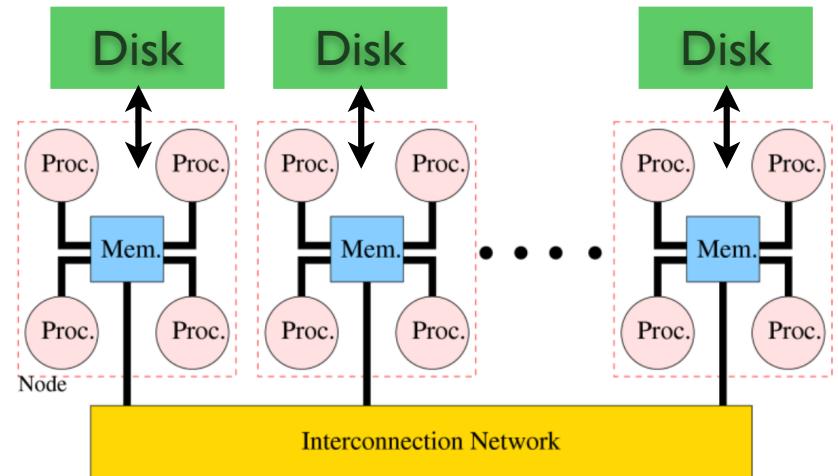
- Bandwidth:
 - Single SSD: 0.5 GB/s
 - Large RAID: 1 GB/s
 - DDR4 RAM > 20 GB/s (peak)
- Reading 1 (resp. 100) TB data needs:
 - 33 min resp. 5.5h
 - 17 min resp. 2.8 h
 - ~ 1 min resp. resp. 1:40 h

Is Input/Output the bottleneck?

- Times below are only a bottleneck if computations are fast
 - 33 min resp. 55h
 - 17 min resp. 28 h
 - ~ 1 min resp resp. 1:40 h
- Alternative: external-memory frameworks and algorithms.
 - E.g. <https://dask.org/> Pandas frames too large for RAM

What if my data is big?

- Solutions:
 - Up to a point expensive hardware (e.g. RAIDs with SSDs)
 - Cluster with local disks & HDFS
 - ⇒ You need to parallelize



What if my solution is too slow?

- General ideas:
 - Change the problem
 - Refactor
 - Compile
 - Optimize
 - Parallelize

Change the problem

- E.g. Descriptive statistics with MapReduce:
Is the median needed?
- Identifying Duplicates: Are a few false positive duplicates acceptable?
- If a median (or quartiles are needed): are approximate estimates acceptable?
- If a classifier is used: Is a simpler algorithm (e.g. k-nearest neighbor instead of deep learning) acceptable?

Refactor

- Generally: short Python programs are fast Python programs.
- Use libraries: E.g. numpy and others offering C/C++-speed

```
def nearestCentroid(datum, centroids):  
    # norm(a-b) is Euclidean distance, matrix - vector computes difference  
    # for all rows of matrix  
    dist = np.linalg.norm(centroids - datum, axis=1)  
    return np.argmin(dist), np.min(dist)
```

- Might make some parallelization trivial

```
import mkl  
import numpy as np  
  
mkl.set_num_threads(4)  
np.dot(x,y)
```

Compile your Python Code

Cython

<https://cython.org>

```
def primes(int nb_primes):
    cdef int n, i, len_p
    cdef int p[1000]
    if nb_primes > 1000:
        nb_primes = 1000

    len_p = 0 # The current number of elements in p.
    n = 2
    while len_p < nb_primes:
        # Is n prime?
        for i in p[:len_p]:
            if n % i == 0:
                break

        # If no break in the loop, we have a prime.
        else:
            p[len_p] = n
            len_p += 1
            n += 1

    # Let's return the result in a python list:
    result_as_list = [prime for prime in p[:len_p]]
    return result_as_list
```

Some parallelization support OpenMP (not clusters though)
Allows calling of C libraries, use of C-structs etc.

Numba

<https://numba.pydata.org/>

```
from numba import jit
import random

@jit(nopython=True)
def monte_carlo_pi(nsamples):
    acc = 0
    for i in range(nsamples):
        x = random.random()
        y = random.random()
        if (x ** 2 + y ** 2) < 1.0:
            acc += 1
    return 4.0 * acc / nsamples
```

```
@numba.jit(nopython=True, parallel=True)
def logistic_regression(Y, X, w, iterations):
    for i in range(iterations):
        w -= np.dot(((1.0 /
                      (1.0 + np.exp(-Y * np.dot(X, w) -
                                     1.0) * Y)), X))
    return w
```

Also parallelizes incl. GPUs (not clusters though)

Optimize

- Measure running time and identify hot spots
- Better data structures and algorithms usually beats code optimization
 - Constants matter
 - Latency/memory hierarchy matters
 - ⇒ Lower computational complexity does not always win
- Sometimes optimizing makes sense

Parallelize

- Based on structure of workload: data flow, computational effort, hardware ...
- Measure Amdahl's f
- SIMD instructions
- Multi-threading
- Message passing
- MapReduce, Spark, ... and many big data computational frameworks.
- Question: Are there specialized frameworks for your problem?
- Avoid oversubscription ...

Parallelize

- Achieve scaling via:
 - Dedicated hardware: GPU, Tensor Units for deep learning
 - Clusters
 - Clouds ...
 - *Volunteer Clouds (Folding at home...)*

Outlook

- Very active research field
 - Intel: <https://github.com/IntelPython>,
<https://www.youtube.com/watch?v=HKjM3pelNtw> (SciPy 2018 talk)
 - Amazon–Uber: <https://arrow.apache.org/>;
see <https://streamdata.io/blog/open-source-apache-big-data-projects/>
 - Tool provider: Anaconda (e.g. Numba)
with Intel, Nvidia and AMD

Q&A from Chat

Q: Where to find hash functions? Defined your own? A: Good hash functions are a science and art in itself. I always took the state-of-the-art for a specific task from the literature. Note that there is a tradeoff between speed and quality.



Tupelo Honey

Cassandra Wilson – Closer To You: The Pop Side



Axis: Bold As Love

Joan Osborne – How Sweet It Is



Tender Love

Me'Shell Ndegéocello – Ventriloquism



Speak Like A Child

The Style Council – Introducing The Style Council



Gradual Return

Trixie Whitley – Fourth Corner



Everybody Is A Star

Joan Osborne – How Sweet It Is